



# Microsoft Azure Training Day: Data and Analytics

## Azure Cosmos DB

Marco Parenzan (twitter: @marco\_parenzan)  
Solution Sales Specialist @ Insight (email: marco.parenzan@insight.com)  
Microsoft Azure MVP (mvpid: 5000823)  
Community Lead @ 1nn0va (url: <https://www.facebook.com/1nn0va>)



# WHAT IS AZURE COSMOS DB

A globally distributed, massively scalable, multi-model database service



Table API



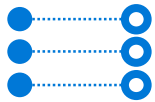
SQL



Cosmos DB's API for MongoDB



**Gremlin**  
 $G = (V, E)$



Key-value



Column-family



Document



Graph

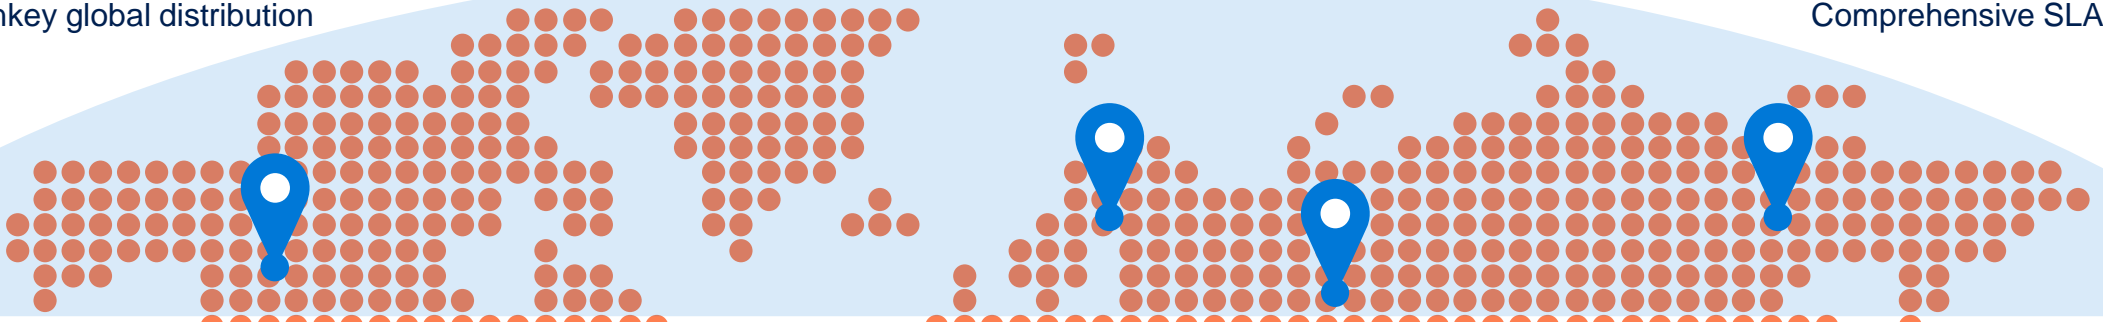
Elastic scale out  
of storage & throughput

Guaranteed low latency at the 99<sup>th</sup> percentile

Five well-defined consistency models

Turnkey global distribution

Comprehensive SLAs

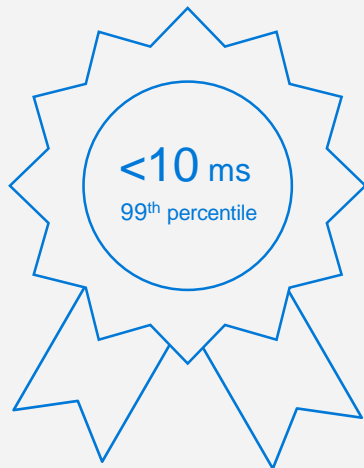


# Comprehensive SLA's

Run your app on world-class infrastructure.

Azure Cosmos DB is the only service with financially-backed SLAs for millisecond latency at the 99th percentile, 99.999% HA and guaranteed throughput and consistency

## Latency



## HA



## Throughput



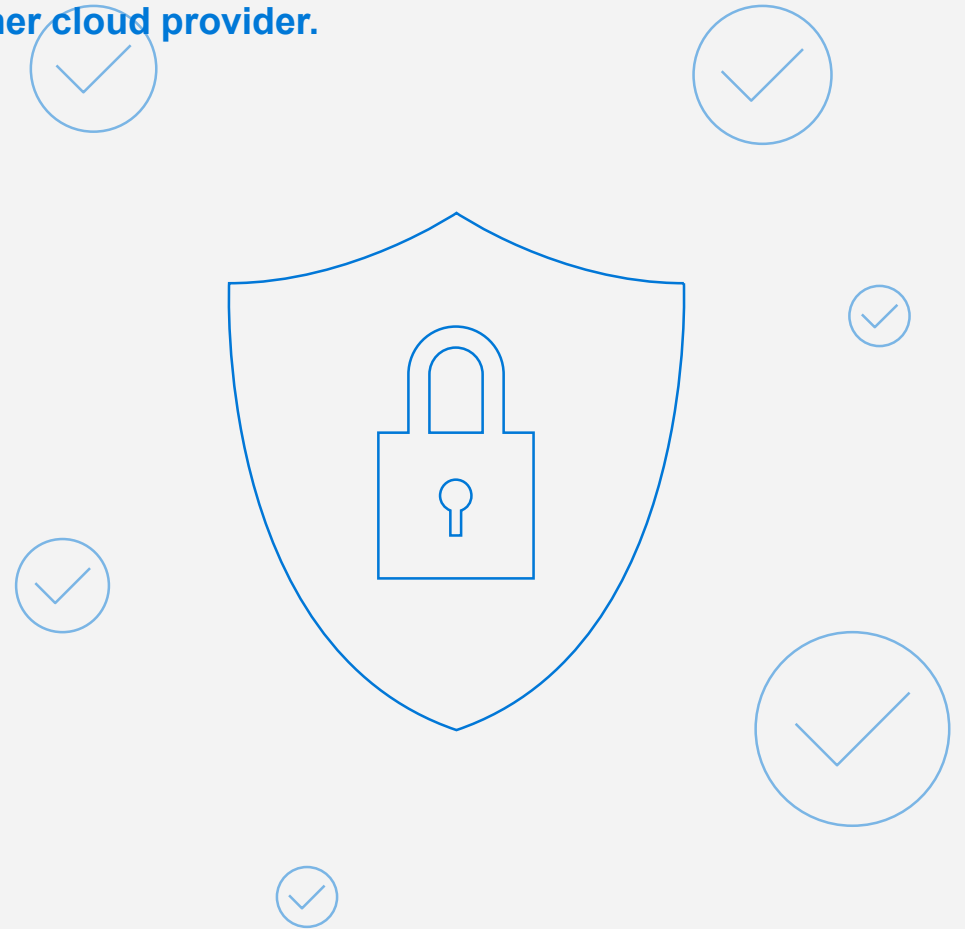
## Consistency



# Trust your Data to Industry-Leading Security & Compliance

**Azure is the world's most trusted cloud, with more certifications than any other cloud provider.**

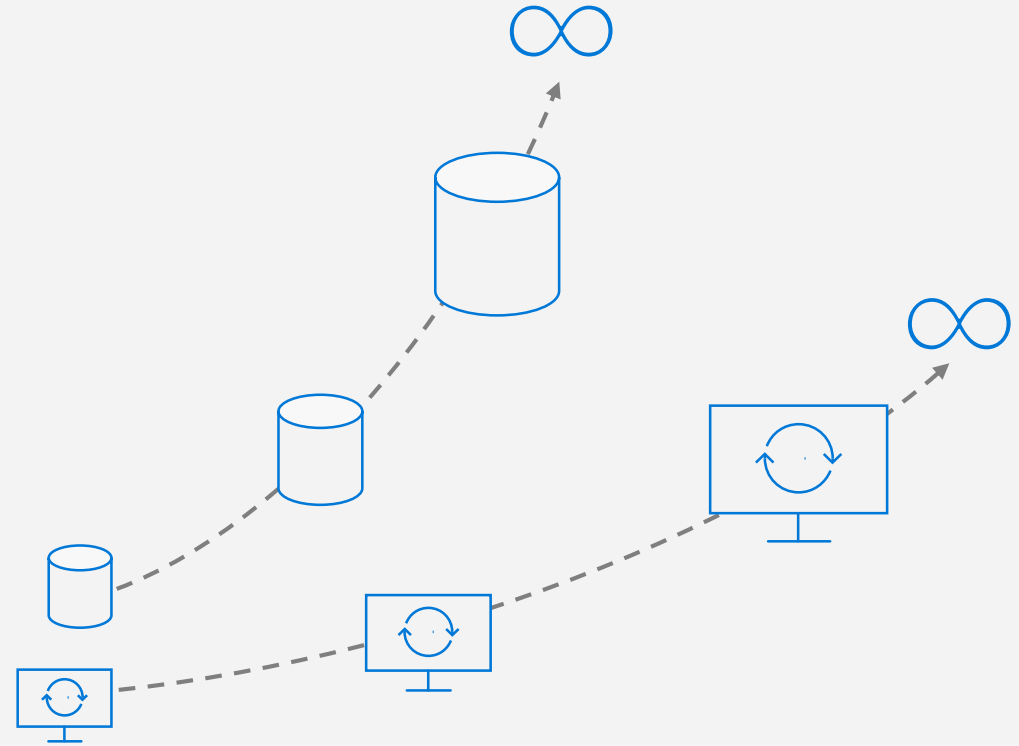
- Enterprise grade security
- Encryption at Rest and Transit
- Encryption is enabled automatically by default
- Comprehensive Azure compliance certification



# Elastically Scale Storage and Throughput

**Independently and elastically scale storage and throughput across regions – even during unpredictable traffic bursts – with a database that adapts to your app's needs.**

- Elastically scale throughput from 10 to 100s of millions of requests/sec across multiple regions
- Support for requests/sec for different workloads
- Pay only for the throughput and storage you need

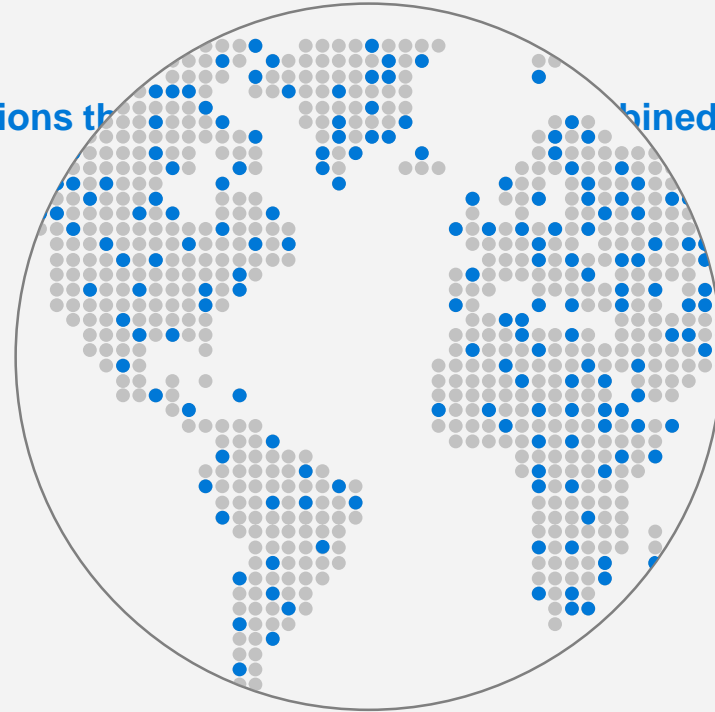


# Turnkey Global Distribution

Put your data where your users are in minutes

Automatically replicate all your data around the world, and across more regions than you can imagine.

- Available in [all Azure regions](#)
- Manual and automatic failover
- Automatic & synchronous multi-region replication
- Configure multiple write regions to further reduce latency and increase availability

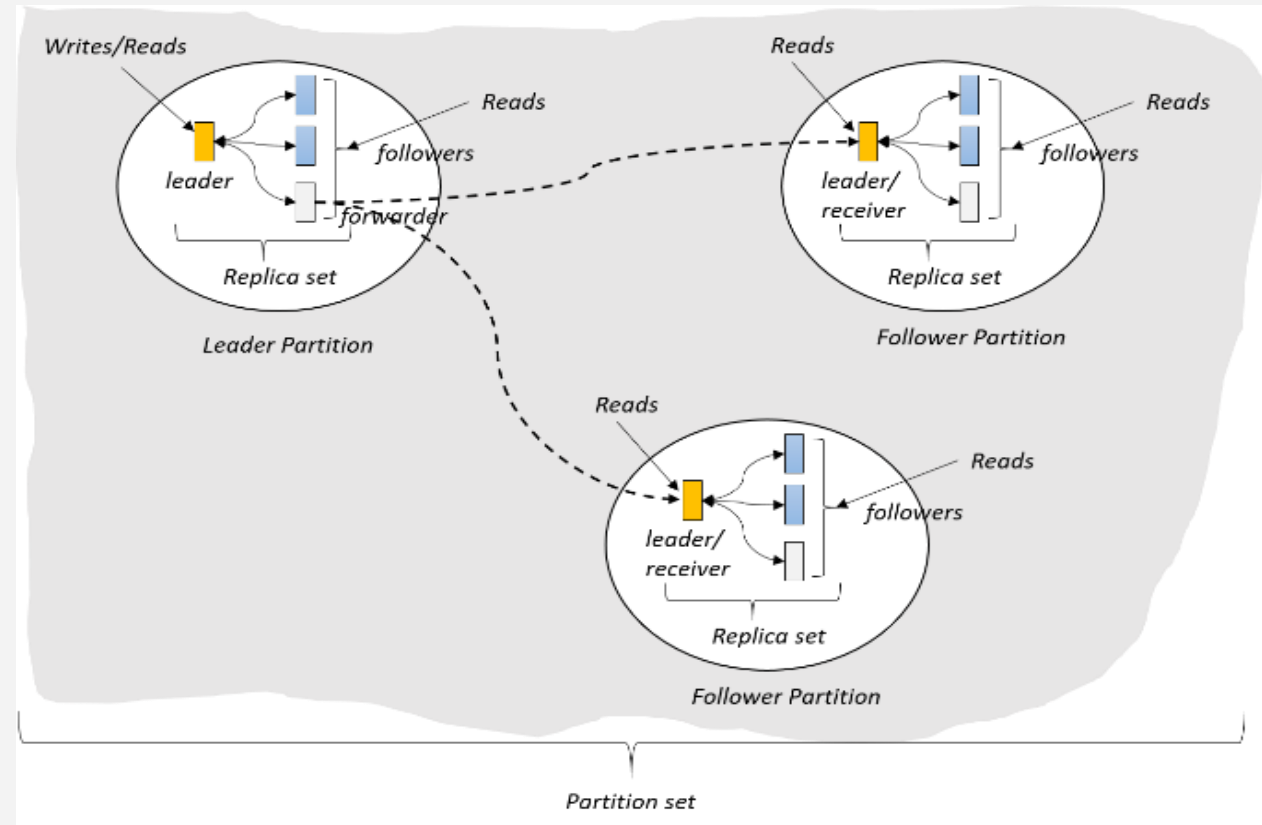


# CONSISTENCY

## BREWER'S CAP THEOREM

Impossible for distributed data store to simultaneously provide more than 2 out of the following 3 guarantees:

- Consistency
- Availability
- Partition Tolerance



# Five Well-Defined Consistency Models

## Choose the best consistency model for your app

Offers five consistency models

Provides control over performance-consistency tradeoffs, backed by comprehensive SLAs.

An intuitive programming model offering low latency and high availability for your planet-scale app.



**Strong**



**Bounded-stateless**



**Session**



**Consistent prefix**



**Eventual**





# Multiple Data Models and API's

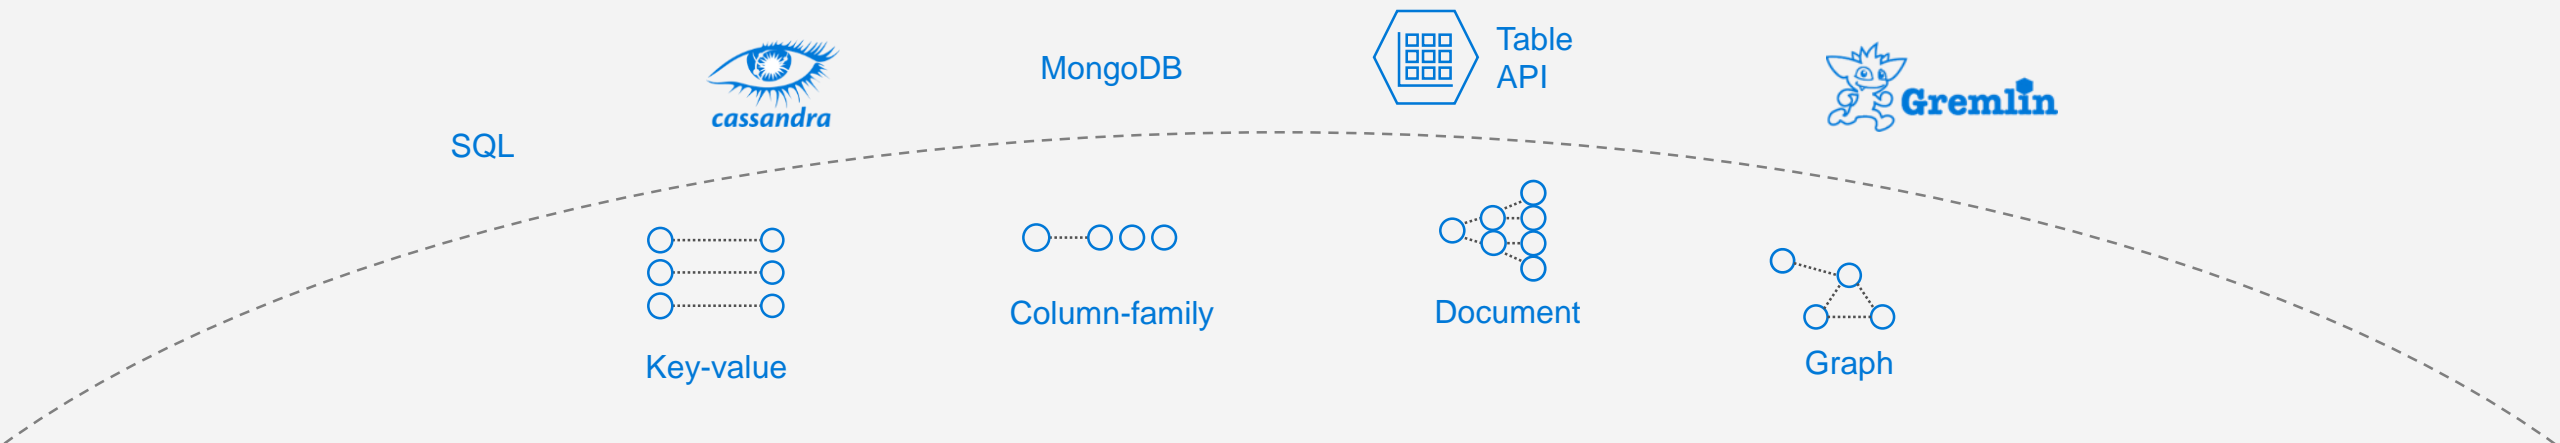
**Use the model that fits your requirements, and the apis, tools, and frameworks you prefer**

Cosmos DB offers a multitude of APIs to access and query data including, SQL, various popular OSS APIs, and native support for NoSQL workloads.

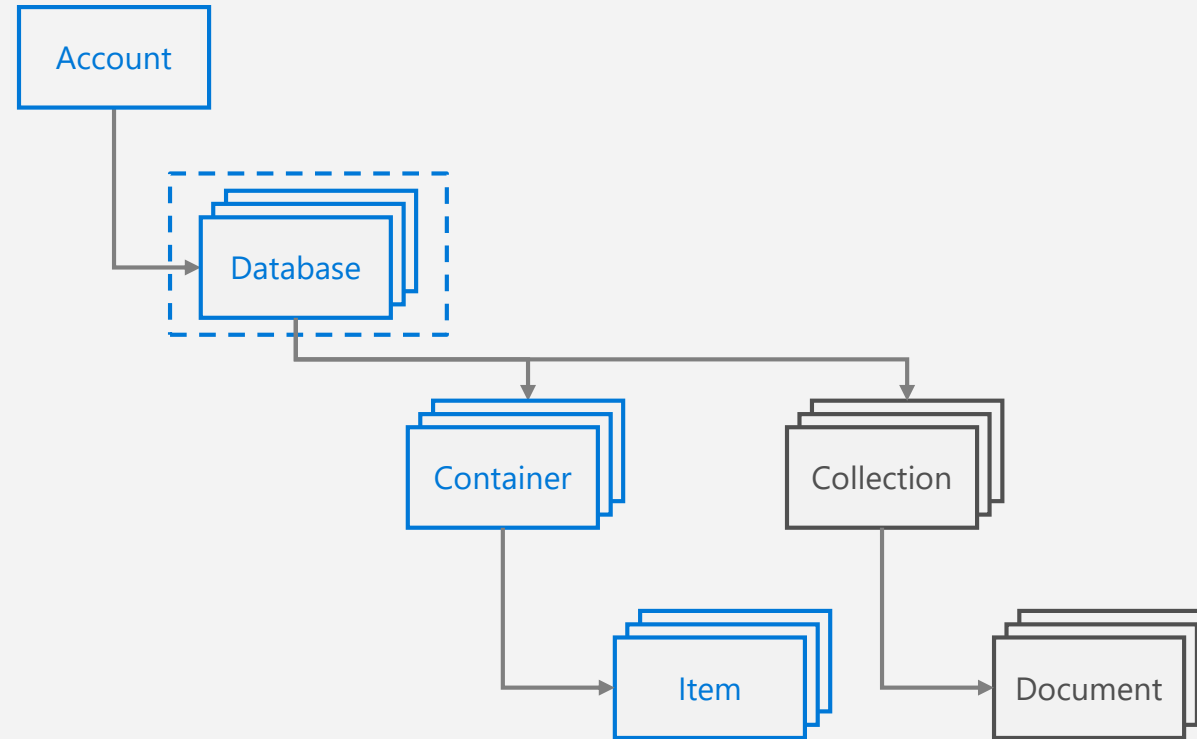
Use key-value, columnar, graph, and document data

Data is automatically indexed, with no schema or secondary indexes required

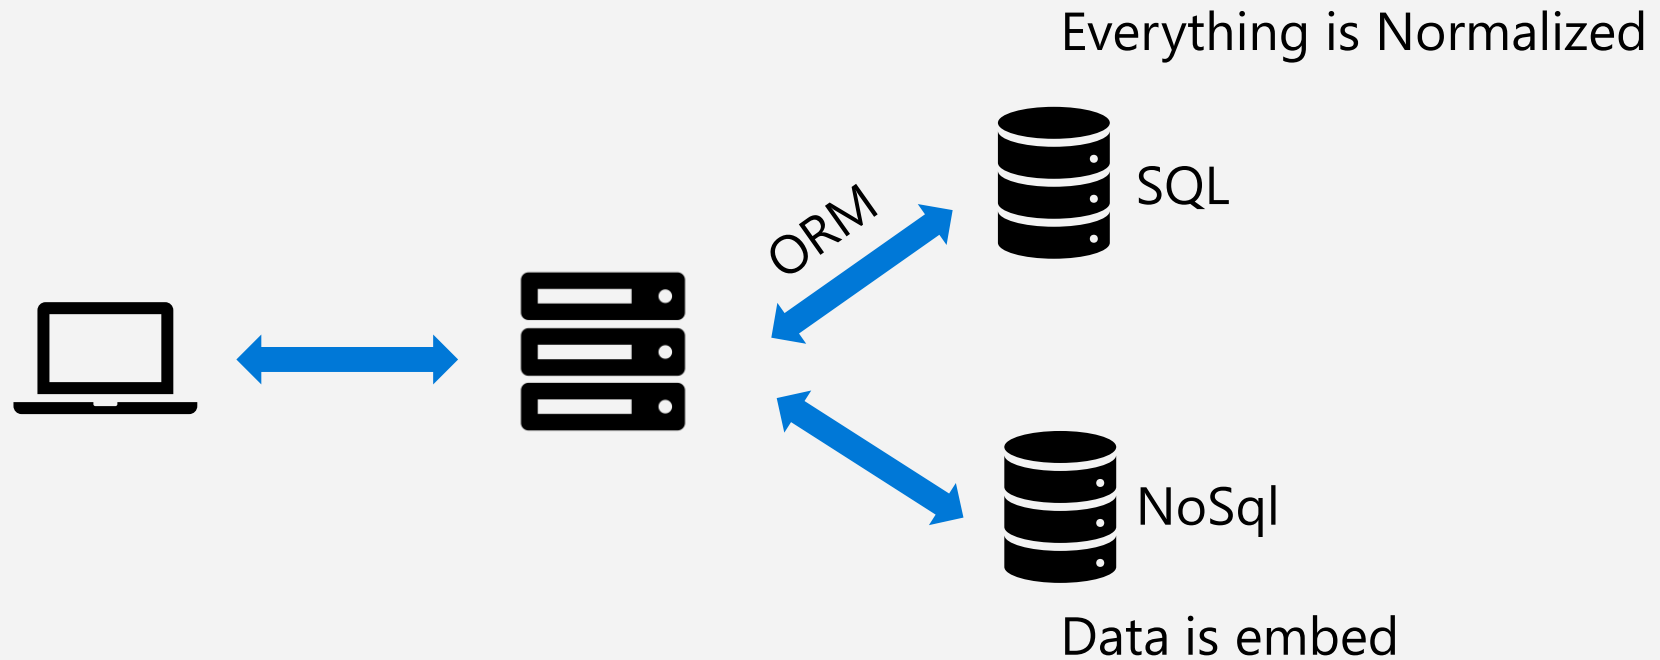
Blazing fast queries with no lag



# RESOURCE MODEL

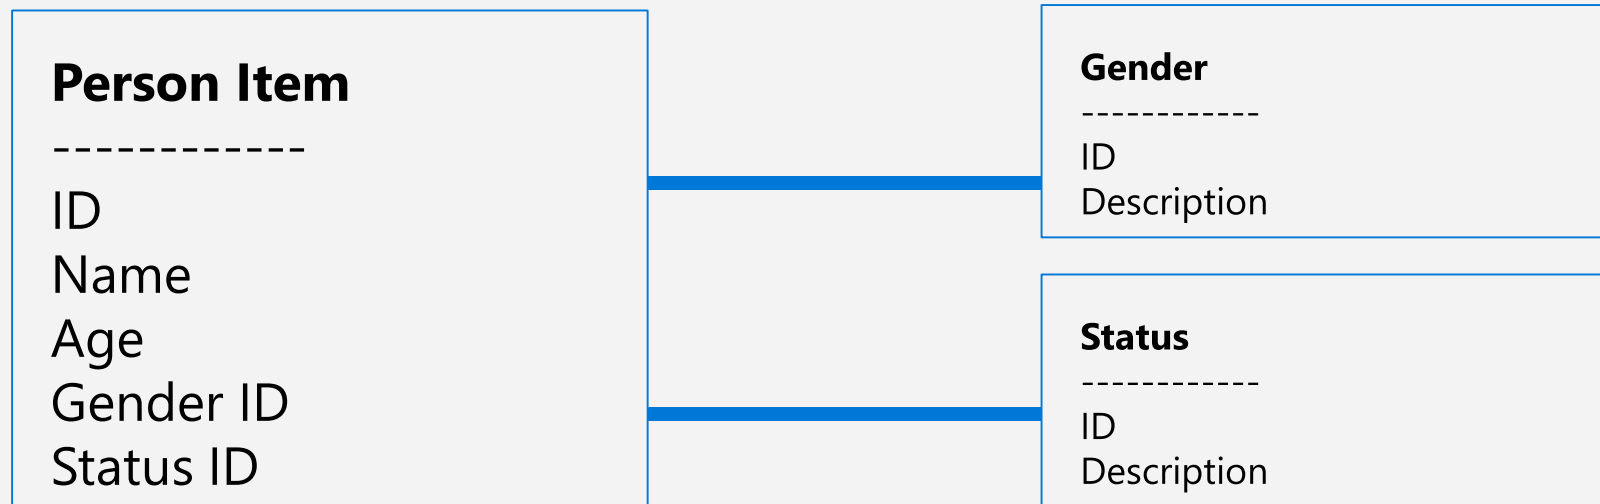


# EMBED vs REFERENCE



# EMBED vs REFERENCE

## Relational modelling



## Non-relational modeling

```
{
  "ID": 1,
  "Name": "John",
  "Age": "42",
  "GenderID": 1,
  "GenderDescription": "male"
  "StatusID": 2,
  "StatusDescription": "married"
}
```

# EMBED vs REFERENCE

When to embed?

When to reference?

When to put different types in the same collection?

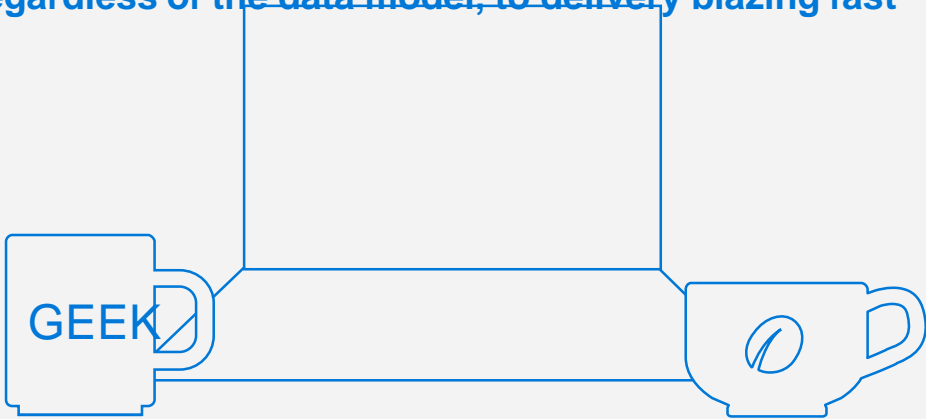
Sometimes the best choice is to do both

- Reference less frequently used data
- Embed more frequently used data

# Handle any Data with no Schema or Indexing Required

Azure Cosmos DB's schema-less service automatically indexes all your data, regardless of the data model, to delivery blazing fast queries.

- Automatic index management
- Synchronous auto-indexing
- Freedom from schema + index management
- Works across every data model
- Ingest and serve data back out in milliseconds



Item	Color	Microwave safe	Liquid capacity	CPU	Memory	Storage
Geek mug	Graphite	Yes	16ox	???	???	???
Coffee Bean mug	Tan	No	12oz	???	???	???
Surface book	Gray	???	???	3.4 GHz Intel Skylake Core i7-6600U	16GB	1 TB SSD

# Index POLICIES

## CUSTOM INDEXING POLICIES

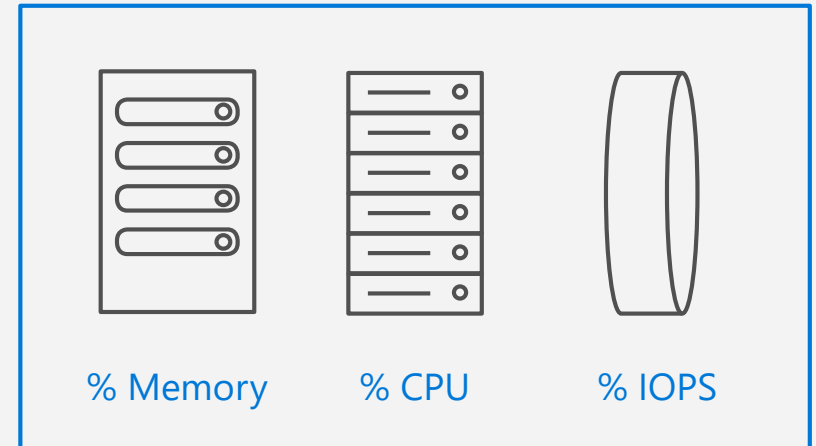
Though all Azure Cosmos DB data is indexed by default, you can specify a custom indexing policy for your collections. Custom indexing policies allow you to design and customize the shape of your index while maintaining schema flexibility.

- Define trade-offs between storage, write and query performance, and query consistency
- Include or exclude documents and paths to and from the index
- Configure various index types

```
{
  "automatic": true,
  "indexingMode": "Consistent",
  "includedPaths": [{
    "path": "/*",
    "indexes": [{
      "kind": "Hash",
      "dataType": "String",
      "precision": -1
    }, {
      "kind": "Range",
      "dataType": "Number",
      "precision": -1
    }, {
      "kind": "Spatial",
      "dataType": "Point"
    }
  ]
}, {
  "excludedPaths": [{
    "path": "/nonIndexedContent/*"
  }]
}
```

# Request Units

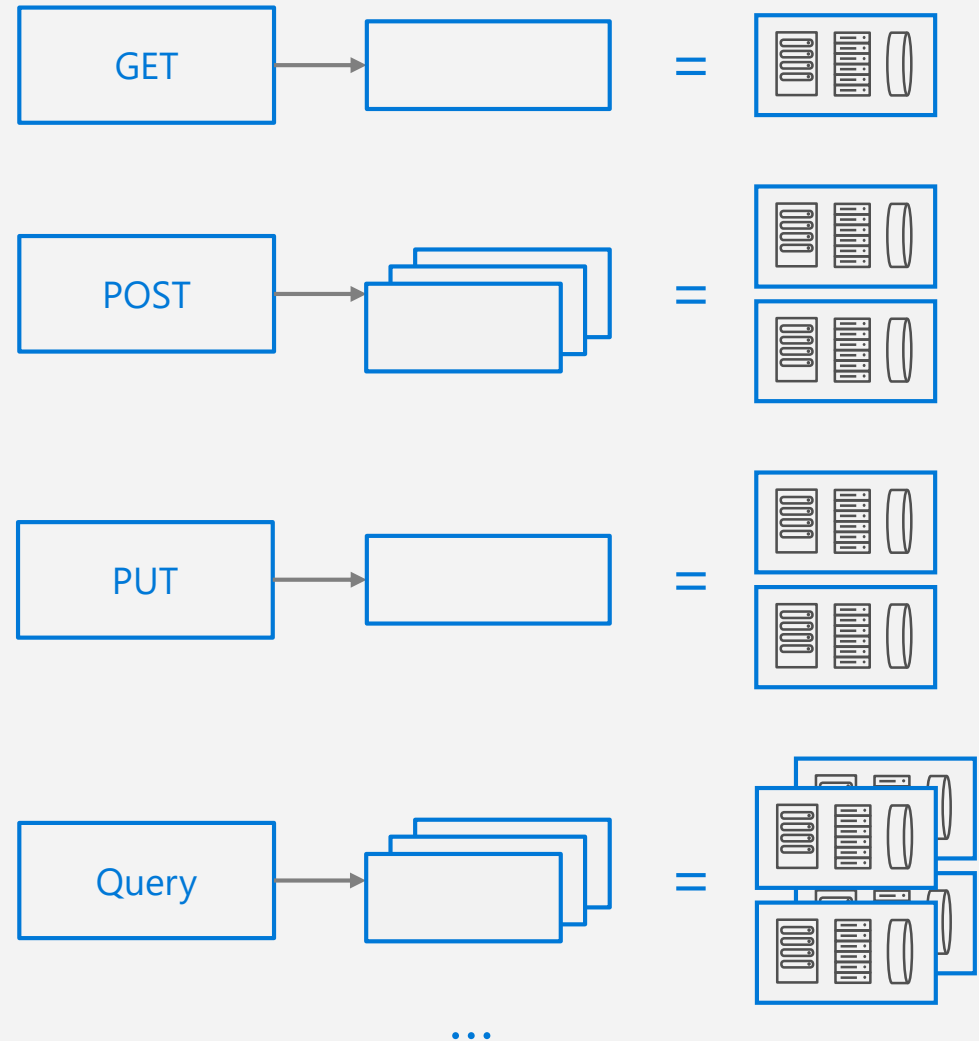
- Request Units (RUs) is a rate-based currency
- Abstracts physical resources for performing requests
- Key to multi-tenancy, SLAs, and COGS efficiency
- Foreground and background activities





# Request Units

- Normalized across various access methods
- 1 RU = 1 read of 1 KB document
- Each request consumes fixed RUs
- Applies to reads, writes, query, and stored procedures



# Request Unit Pricing Example

## Storage Cost

Avg Record Size (KB)	1
Number of Records	100,000,000
Total Storage (GB)	100
Monthly Cost per GB	\$0.25
Expected Monthly Cost for Storage	\$25.00

## Throughput Cost

Operation Type	Number of Requests per Second	Avg RU's per Request	RU's Needed
Create	100	5	500
Read	400	1	400

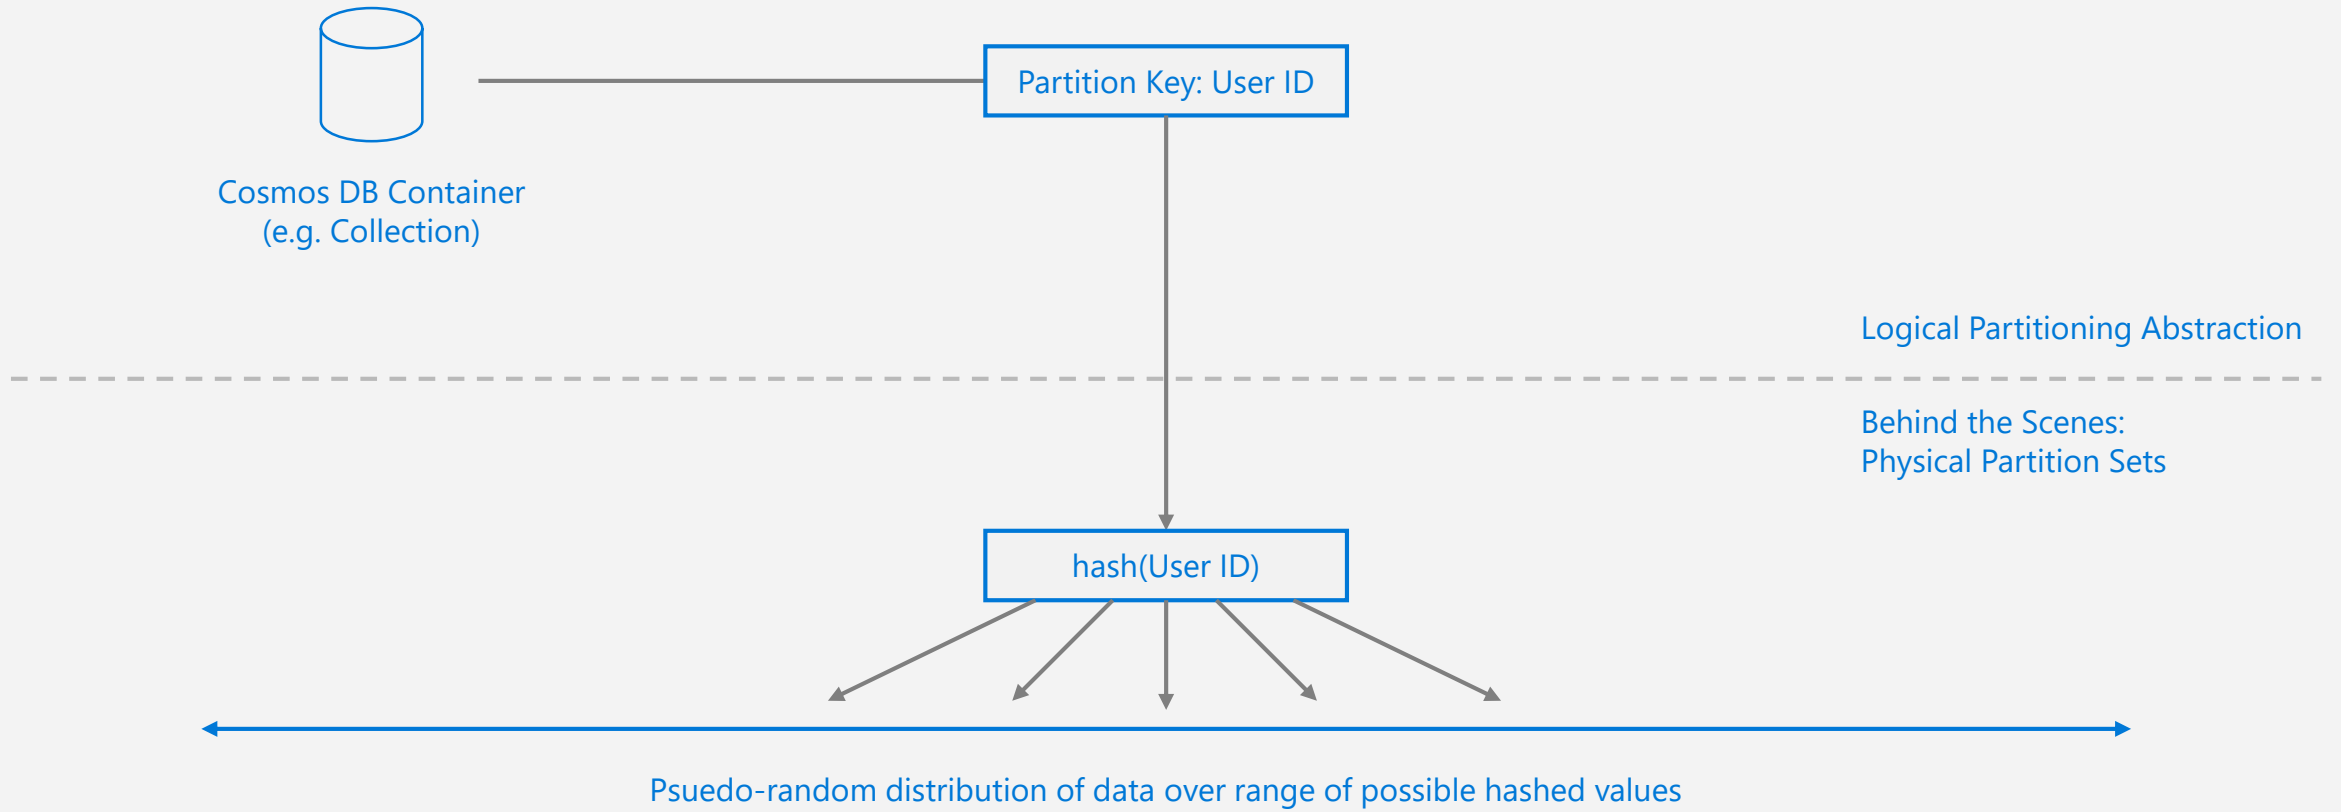
Total RU/sec	900
Monthly Cost per 100 RU/sec	\$6.00
Expected Monthly Cost for Throughput	\$54.00

## Total Monthly Cost

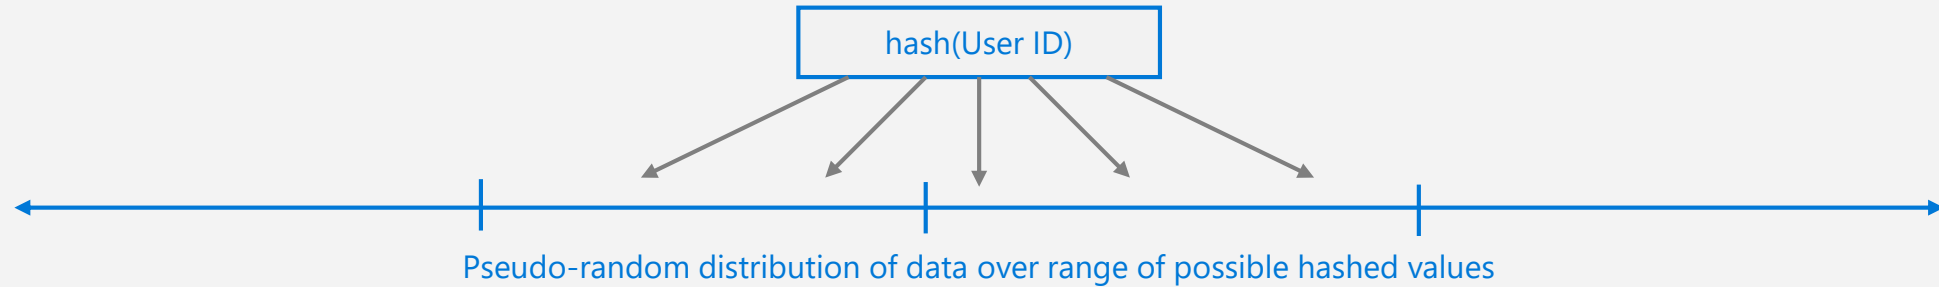
$$\begin{aligned} \text{[Total Monthly Cost]} &= \text{[Monthly Cost for Storage]} + \text{[Monthly Cost for Throughput]} \\ &= \$25 \quad \quad \quad + \$54 \\ &= \$79 \text{ per month} \end{aligned}$$

\* pricing may vary by region; for up-to-date pricing, see: <https://azure.microsoft.com/pricing/details/cosmos-db/>

# Partitions

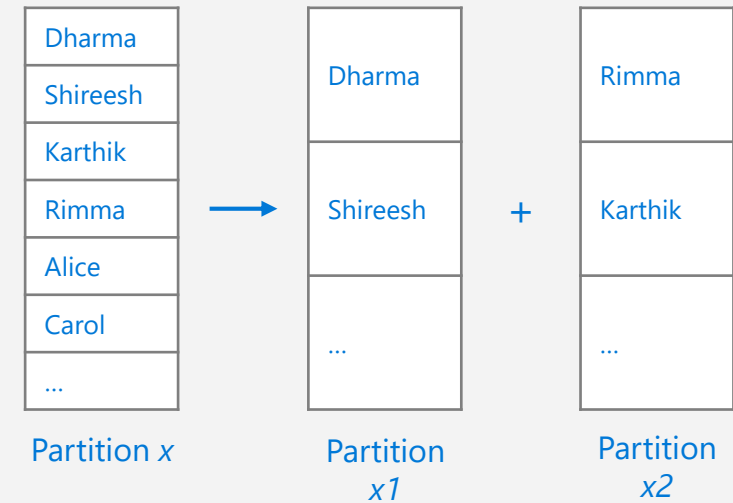


# Partitions



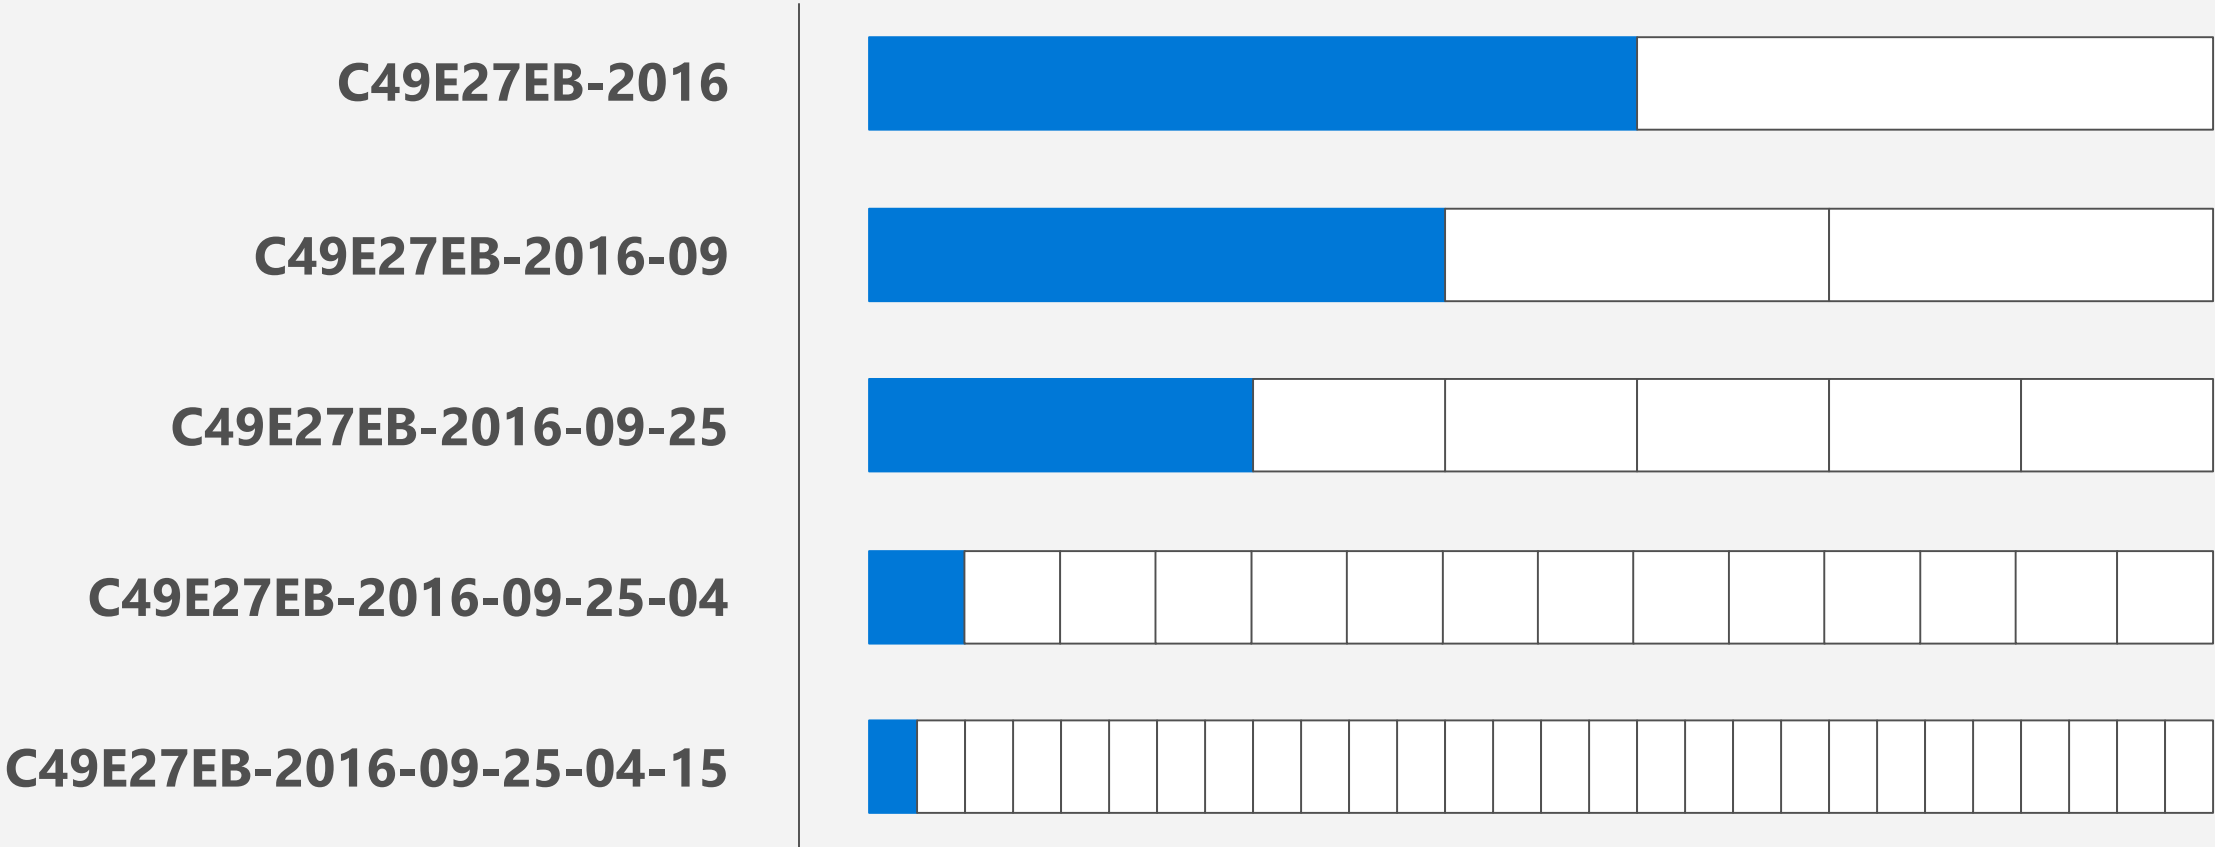
Partition Ranges can be dynamically sub-divided to seamlessly grow database as the application grows while simultaneously maintaining high availability.

**Partition management is fully managed** by Azure Cosmos DB, so you don't have to write code or manage your partitions.



# Partition Key Scenario

Interaction that occurred on: September 25, 2016 at 4:15 AM UTC



# Partitions

## Best Practices: Design Goals for Choosing a Good Partition Key

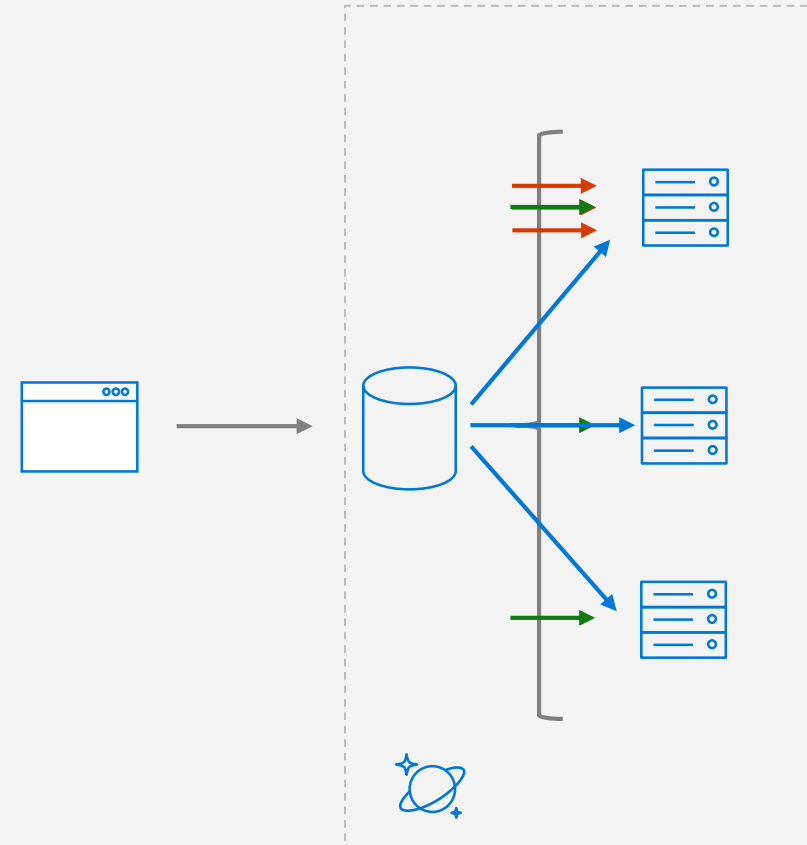
- Distribute the overall request + storage volume
  - Avoid "hot" partition keys
- Partition Key is scope for multi-record transactions and routing queries
  - Queries can be intelligently routed via partition key
  - Omitting partition key on query requires fan-out

## Steps for Success

- Ballpark scale needs (size/throughput)
- Understand the workload
- # of reads/sec vs writes per sec
  - Use pareto principal (80/20 rule) to help optimize bulk of workload
  - For reads – understand top 3-5 queries (look for common filters)
  - For writes – understand transactional needs

## General Tips

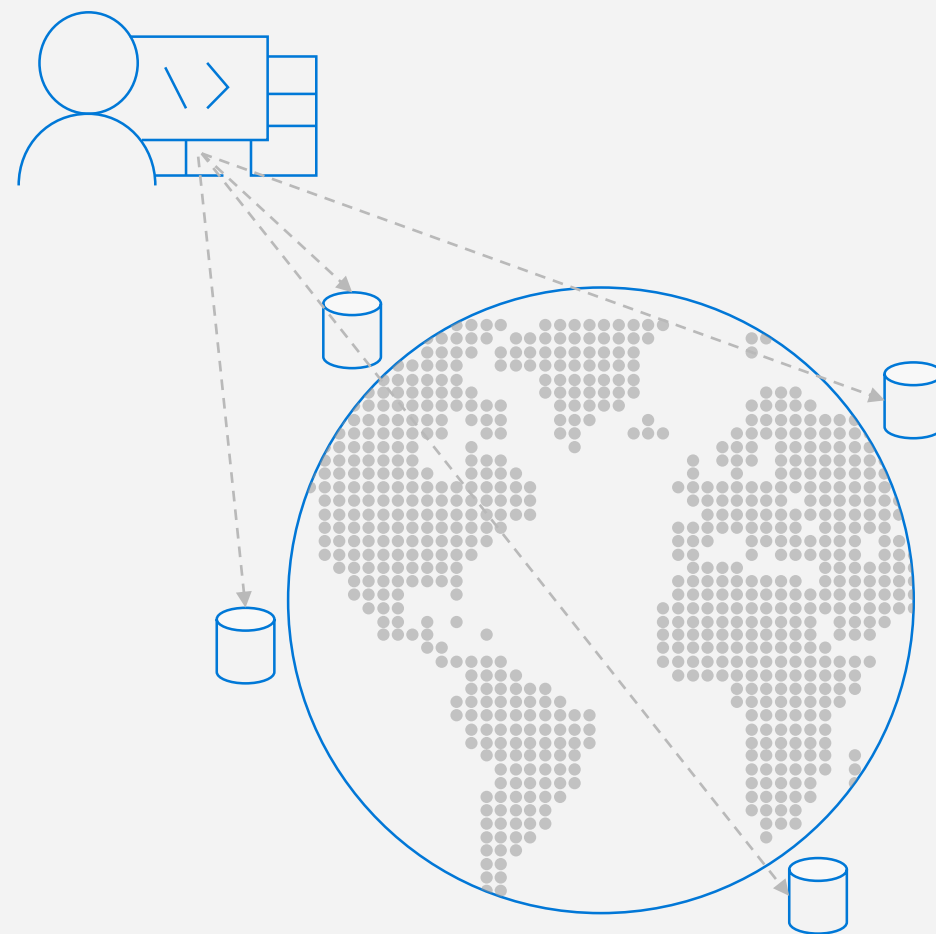
- Build a POC to strengthen your understanding of the workload and iterate (avoid analyses paralysis)
- Don't be afraid of having too many partition keys
  - Partitions keys are logical
  - More partition keys → more scalability



# SQL SYNTAX

Using the popular query language, SQL, to access semi-structured JSON data.

*This module will reference querying in the context of the SQL API for Azure Cosmos DB.*



# SQL QUERY SYNTAX - WHERE

## FILTERING

**WHERE** supports complex scalar expressions including arithmetic, comparison and logical operators

```
SELECT
    tickets.id,
    tickets.pricePaid
FROM tickets
WHERE
    tickets.pricePaid > 500.00 AND
    tickets.pricePaid <= 1000.00
```



# SQL QUERY SYNTAX - PROJECTION

## JSON PROJECTION

If your workloads require a specific JSON schema, Azure Cosmos DB supports JSON projection within its queries

```
SELECT {  
  "id": tickets.id,  
  "flightNumber": tickets.assignedFlight.flightNumber,  
  "purchase": {  
    "cost": tickets.pricePaid  
  },  
  "stops": [  
    tickets.assignedFlight.origin,  
    tickets.assignedFlight.destination  
  ]  
} AS ticket  
FROM tickets
```



```
[  
  {  
    "ticket": {  
      "id": "6ebe1165836a",  
      "purchase": {  
        "cost": 575.5  
      },  
      "stops": [  
        "SEA",  
        "JFK"  
      ]  
    }  
  }  
]
```

# INTRA-DOCUMENT JOIN

JOIN allows us to merge embedded documents or arrays across multiple documents and returned a flattened result set:

```
SELECT
    tickets.assignedFlight.number,
    tickets.seat,
    requests
FROM
    tickets
JOIN
    requests IN tickets.requests
```



```
[
  {
    "number": "F125", "seat": "12A",
    "requests": "kosher_meal"
  },
  {
    "number": "F125", "seat": "12A",
    "requests": "aisle_seat"
  },
  {
    "number": "F752", "seat": "14C",
    "requests": "early_boarding"
  },
  {
    "number": "F752", "seat": "14C",
    "requests": "window_seat"
  }
]
```

# Paginated Query Results

## Pagination with hasNext() in Java:

```
Iterator<Document> documents = client.queryDocuments(  
    collectionLink,  
    queryString,  
    options  
) .getQueryIterator();  
  
while(documents.hasNext()) {  
    Document current = documents.next();  
}
```



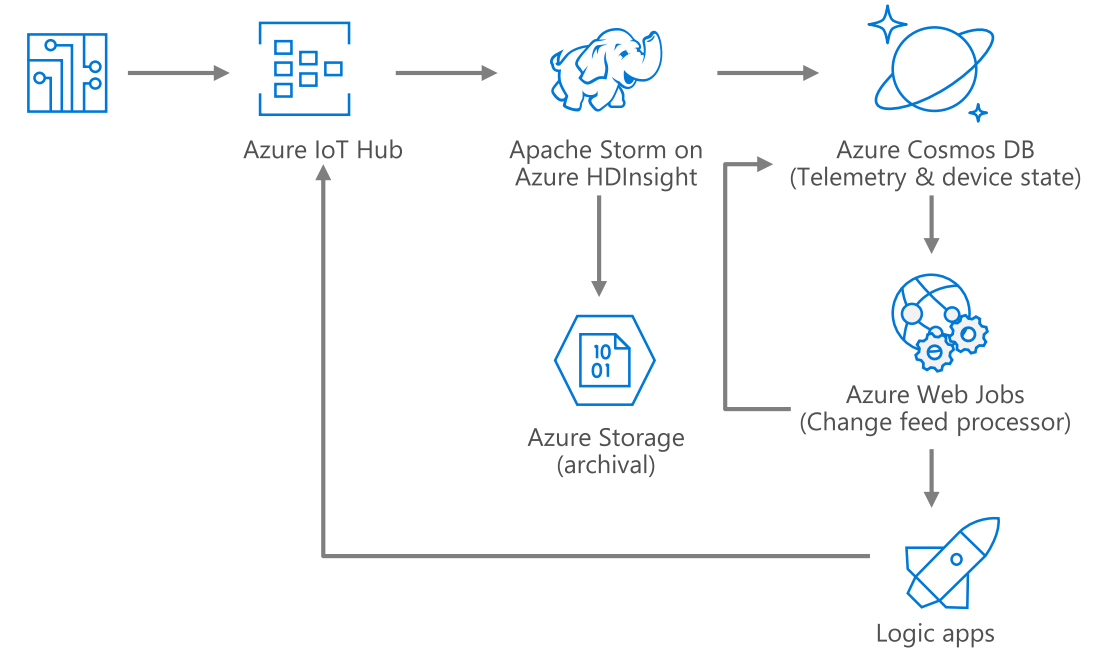
# Massive Scale Telemetry Stores for IOT

Diverse and unpredictable IoT sensor workloads  
require a responsive data platform

Seamless handling of any data output or volume

Data made available immediately, and indexed  
automatically

High writes per second, with stable ingestion and query  
performance

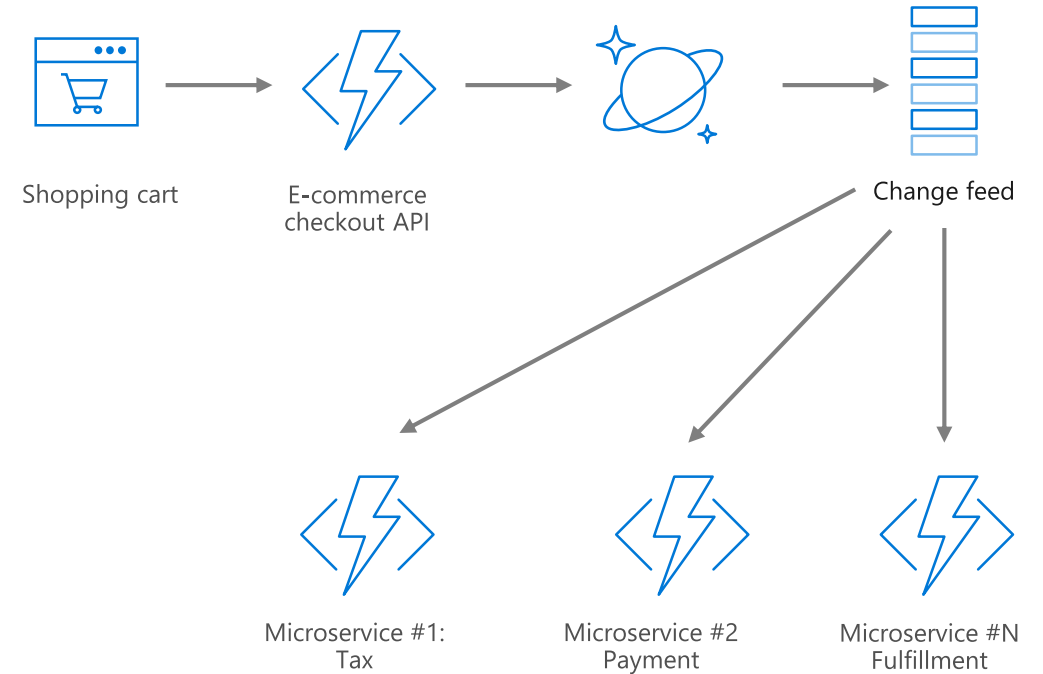


**Honeywell**

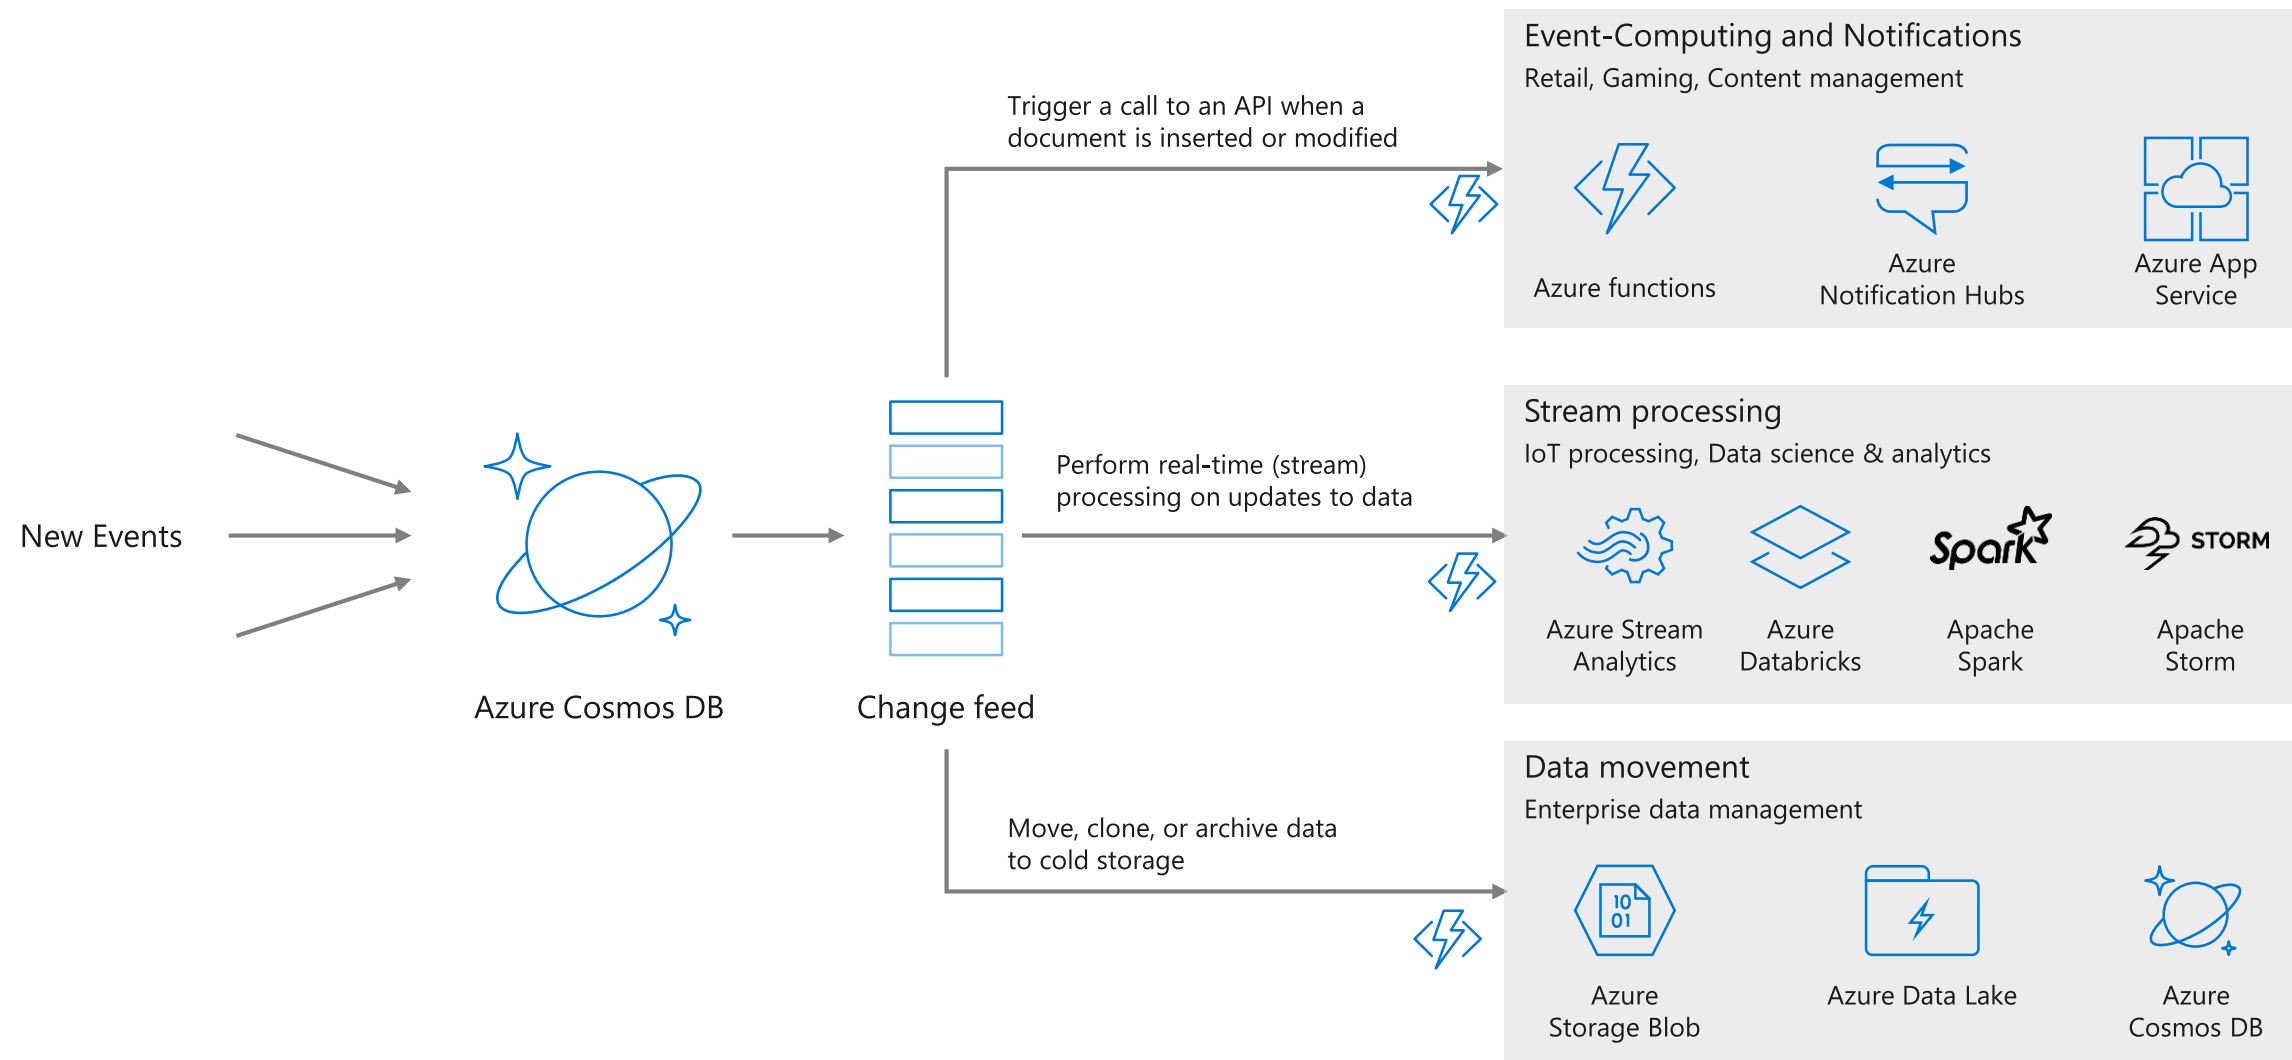


# Event Sourcing & Microservice architecture

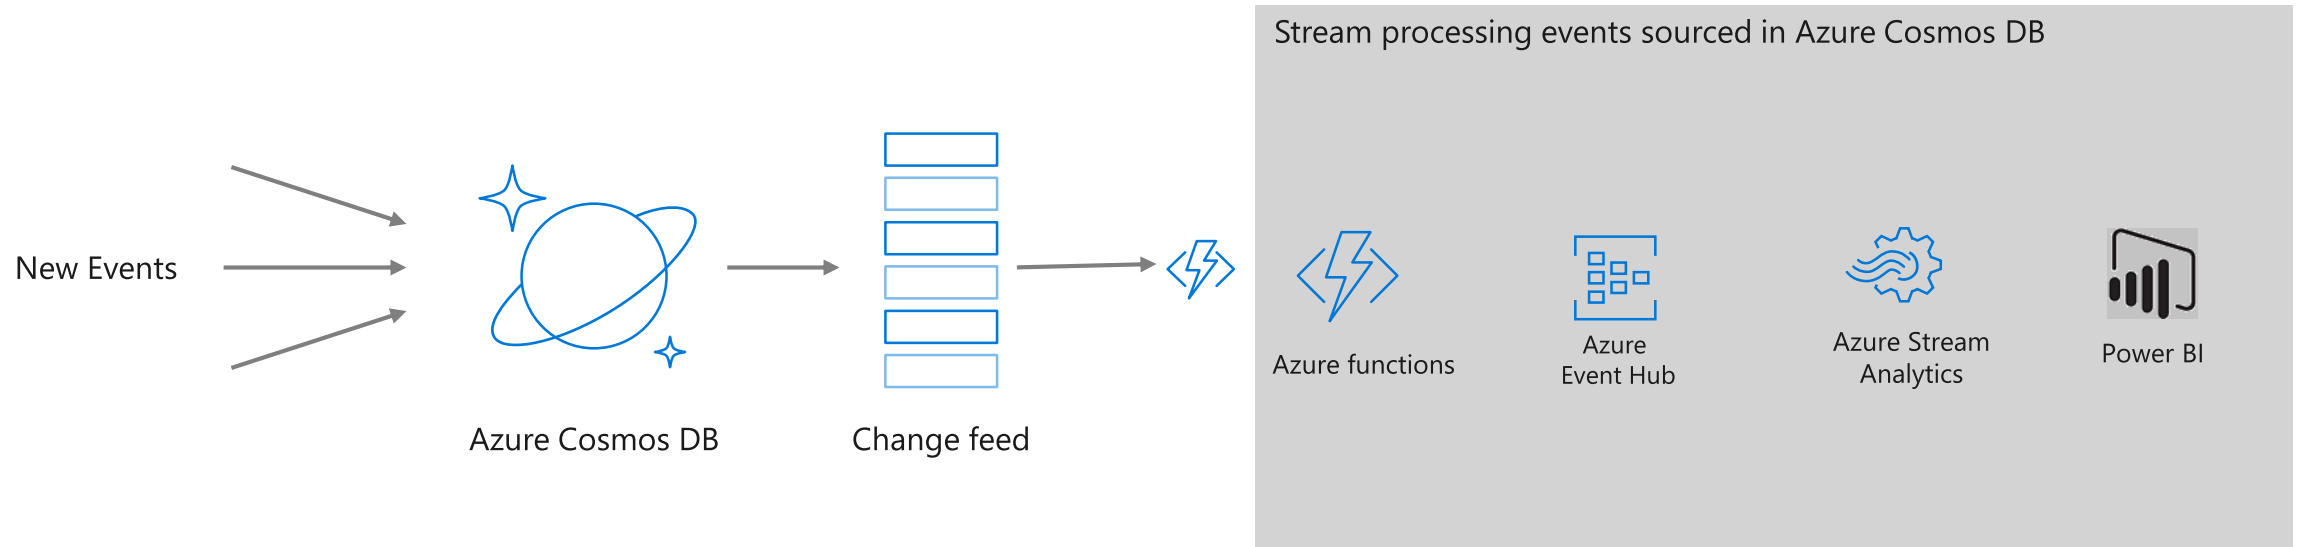
- IoT, gaming, retail and operational logging applications need to track and respond to tremendous amount of data being ingested, modified or removed from a globally-scaled database.
- Using Cosmos DB Change Feed:
- Listen for any changes in a container with Azure Functions
- Process ordered list of changed documents
- Trigger downstream microservices, with Cosmos as "source of truth" for all events in system



# Change Feed Scenarios



# The Architecture for Today's Experience



# Ideal for E-commerce, Gaming, and IOT

Maintain service quality during high-traffic periods requiring massive scale and performance.

Instant, elastic scaling handles traffic bursts

Uninterrupted global user experience

Low-latency data access (<10ms at P99) to process large and changing user bases

High availability (99.999%) across multiple data centers

