Microsoft Fabric

Data analytics for the era of AI

Andrea Benedetti

Sr Cloud Architect, Microsoft

/in/abenedetti

@anBenedetti

https://github.com/anbened

General Availability of Microsoft Fabric

Copilot in Microsoft Fabric is rolling out in public preview

Public Preview → over 25,000 orgs

"Analytics for the era of AI"
becoming GA

"Fabric has been our biggest data launch perhaps since SQL server"
Satya Nadella

Timeline

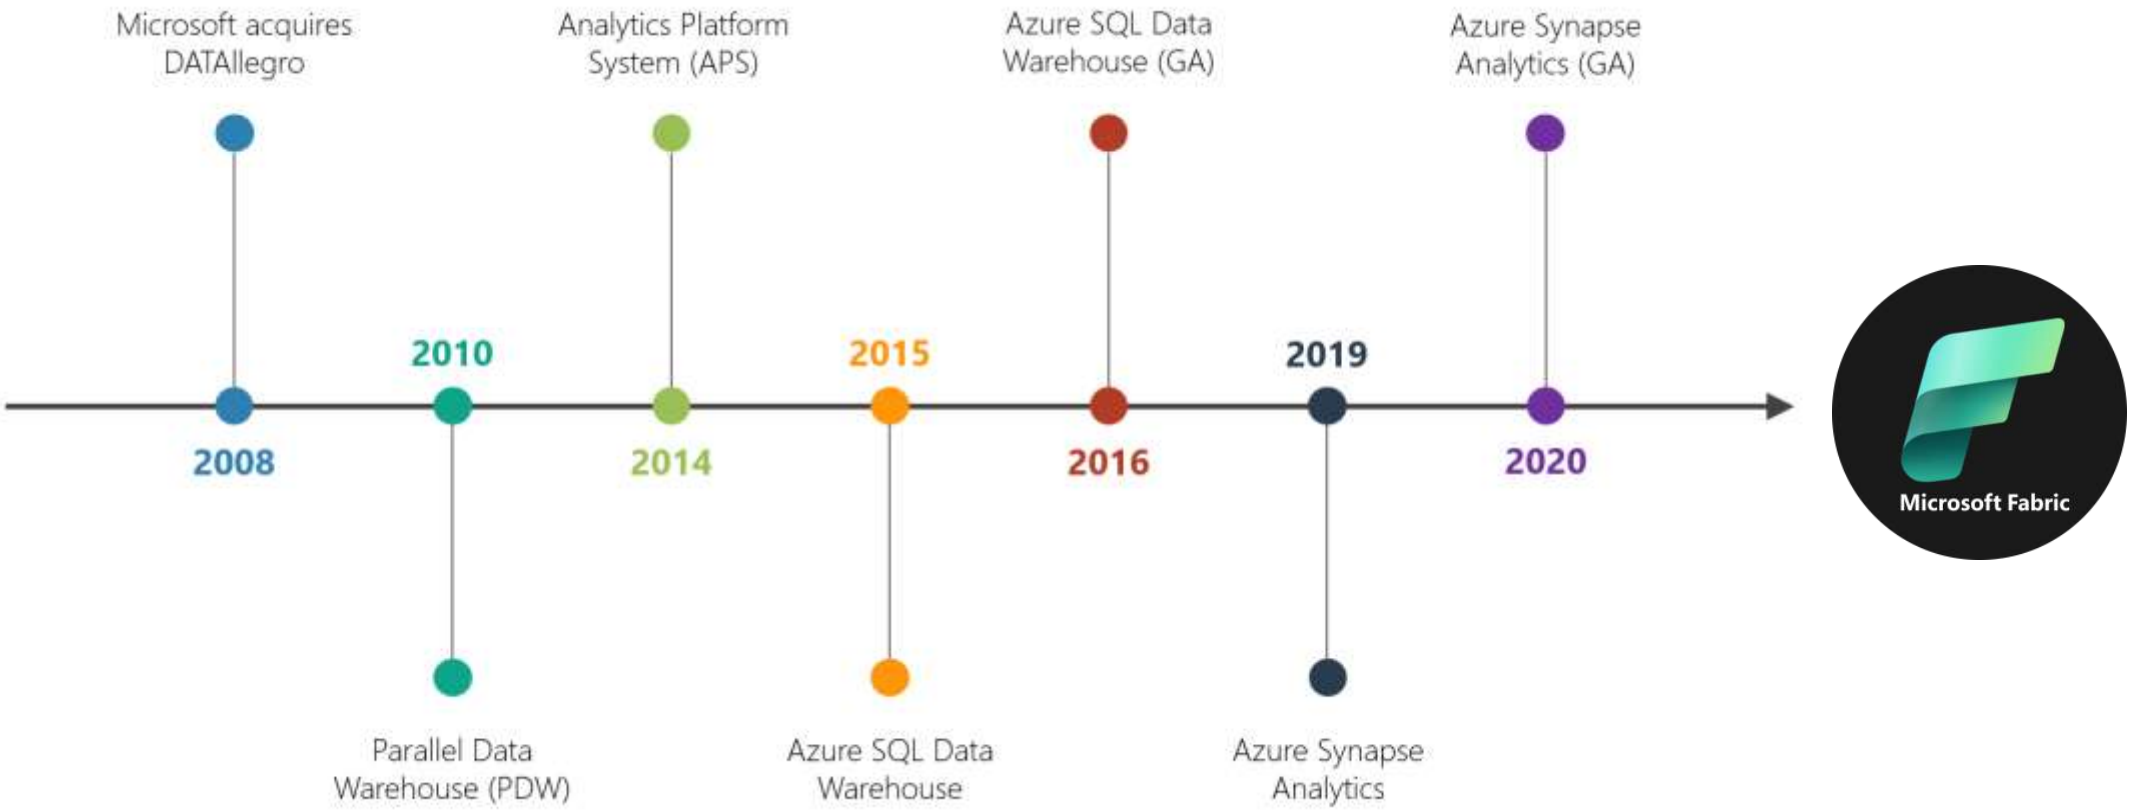

Ref: https://www.taygan.co/blog/2022/01/04/azure-synapse-analytics

Scalable analytics are complex and fragmented

Every analytics project has many subsystems

Every subsystem need a different class of product

Products often come from multiple vendors

Integration at scale across products is complex, fragile, and expensive

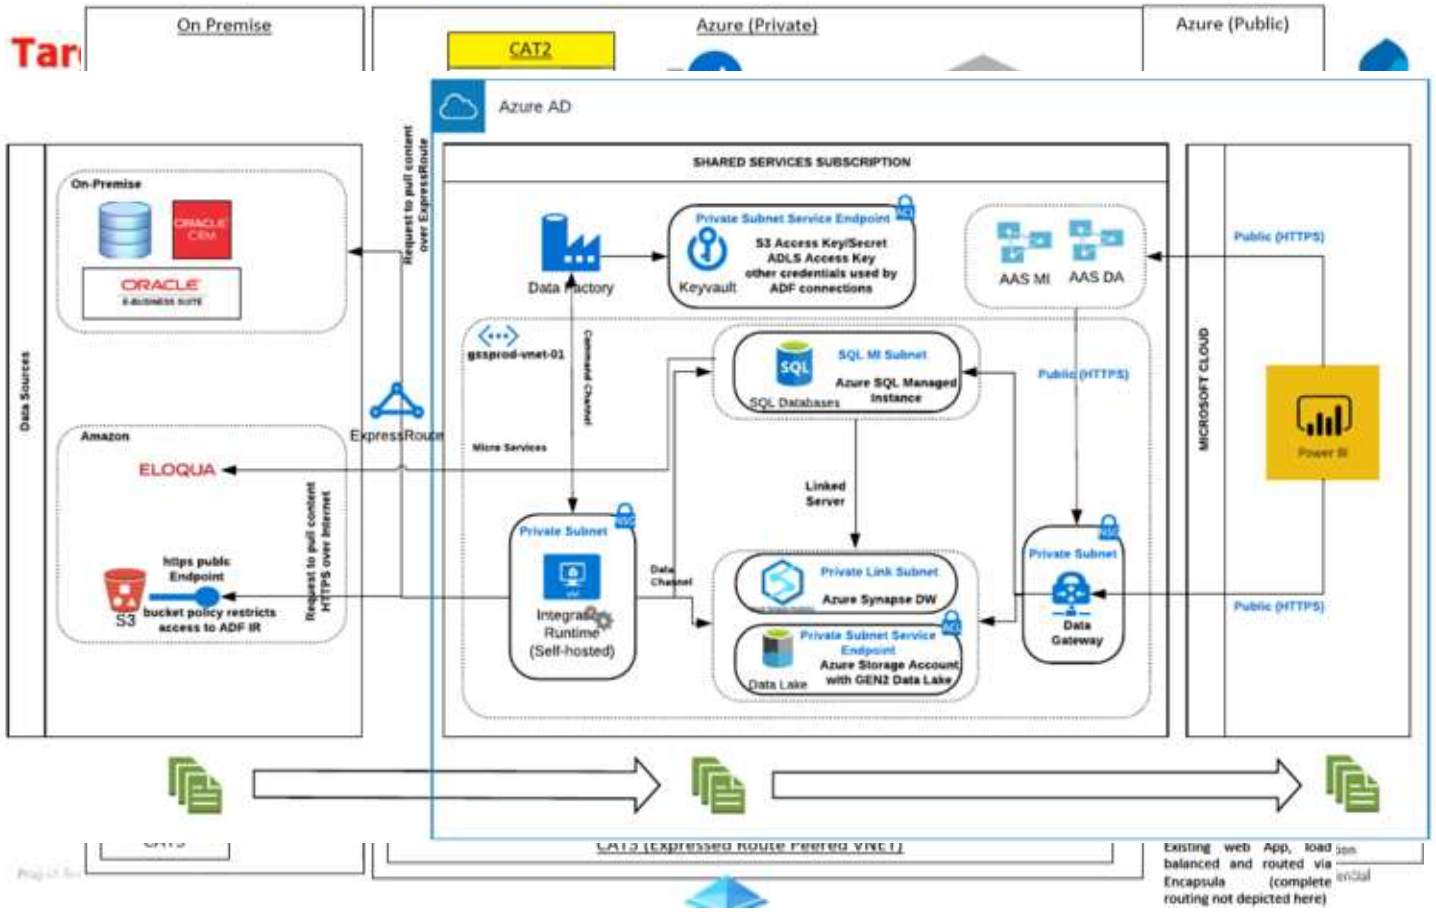

Scalable analytics are complex and fragmented

Every analytics project has many subsystems

Every subsystem need a different class of product

Products often come from multiple vendors

Integration at scale across products is complex, fragile and expensive

> **Simplify,**
> I am the Chief Data Officer and don't want to be the Chief Integration Officer."
>
> Every CDO, Every Enterprise

A silver lining?

Analytics systems have very predictable patterns

Microsoft has all the products with the right scale needed to build a complete analytics system

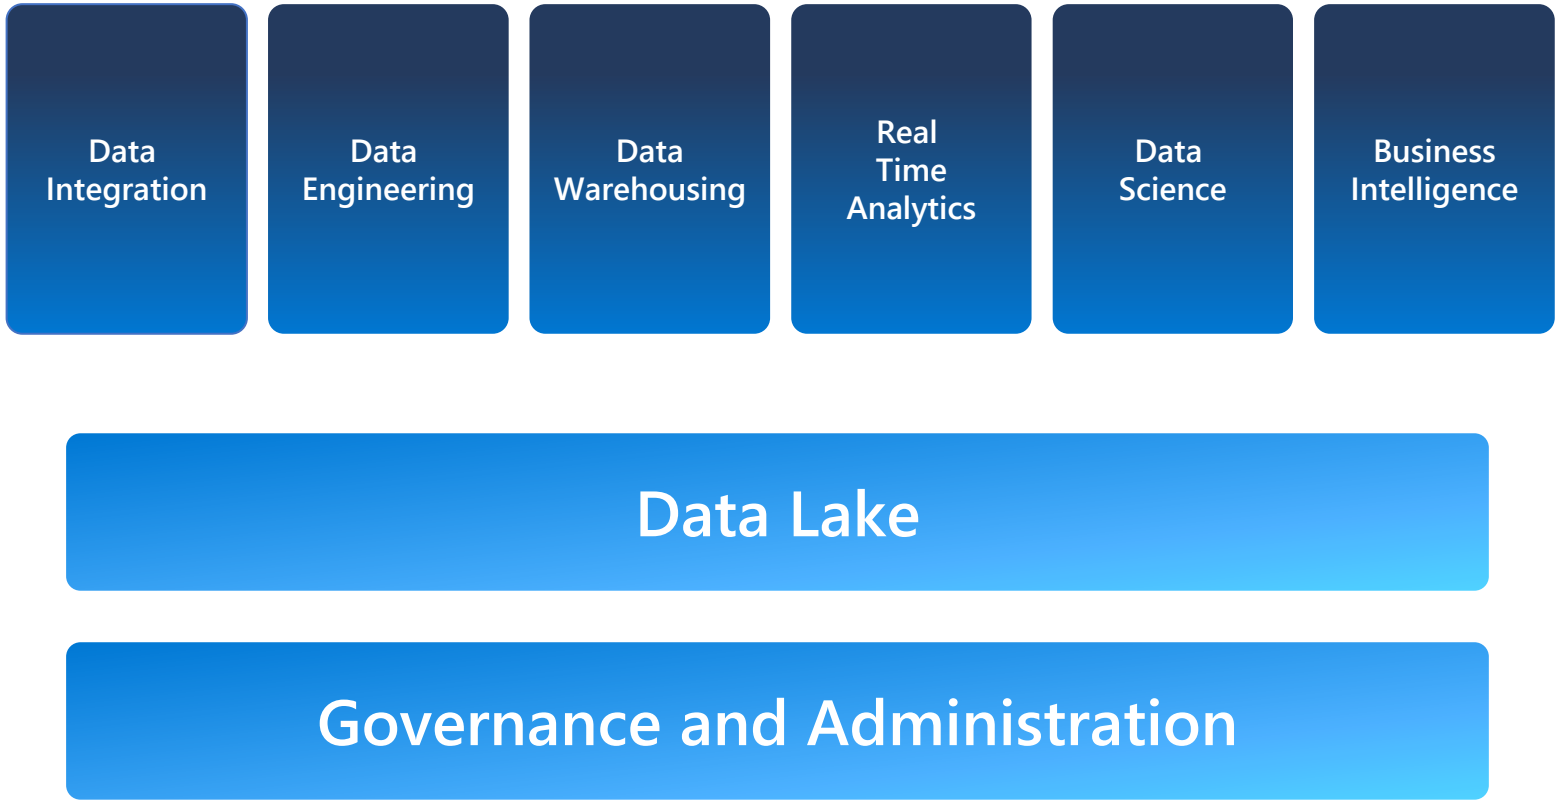

A silver lining?

Analytics systems have very predictable patterns

Microsoft has all the products with the right scale needed to build a complete analytics system

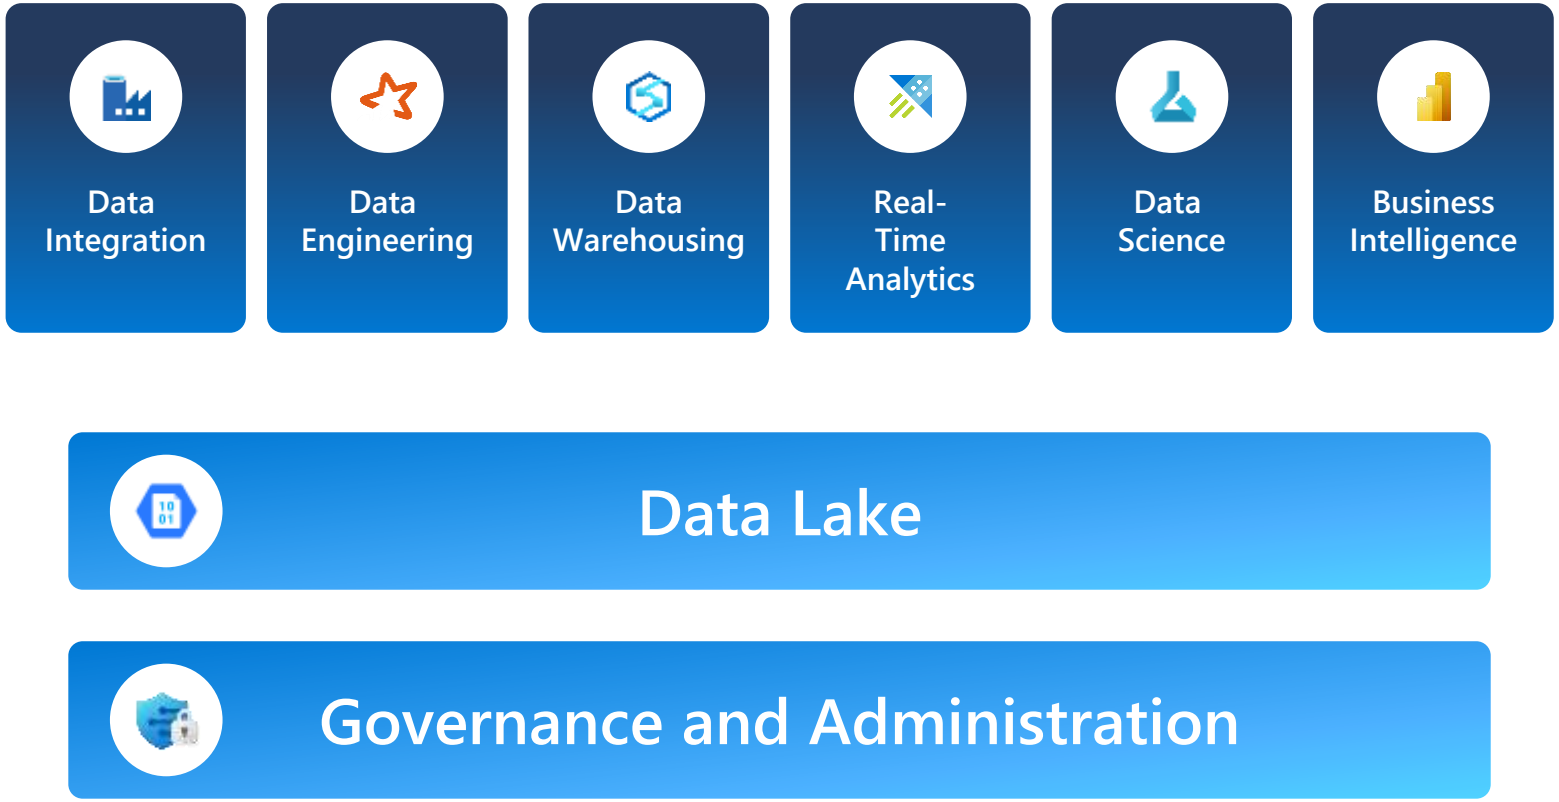

Still far too complex

- Many Products
- Different Experiences
- Proprietary and Open
- Dedicated and Serverless
- PaaS and SaaS
- Different Business Models
- Steep Learning Curves
- Deep Expertise Needed
- High Integration Effort


Power BI


Synapse


Kusto


Azure AI


Data Factory


Spark

What we’re hearing from our customers

How do I unify all our disparate sources of data into one single source of truth?

How do I improve democratization of data and insights with easy-to-use analytics tools?

How do I enable faster decision-making with scale, while optimizing governance and cost?

What is Fabric?

End-to-end analytics data fabric
From the data lake to the business user

Complete Analytics Platform

- Best of Breed
- Unified SaaS Solution
- Low Code Plus Pro Dev

Lake-centric and Open

- OneLake
- One Copy
- Always Synced

Empower Every Office User

- Familiar and Intuitive
- Built Into Office
- Insight to Action

Persistent Security and Governance

- End-to-End Visibility
- Always Governed
- Secure by Default

Microsoft Fabric

Data Integration	Data Engineering	Data Warehouse	Data Science	Real Time Analytics	Business Intelligence	Observability
Data Factory	Synapse	Synapse	Synapse	Synapse	Power BI	Data Activator

Unified data foundation
OneLake

UNIFIED

SaaS product experience | Security and governance | Compute and storage | Business model

End-to-end analytics data fabric
From the data lake to the business user

Microsoft Fabric

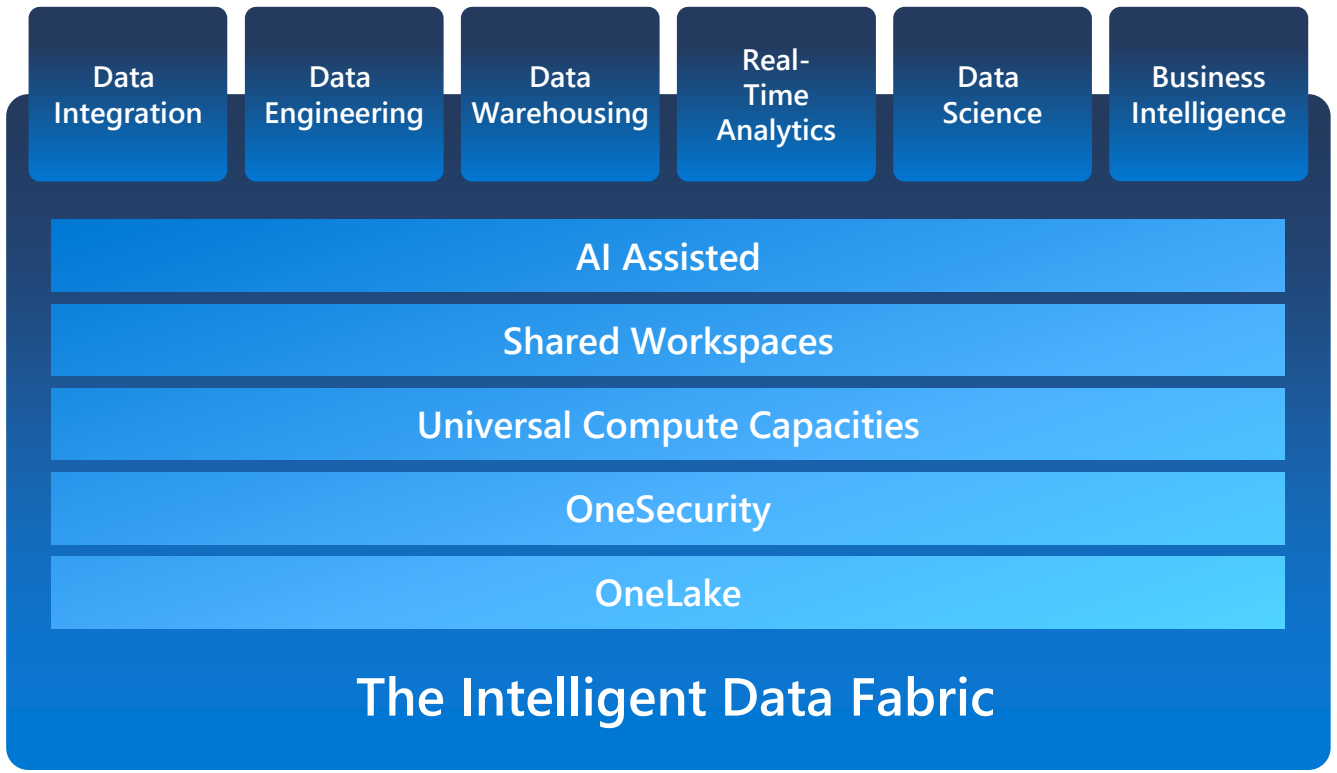

Single...

Onboarding and trials
Sign-on
Navigation model
UX model
Workspace organization
Collaboration experience
Data Lake
Storage format
Data copy for all engines
Security model
CI/CD
Monitoring hub
Data Hub
Governance & compliance

What is Fabric? Recap

- Adds significant functionality to a Power BI tenant
- Compute in the form of Capacity Units (CU) → not a «server»
  - Power BI Premium P1 = F64
- Compute control abstracted away
- Storage = open-souce Delta Lake format

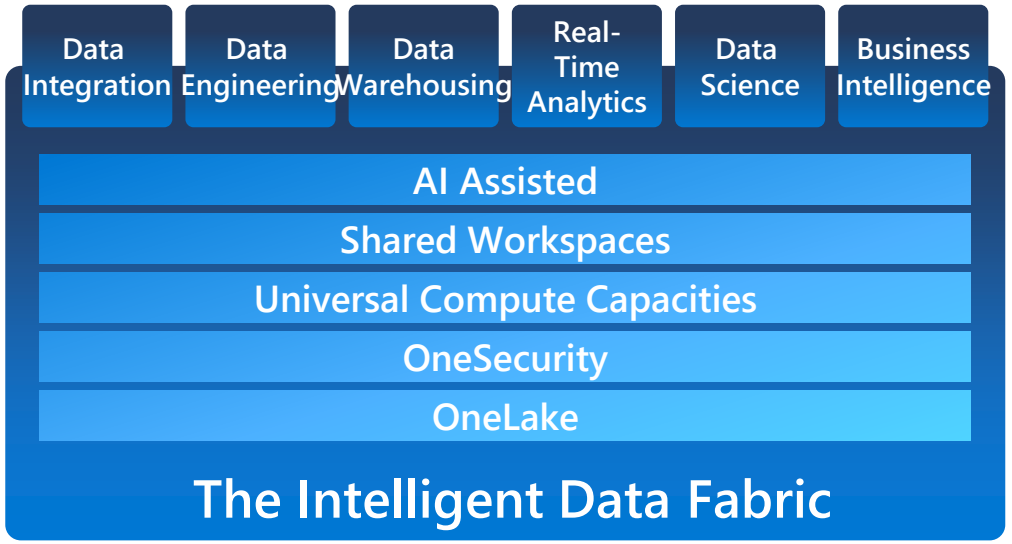

Persona Centric Experience

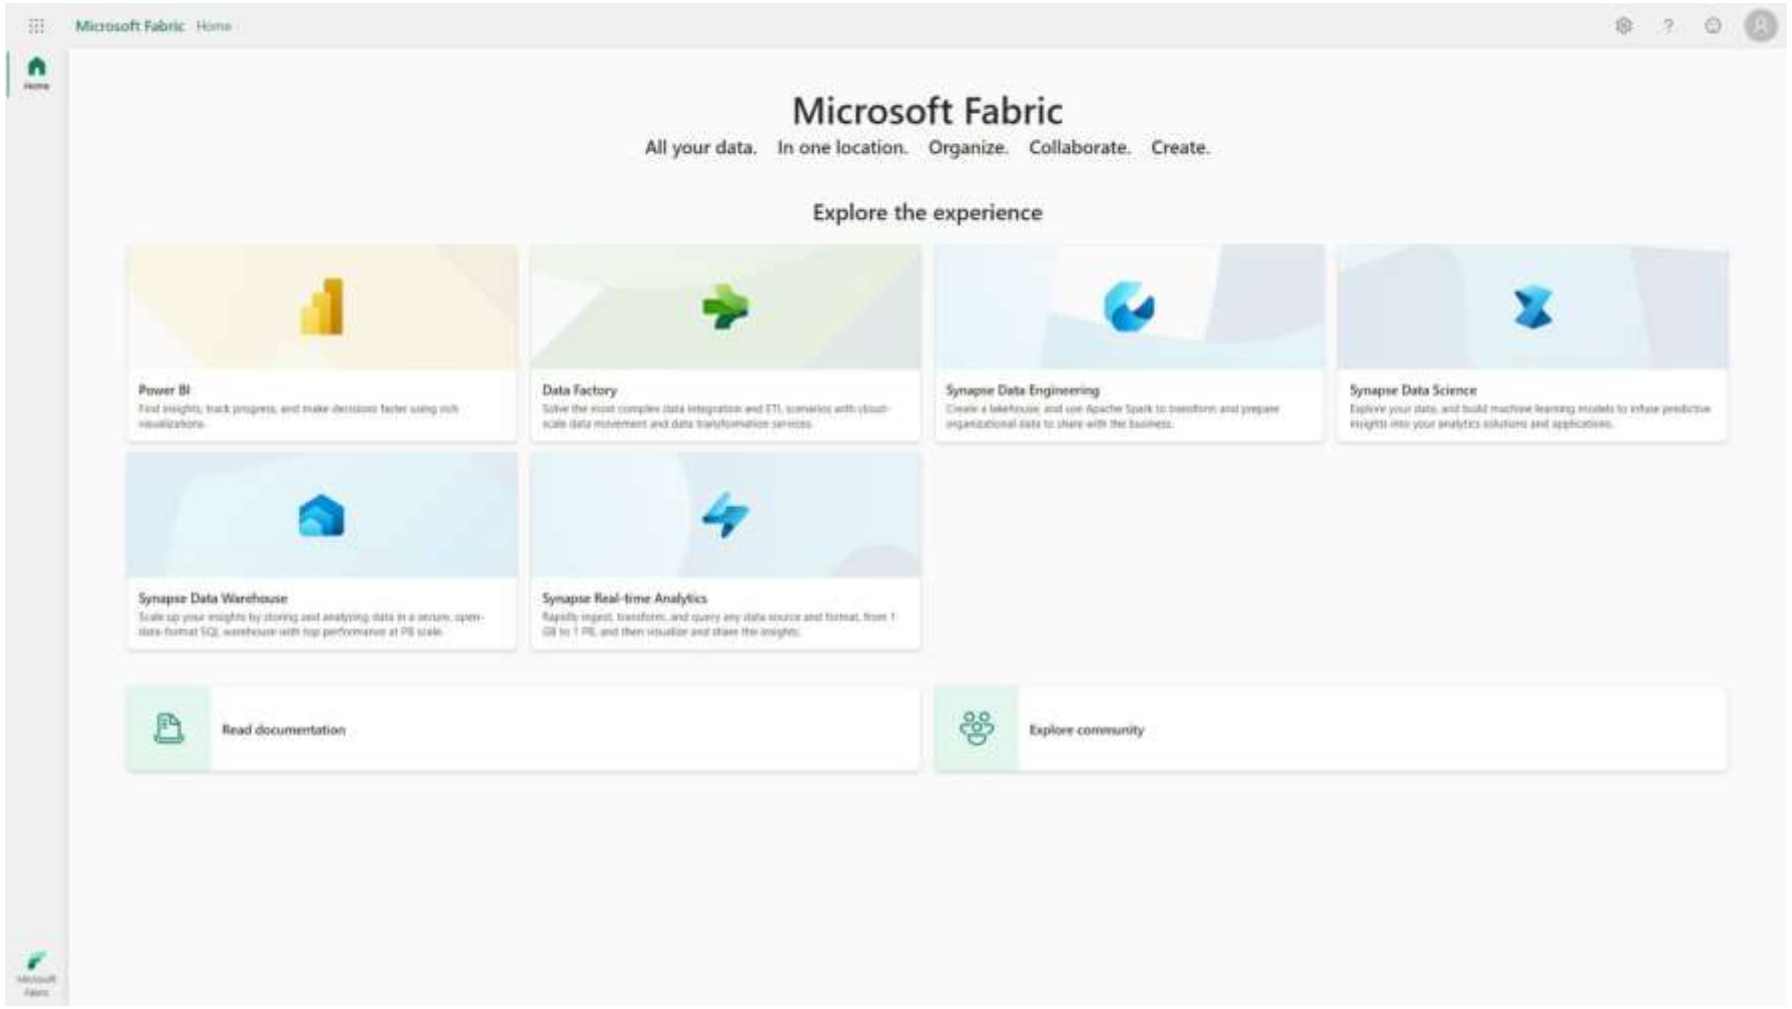

Data Integration

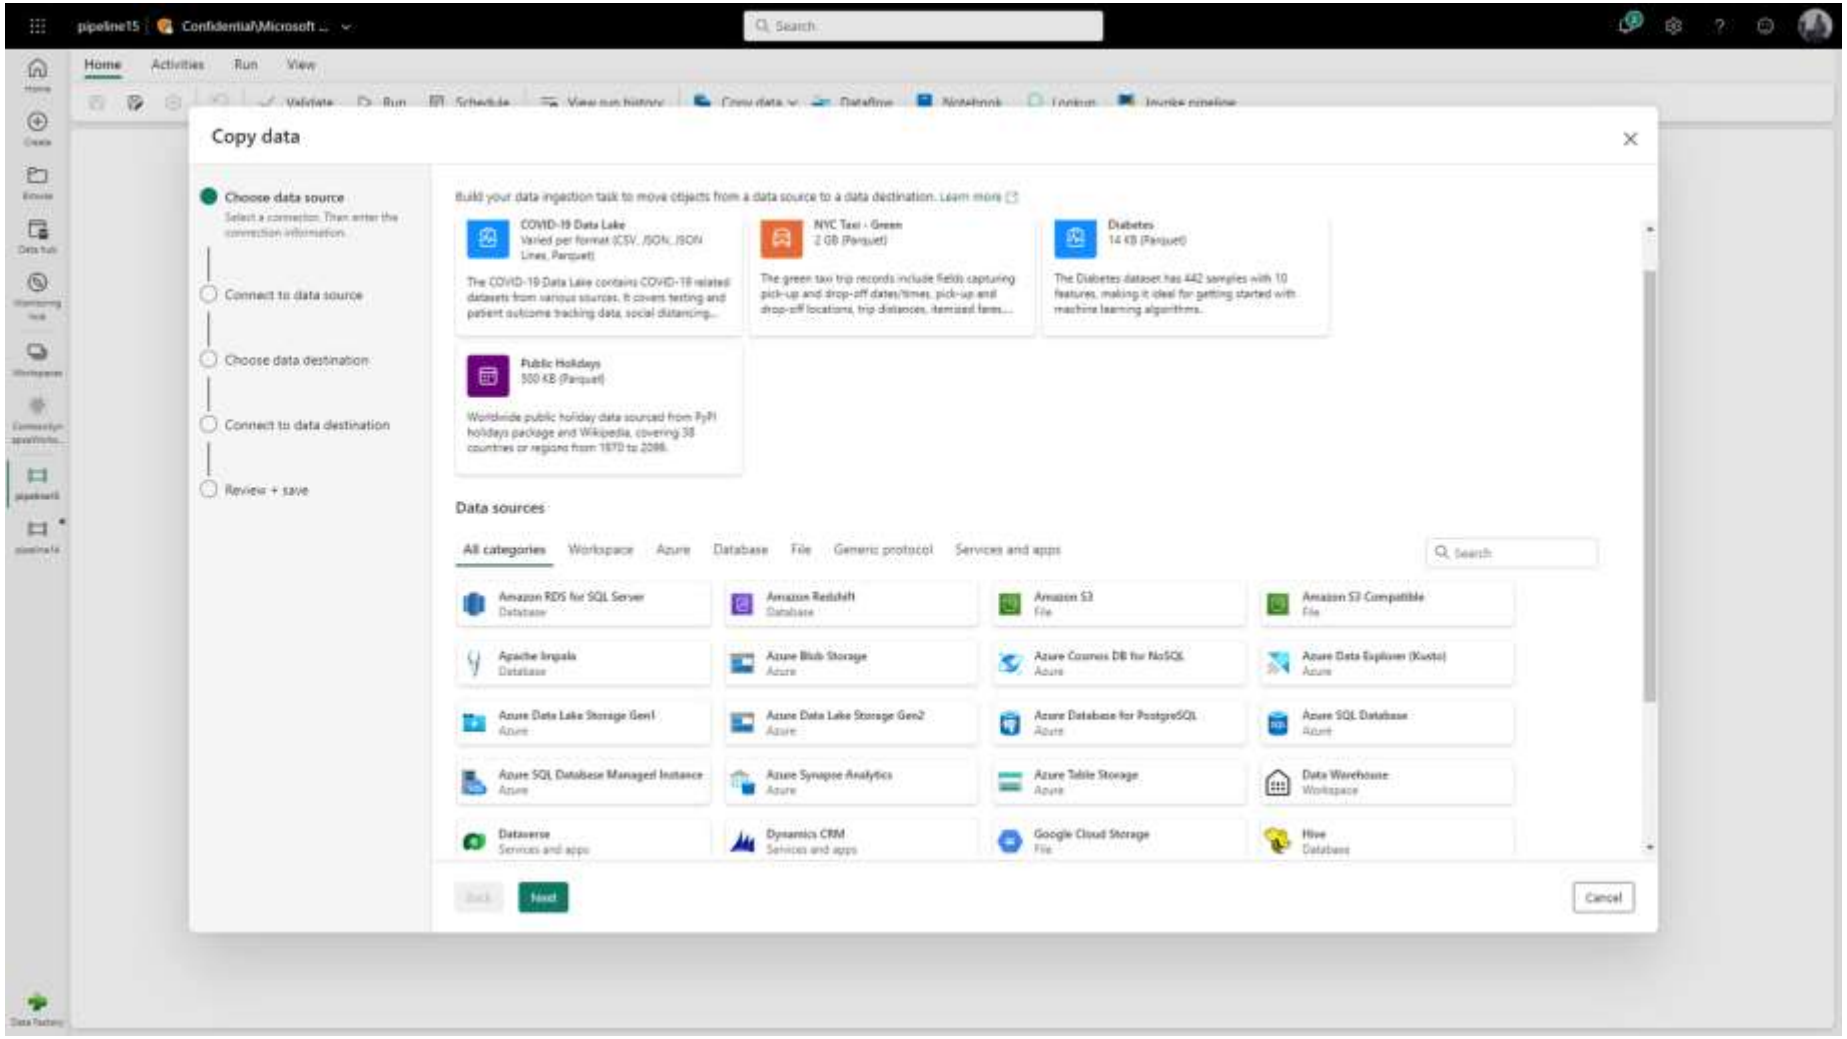

Data Engineering

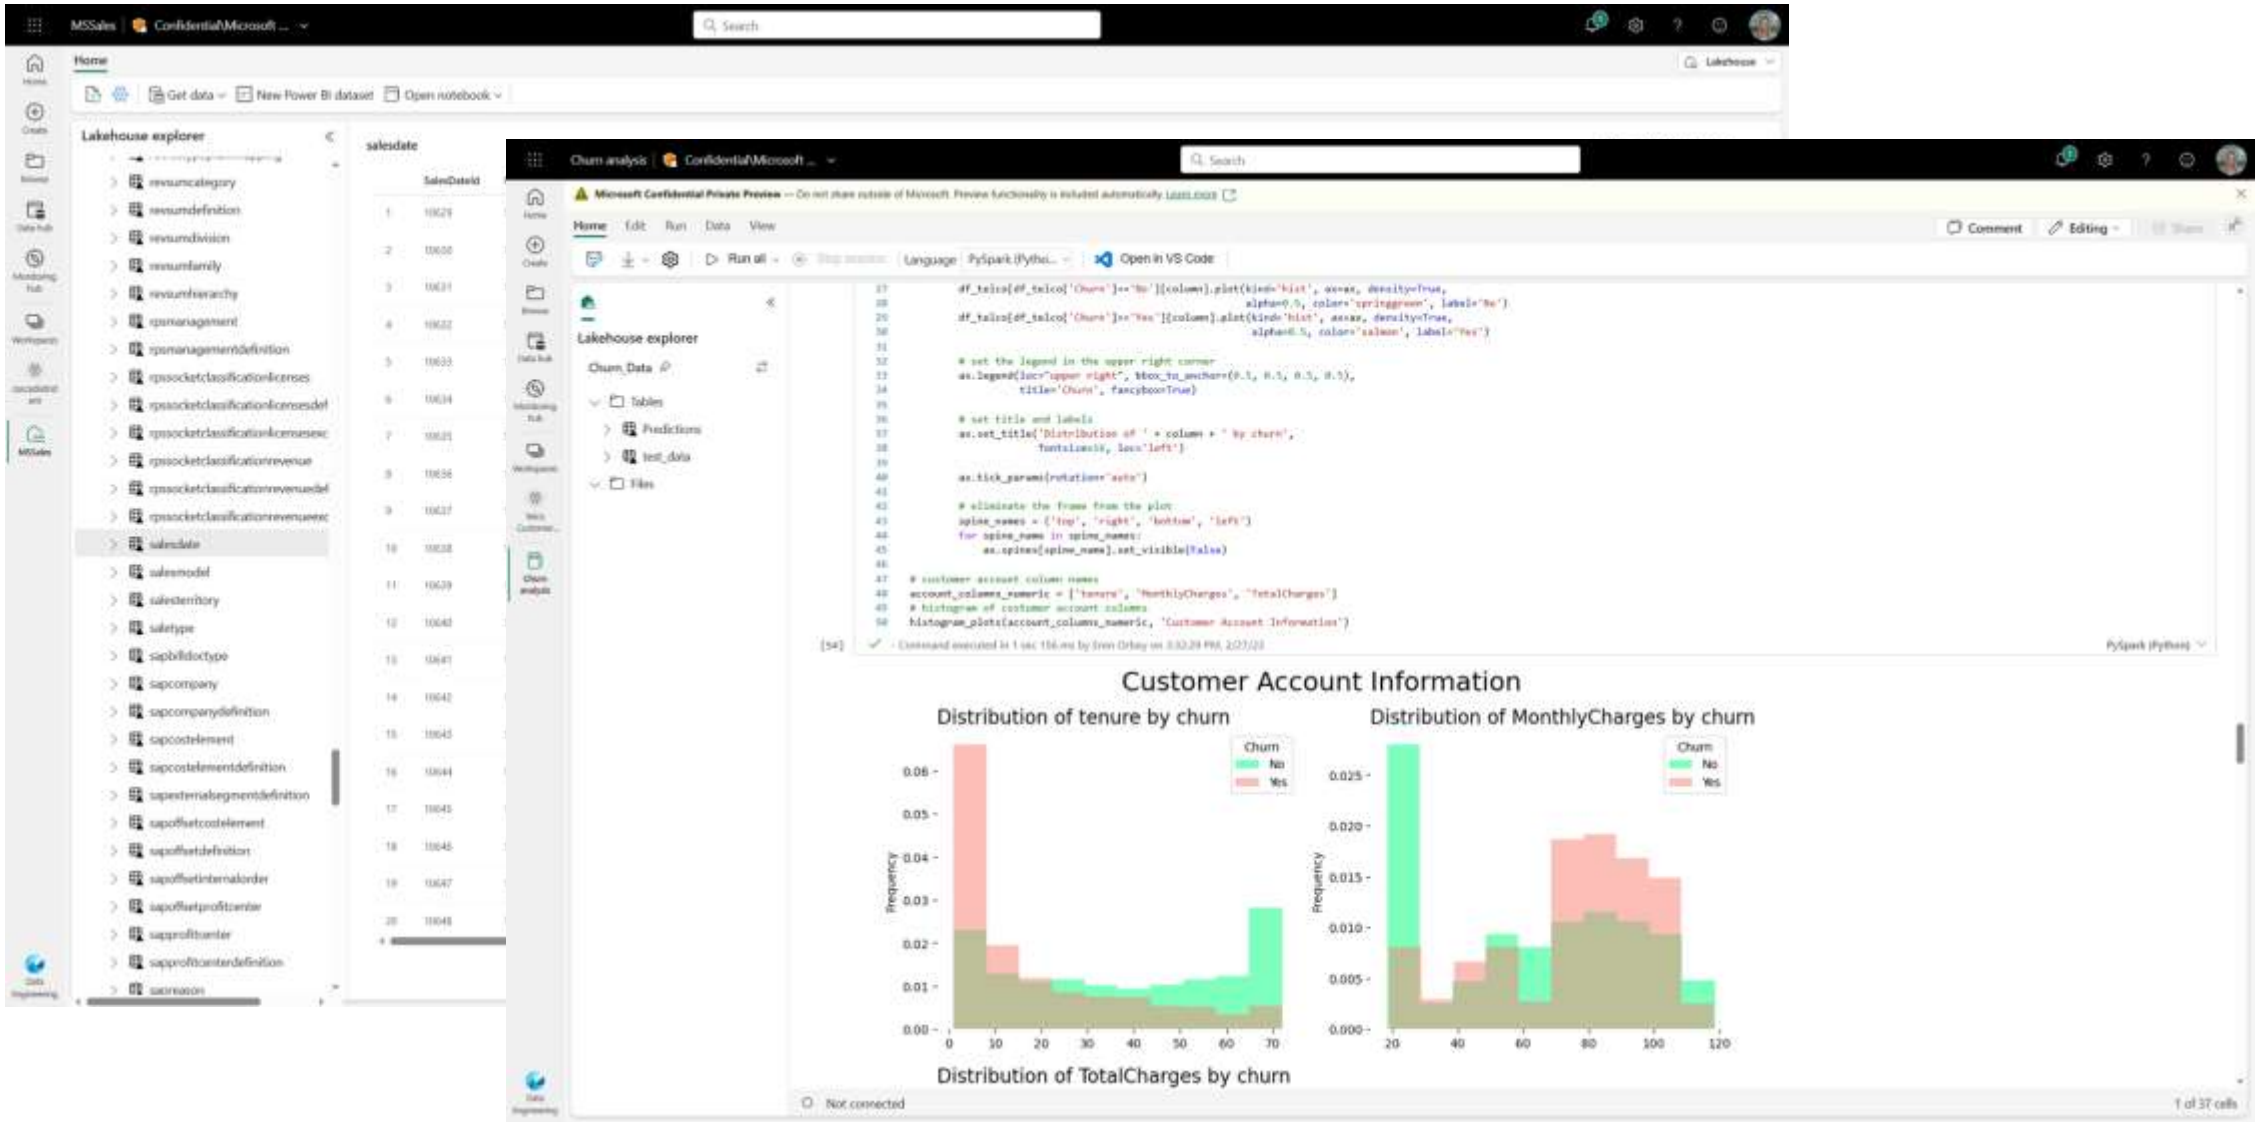

Data Warehouse

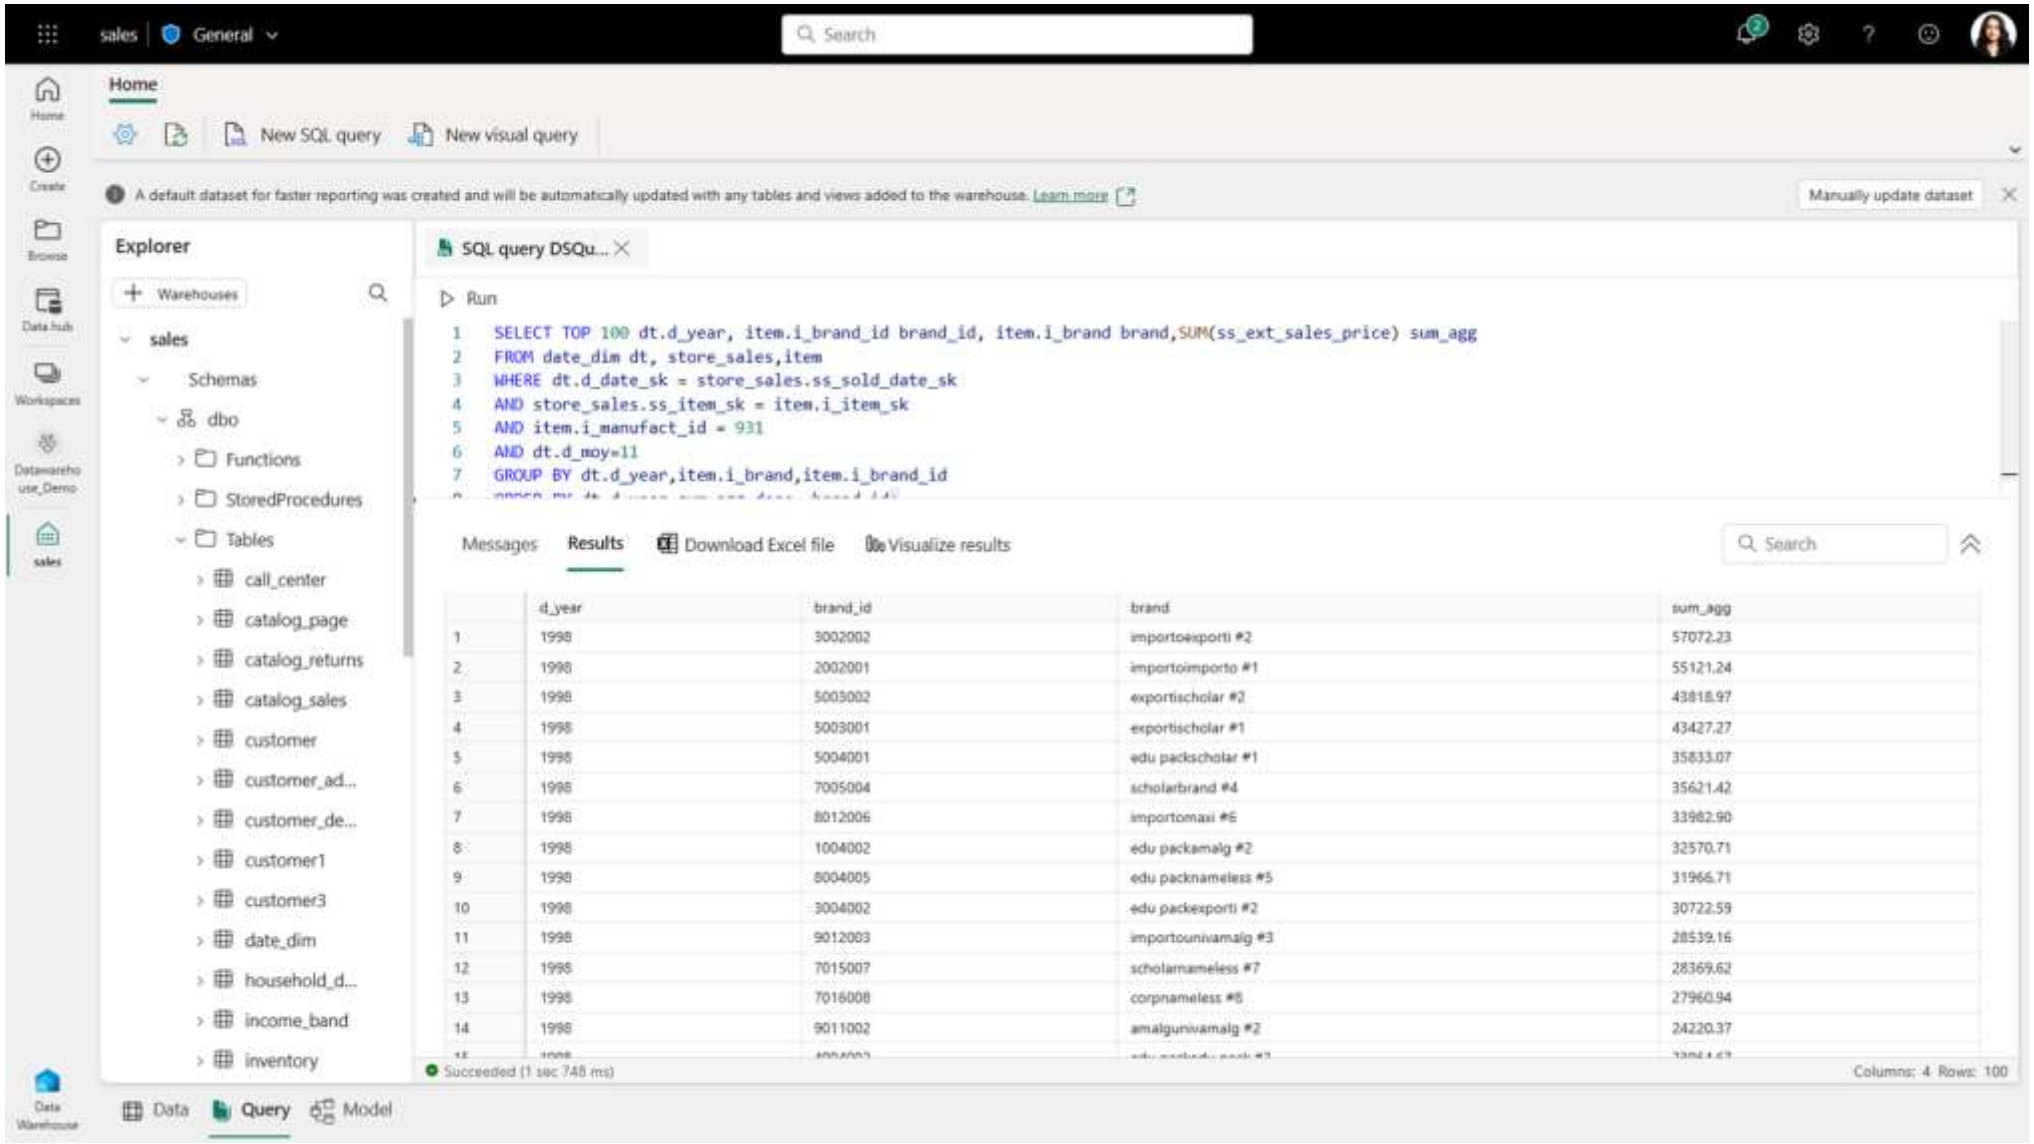

Real-time Analytics

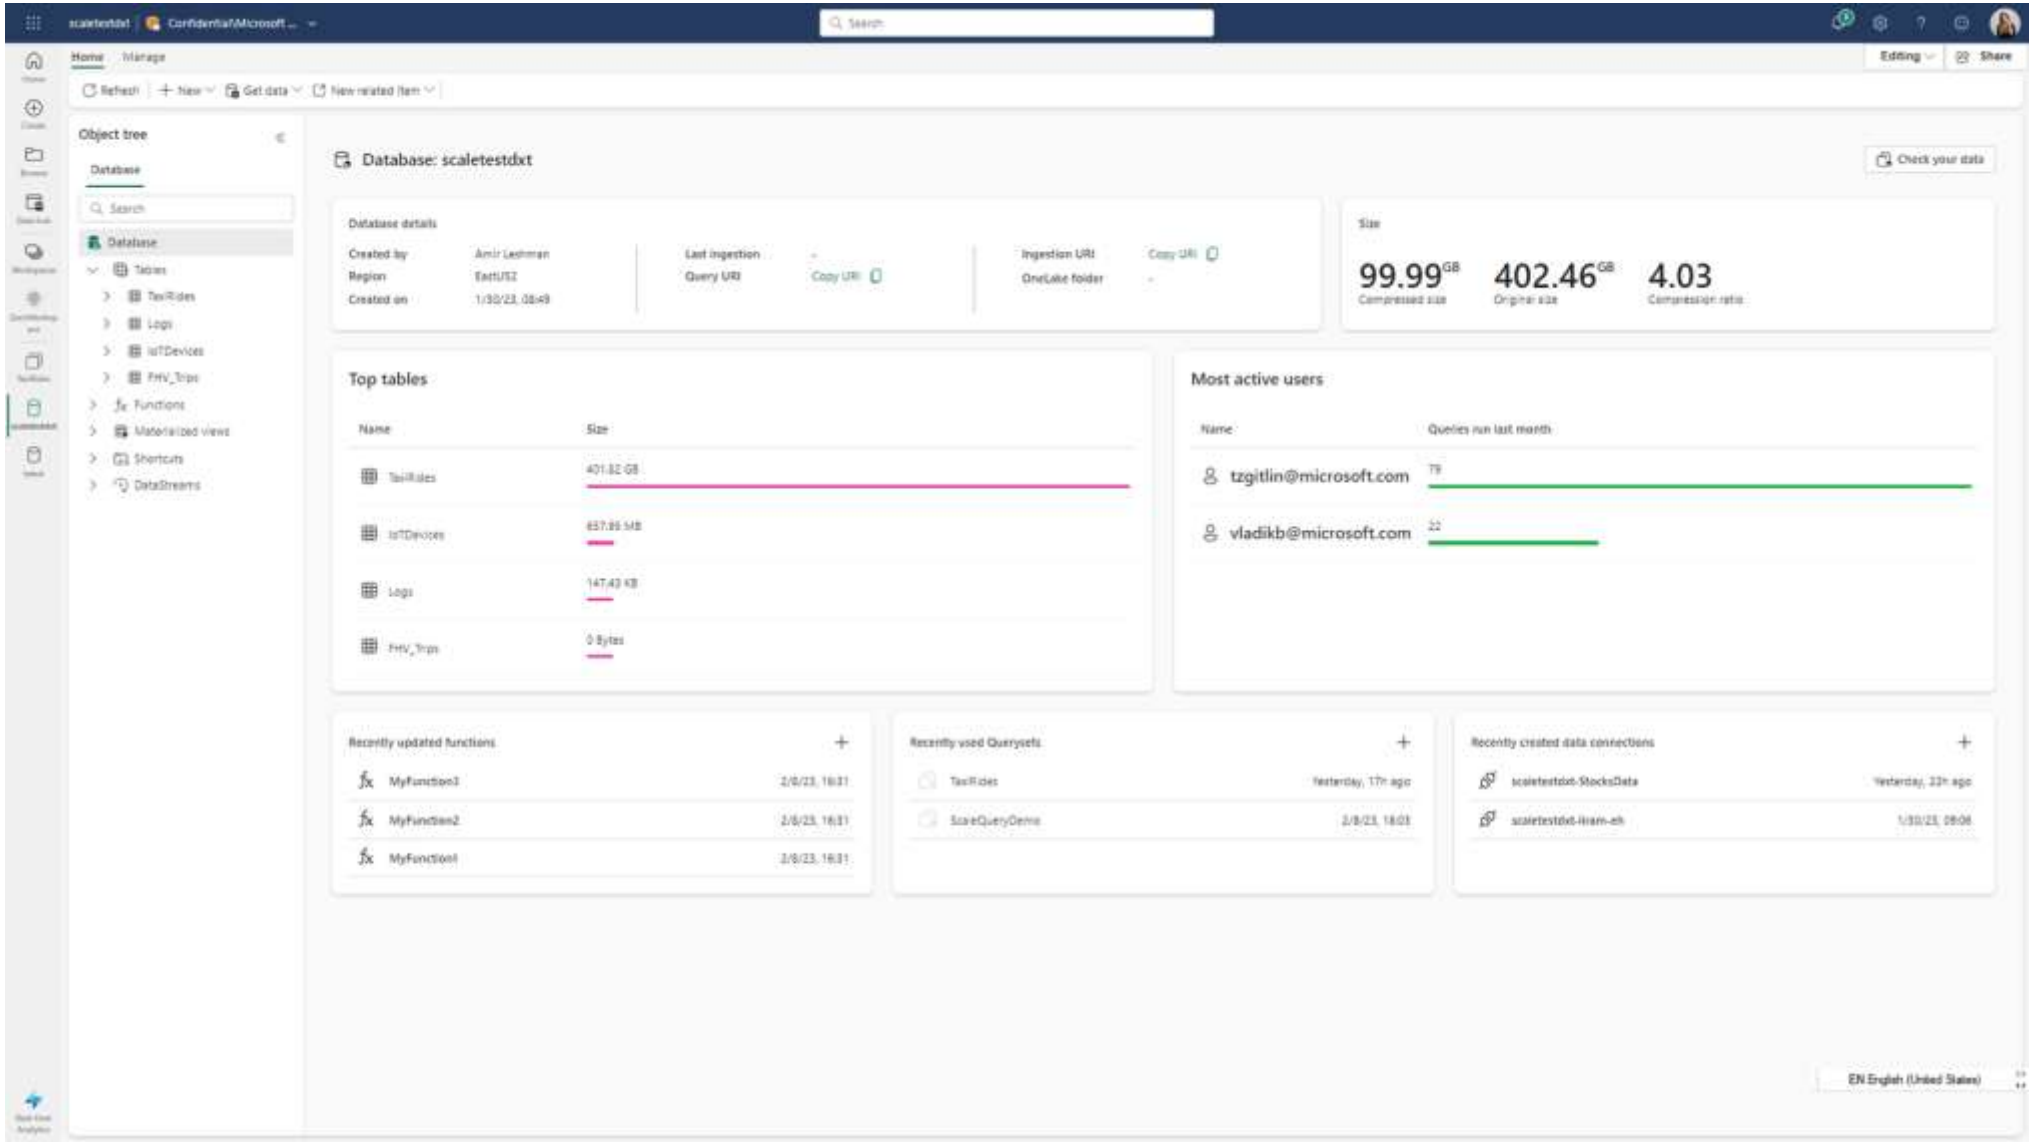

Data Science

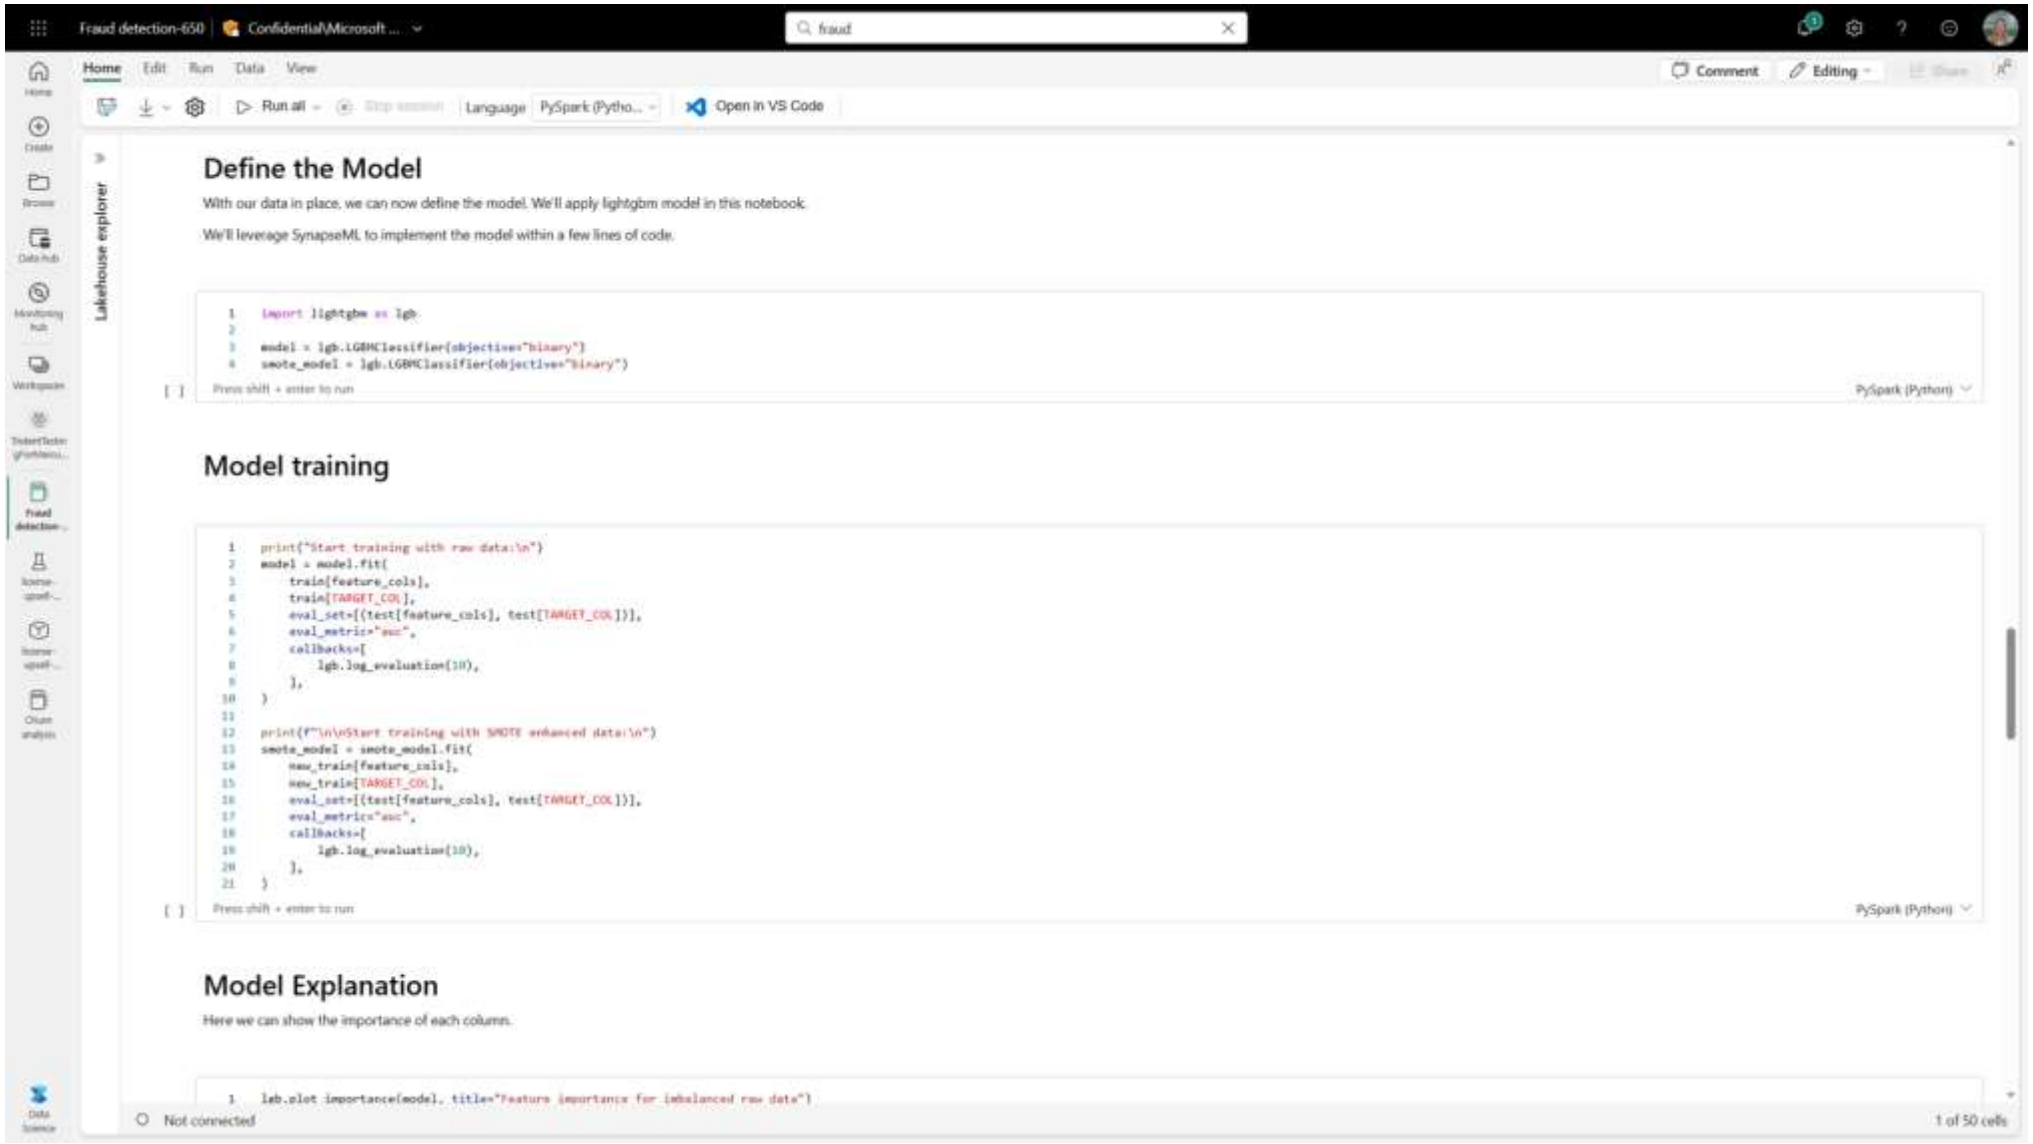

Power BI

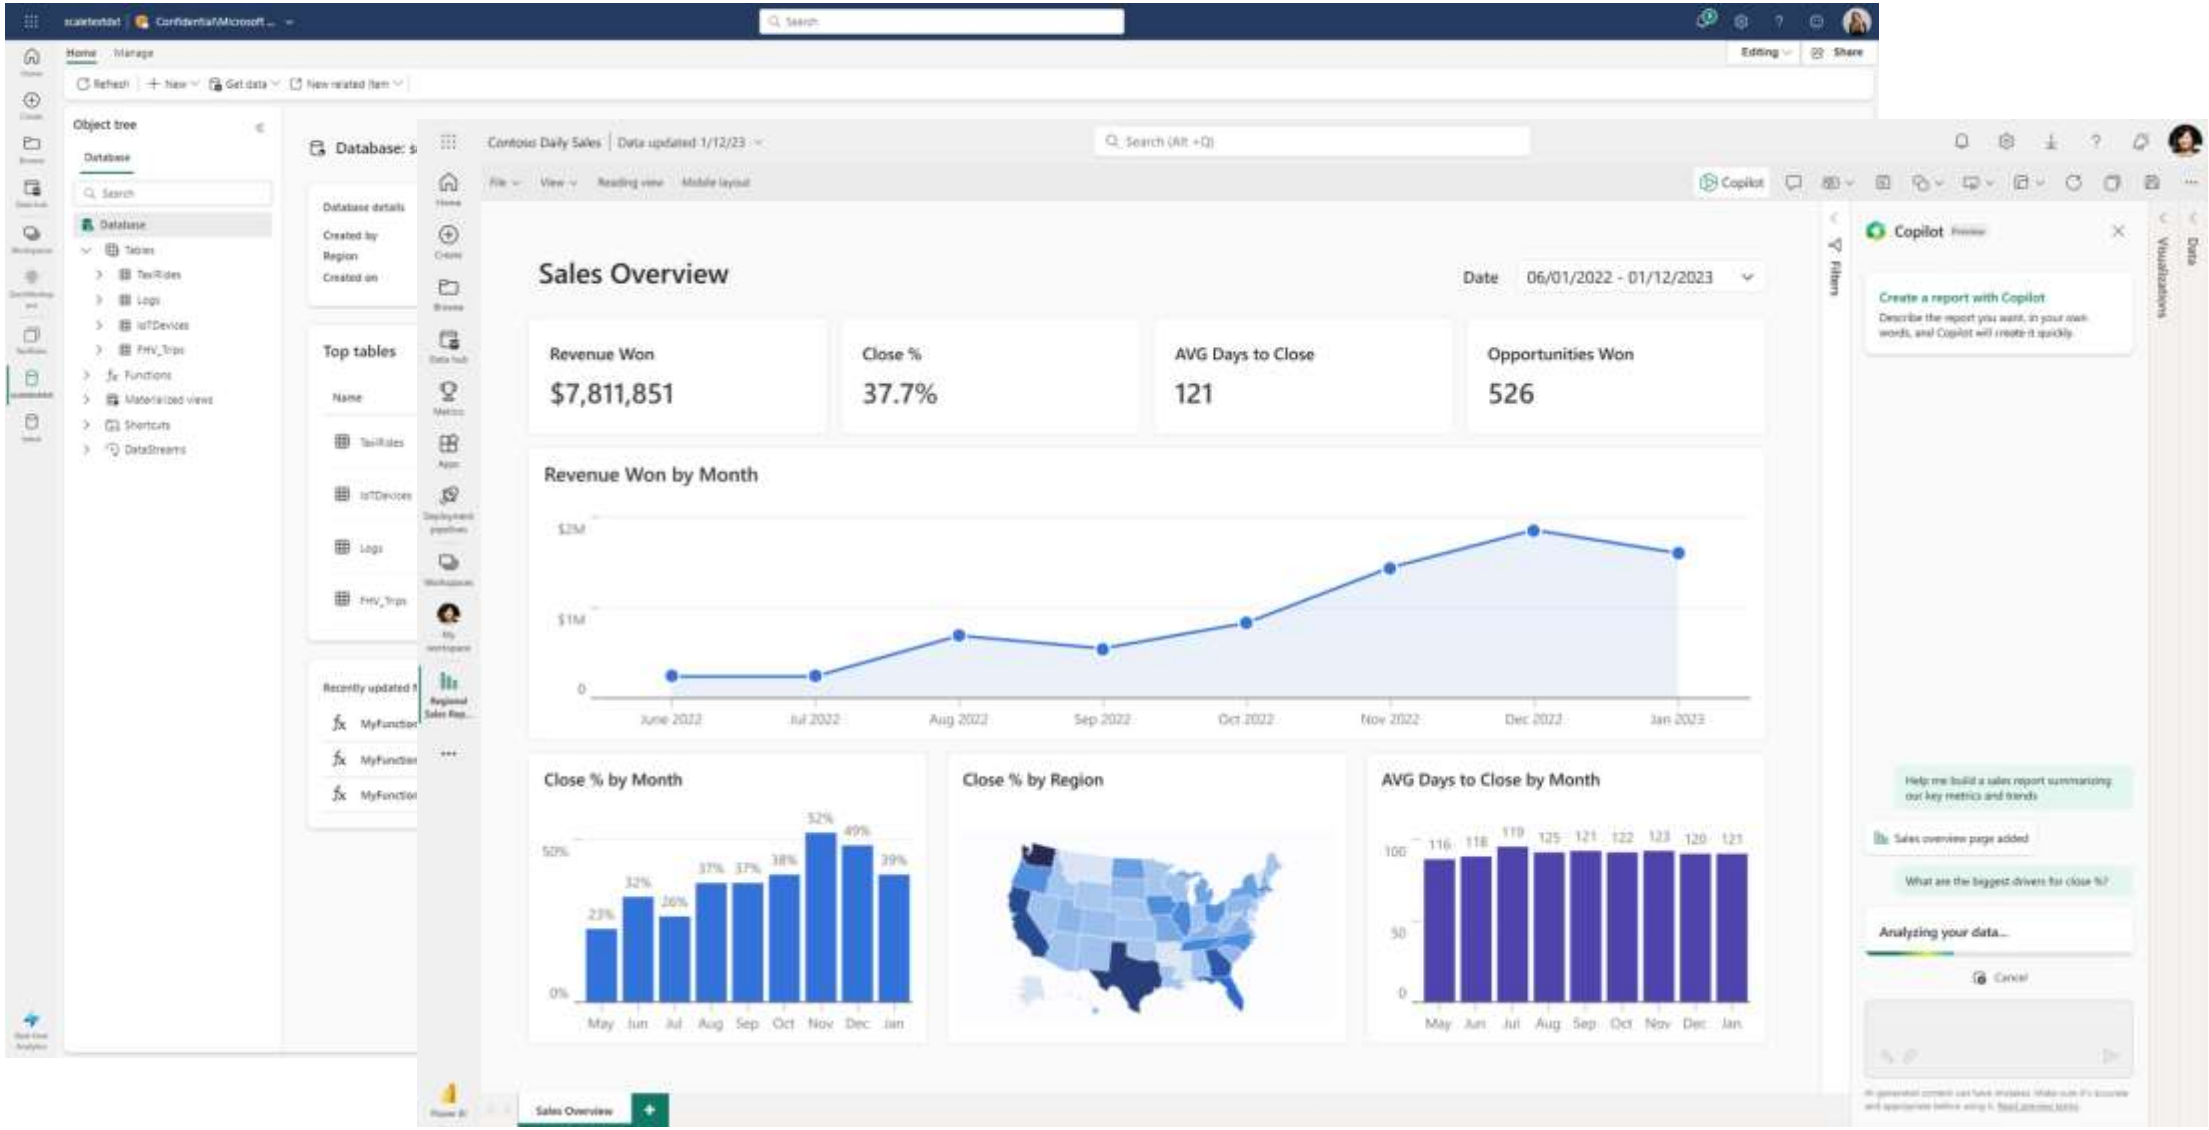
Sales Overview
Date 06/01/2022 - 01/12/2023
Revenue Won
$7,811,851
Close %
37.7%
AVG Days to Close
121
Opportunities Won
526
Revenue Won by Month
Close % by Month
Close % by Region
AVG Days to Close by Month
Copilot
Create a report with Copilot
Analyzing your data...

AI Assisted Creation in Microsoft Fabric

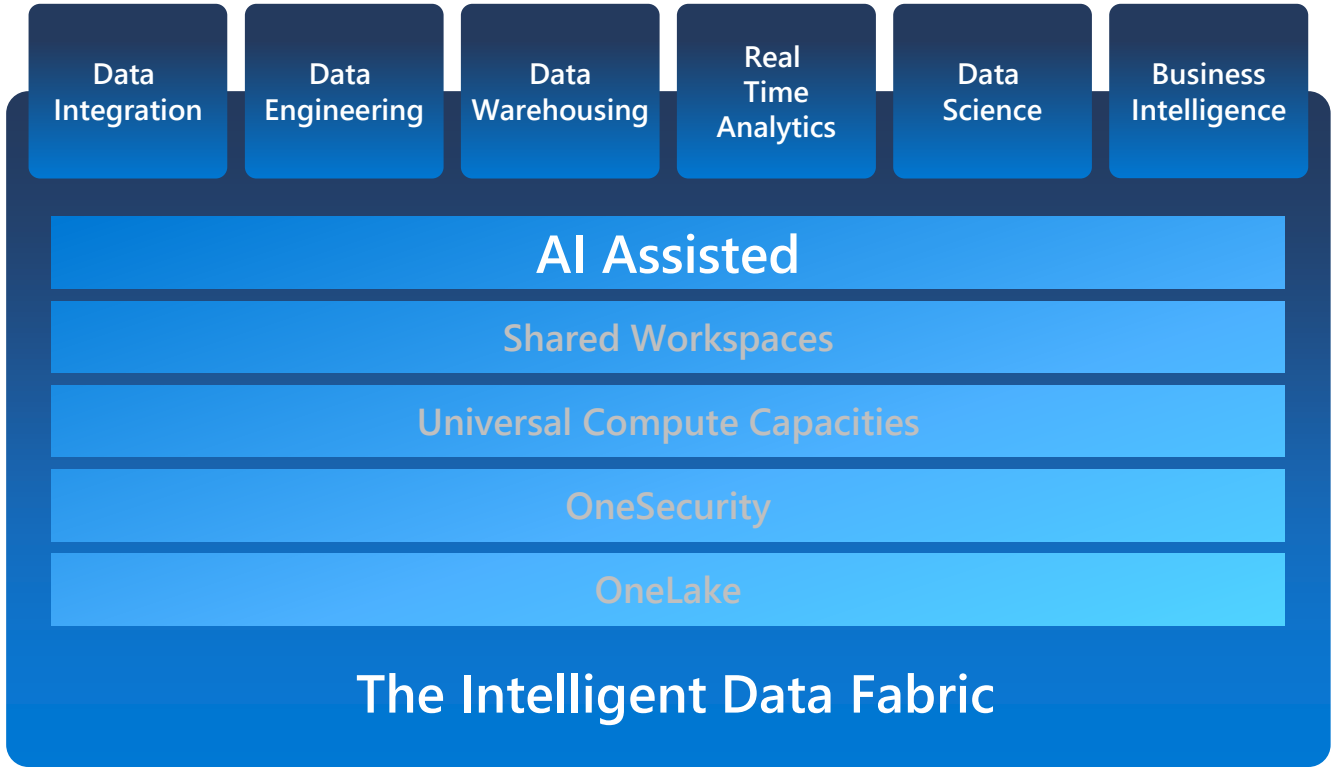

The Fabric platform will include built in Azure Open AI based assistant that will serve all the workloads

First GPT-based feature is already shipping in Power BI - NL2DAX – DAX calculation creation based on natural language prompts

Ongoing major ramp-up for pervasive AOAI based product-wide AI assistance

Lake-centric and open architecture

OneLake for all Data

"The OneDrive for Data"

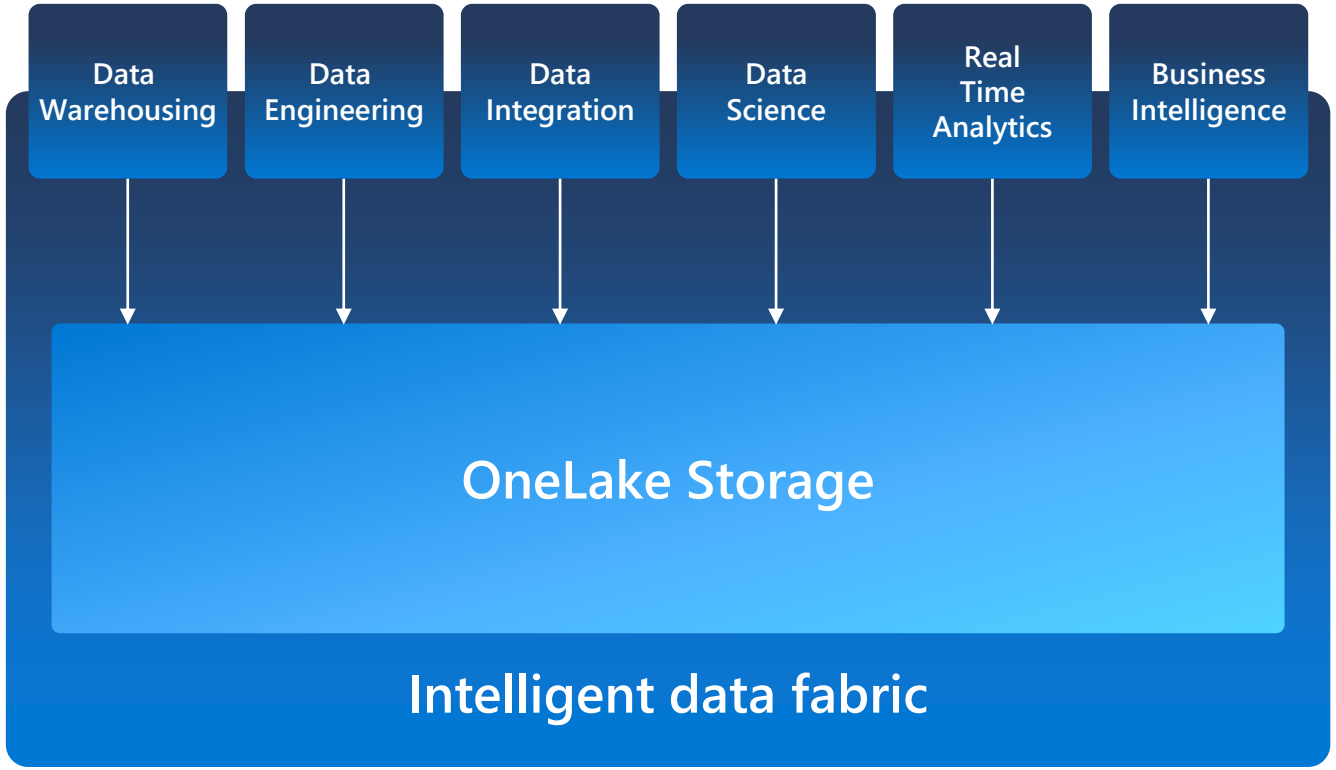

- A single SaaS lake for the whole organization
- Provisioned automatically with the tenant
- All workloads automatically store their data in the OneLake workspace folders
- All the data is organized in an intuitive hierarchical namespace
- The data in OneLake is automatically indexed for discovery, MIP labels, lineage, PII scans, sharing, governance and compliance

OneLake for all domains

A true data mesh across organization domains

Introducing domains as an integral part of Fabric: A domain is a way to logically group together all the data in an organization relevant to an area or field, according to business needs

Domains are defined with domain admins and contributors who can associate workspaces and group them together under a relevant domain

Federated governance can be achieved by delegating settings to domain admins, thus allowing them to achieve more granular control over their business area

Domains simplify discovery and consumption of data across the organization, thus allowing business optimized consumption

Avoid data swamps by endorsing certain data as certified or promoted, thus encouraging reuse.

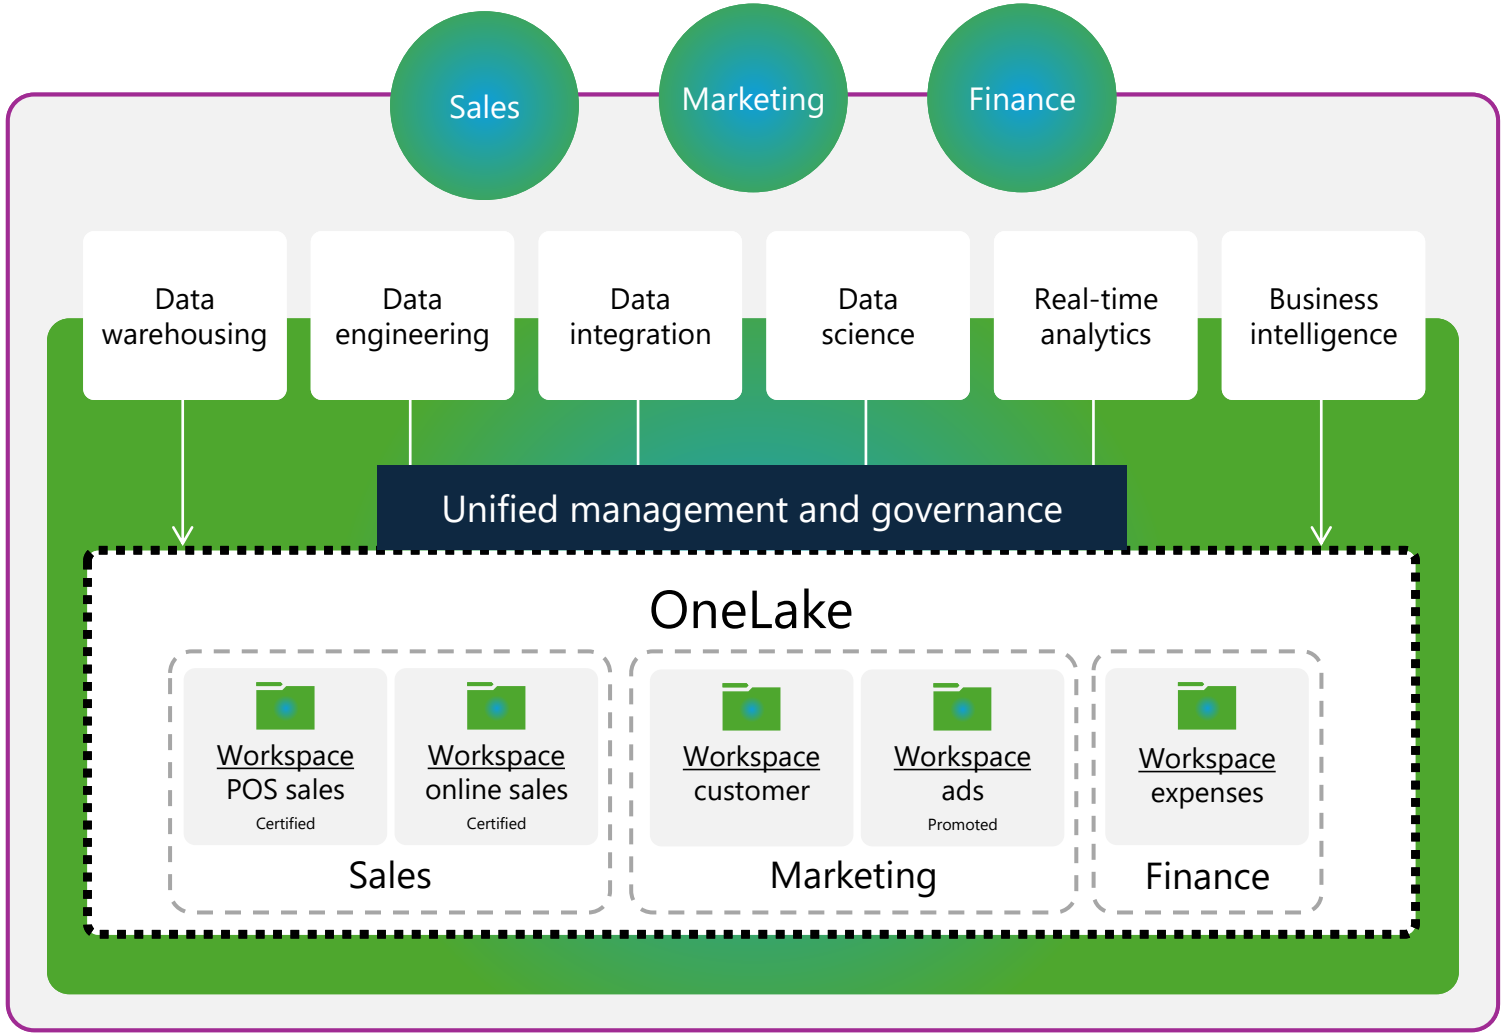

One Copy for all computes

Real separation of compute and storage

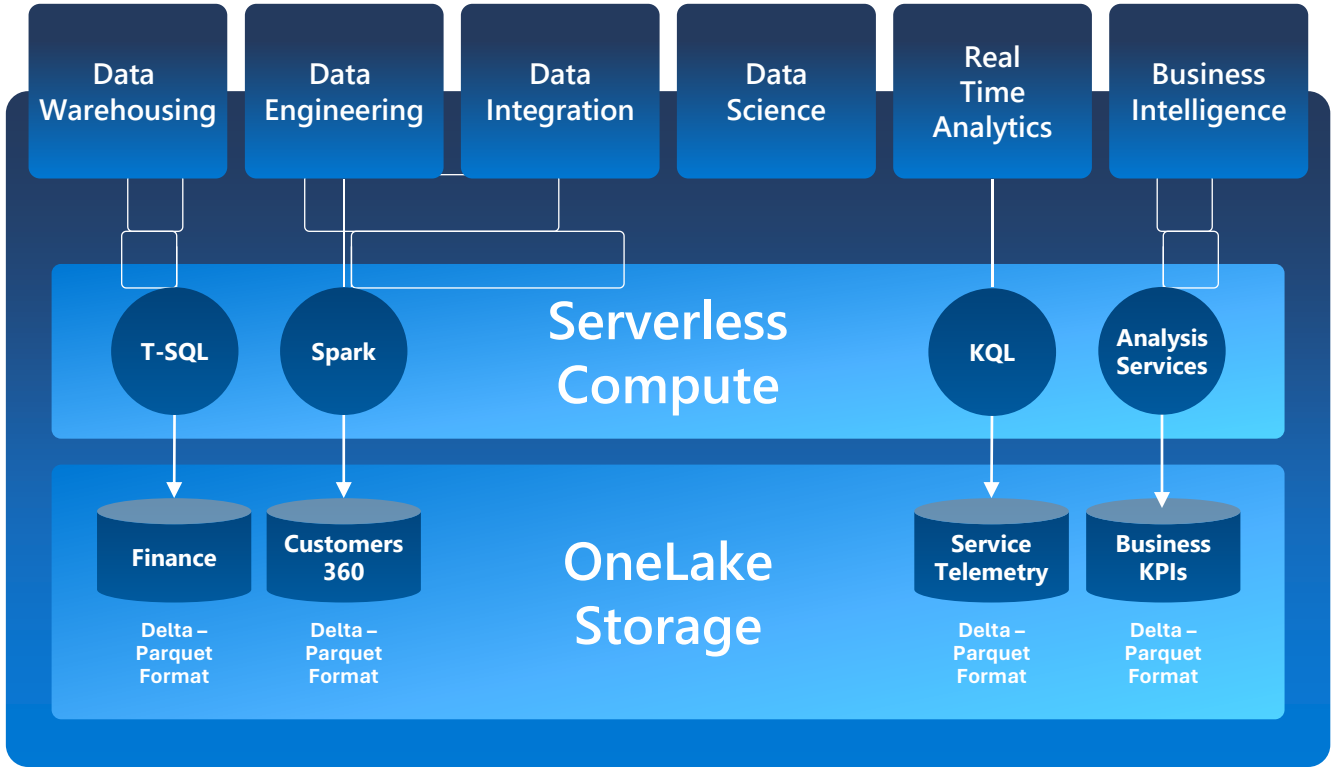

All the compute engines store their data automatically in OneLake

The data is stored in a single common format

Delta – Parquet, an open standards format, is the storage format for all tabular data in Fabric

Once data is stored in the lake, it is directly accessible by all the engines without needing any import/export

All the compute engines have been fully optimized to work with Delta Parquet as their native format

Shared universal security model is enforced across all the engines

Taking One Copy to the next level

Shortcuts

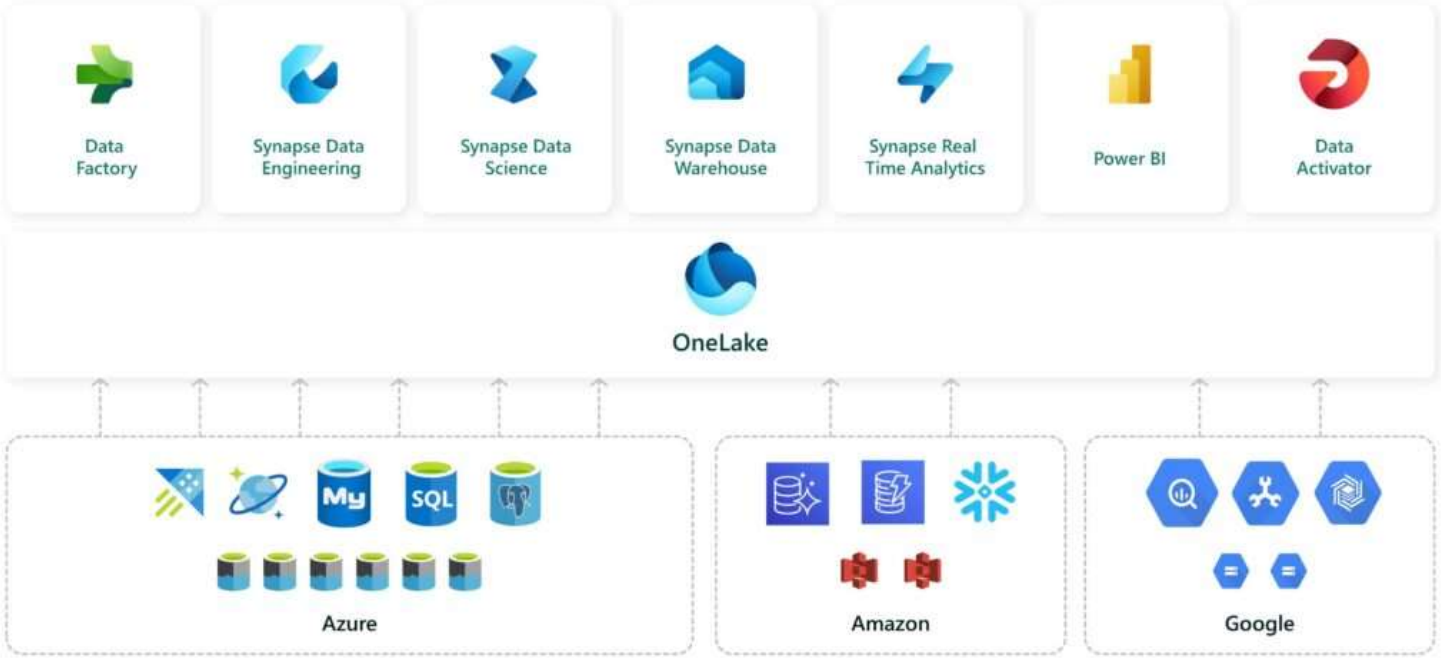

Coming soon

We will soon enable **Azure Cosmos DB**, **Azure SQL DB**, **Snowflake**, and **Mongo DB** customers to use mirroring to access their data in OneLake, with more data sources coming in 2024.

Taking One Copy to the next level

Shortcuts

Sharing data in OneLake is as easy as sharing files in OneDrive, removing the needs for data duplication

With **shortcuts**, data throughout OneLake can be composed together without any data movement

Shortcuts also allow instant linking of data already existing in Azure and in other clouds, without any data duplication and movement, making **OneLake the first multi-cloud data lake**

With support for industry standard APIs, OneLake data can be directly accessed by any application or service

Parquet file format

- Data compression
- Columnar Storage
- Language Agnostics
- Open-source format
- Support for complex data types
- Row-store vs Column-store

→ https://data-mozart.com/parquet-file-format-everything-you-need-to-know/

Data Activator

Trigger actions on all your data, from one place

Trigger alerts & actions from Power BI visuals

Trigger alerts and actions from visuals in Power BI reports

Alert yourself, or define alerts that go to others in your organization.

Supports both report consumers and report creators

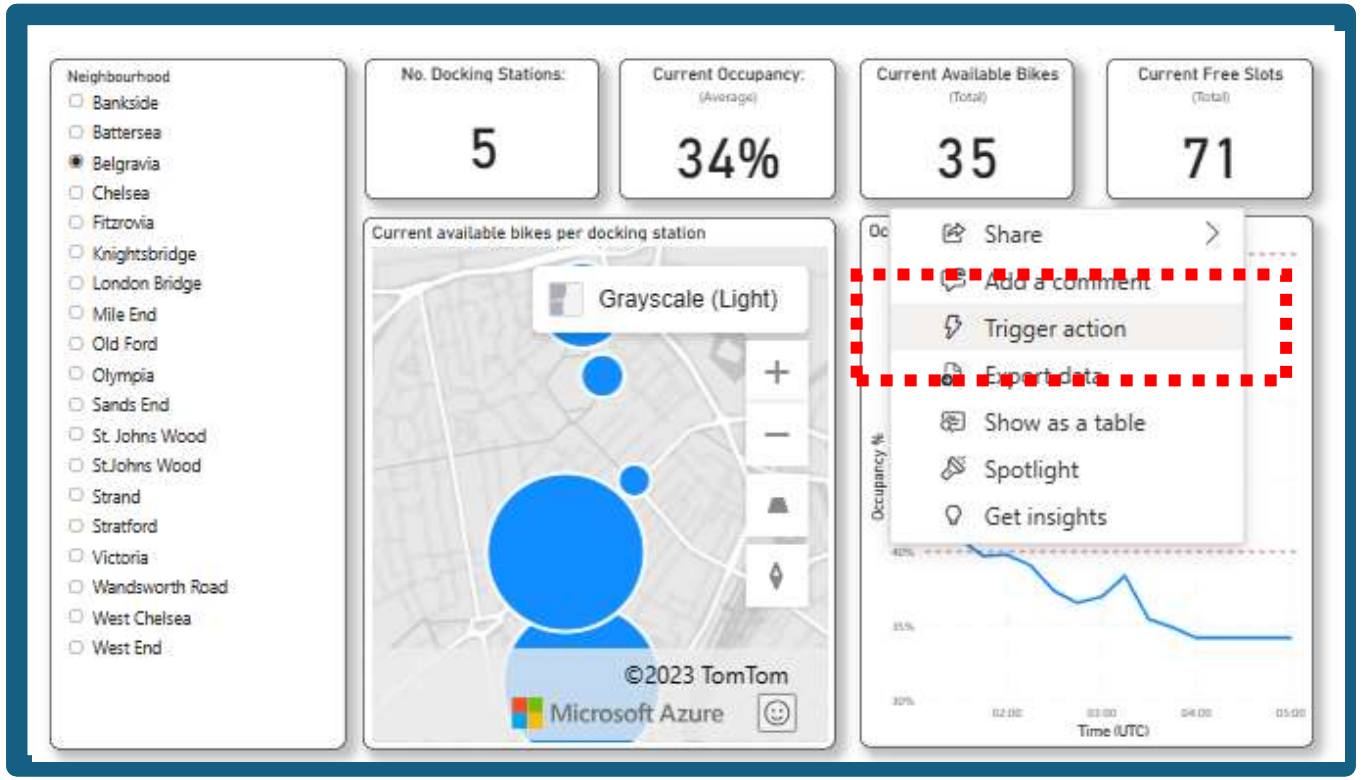

Send alerts, or drive actions via Power Automate

Send alerts via Microsoft Teams, or via email when a trigger condition is met.

Trigger a Power Automate flow to connect to other alerting systems, to log a ticket in a ticketing system, and more

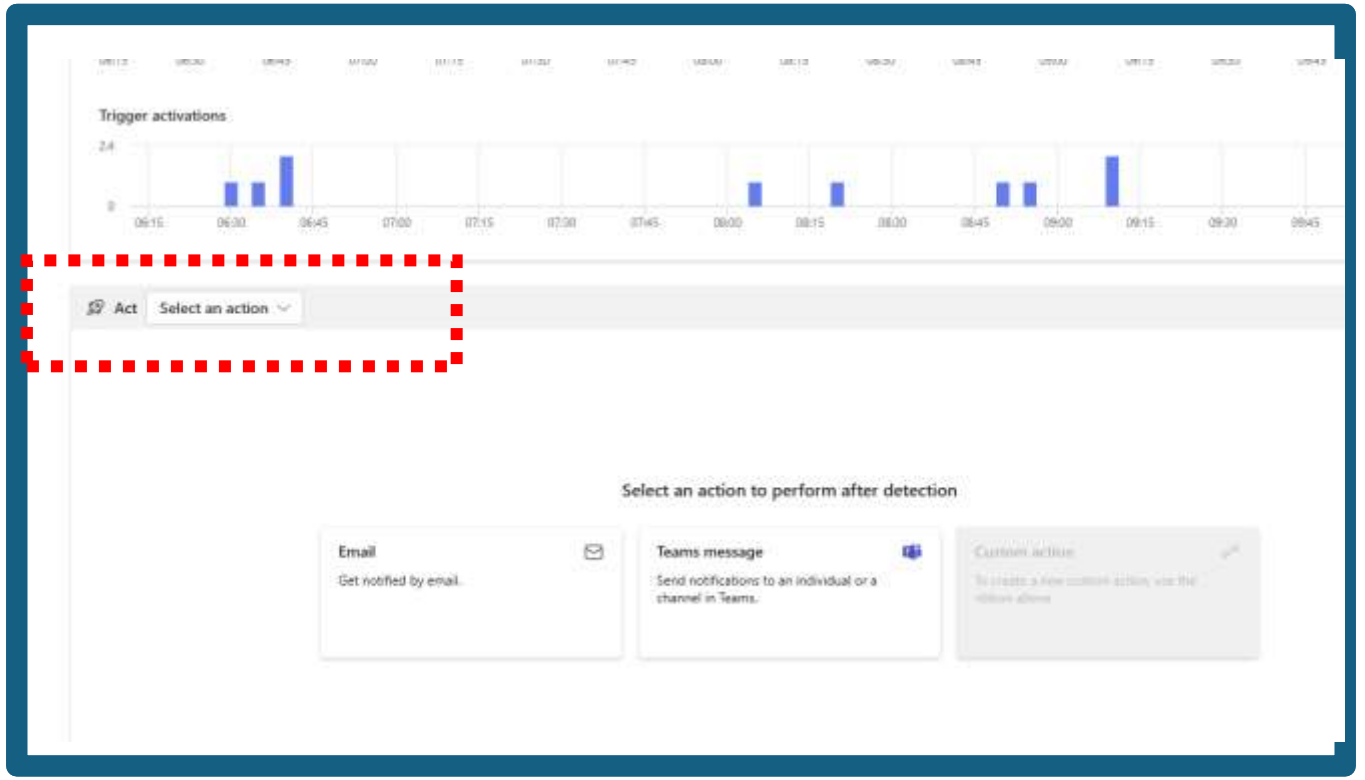

Copilot & Explore

Copilot for...

A new way to transform and analyze data, generate insights, and create visualizations and reports

Copilot and other **generative AI** features in preview

- Data Science and Data Engineering
- Data Factory
- Power BI

Copilot for...

A new way to transform and analyze data, generate insights, and create visualizations and reports

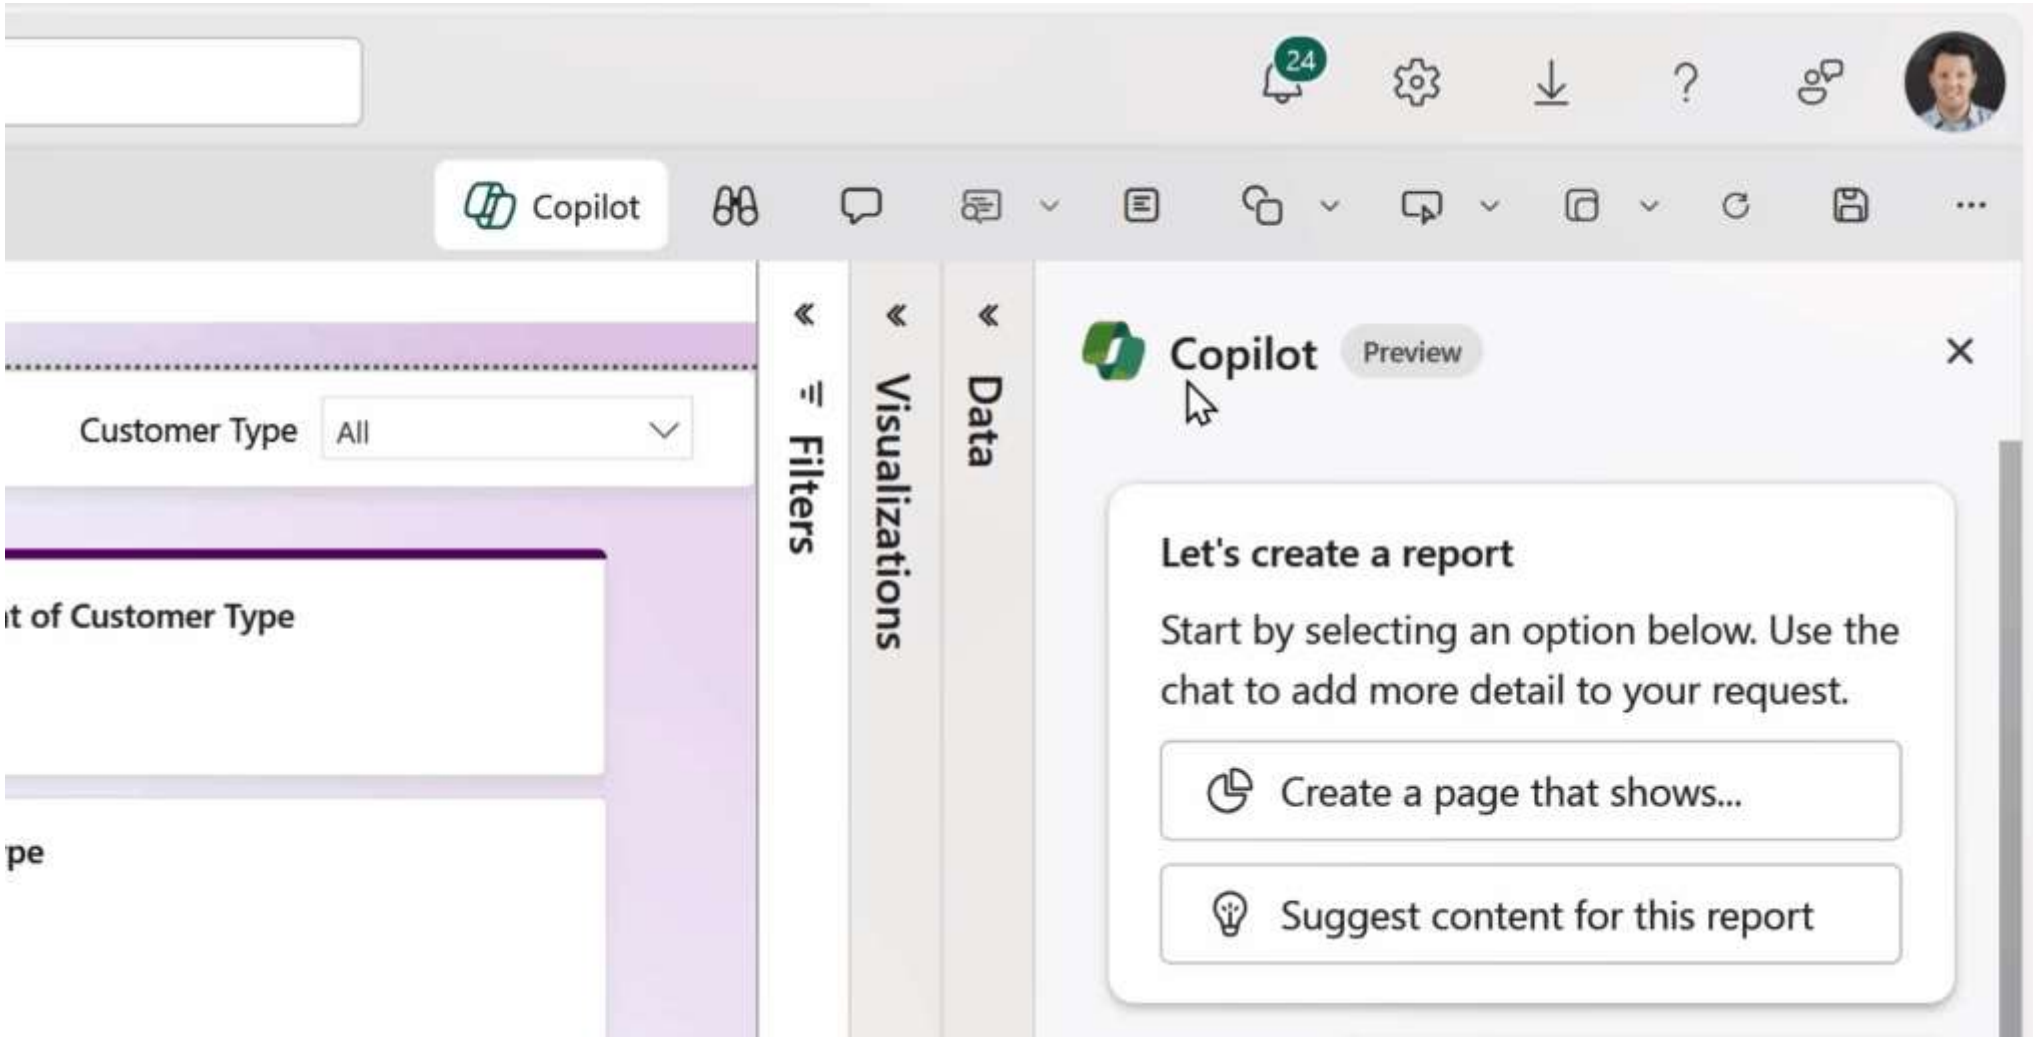

Copilot for...

A new way to transform and analyze data, generate insights, and create visualizations and reports

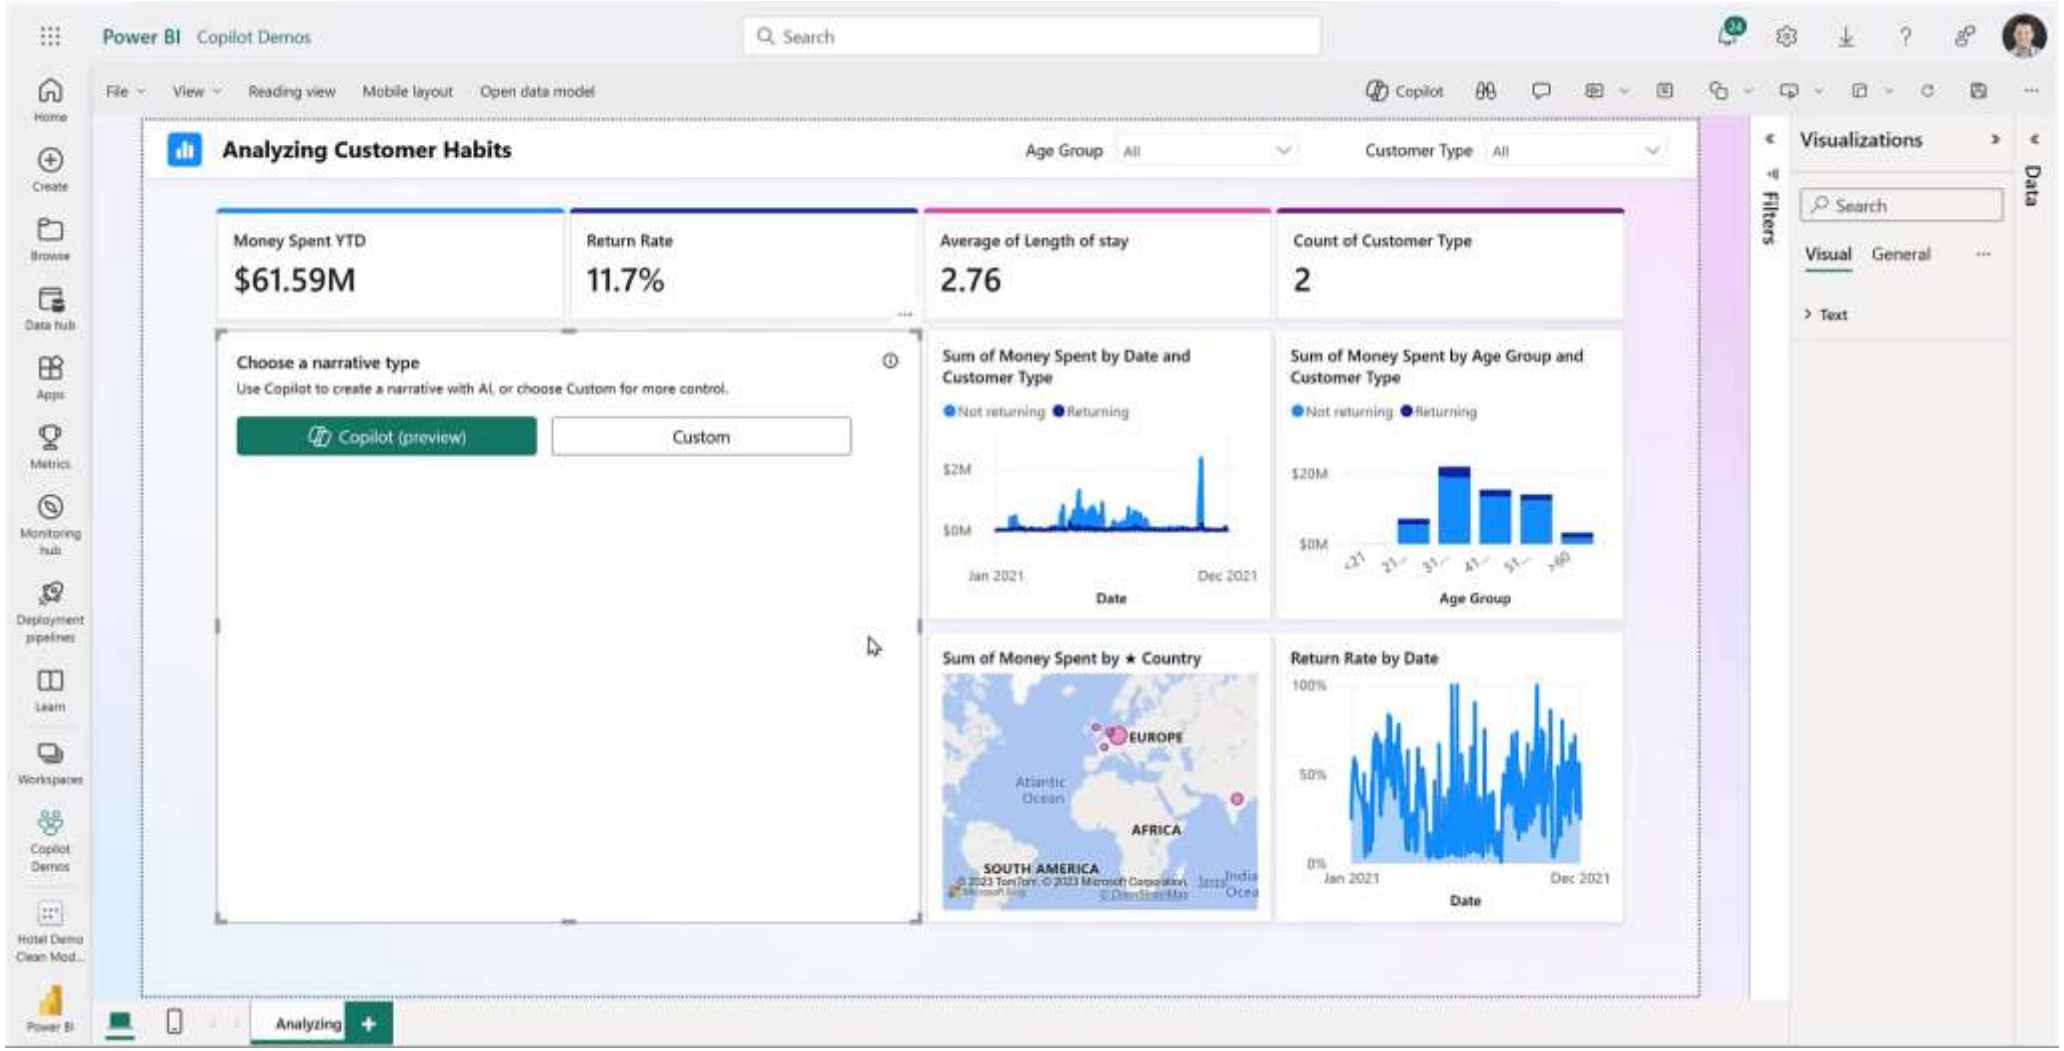

Copilot for Data Science

Copilot Chat Panel

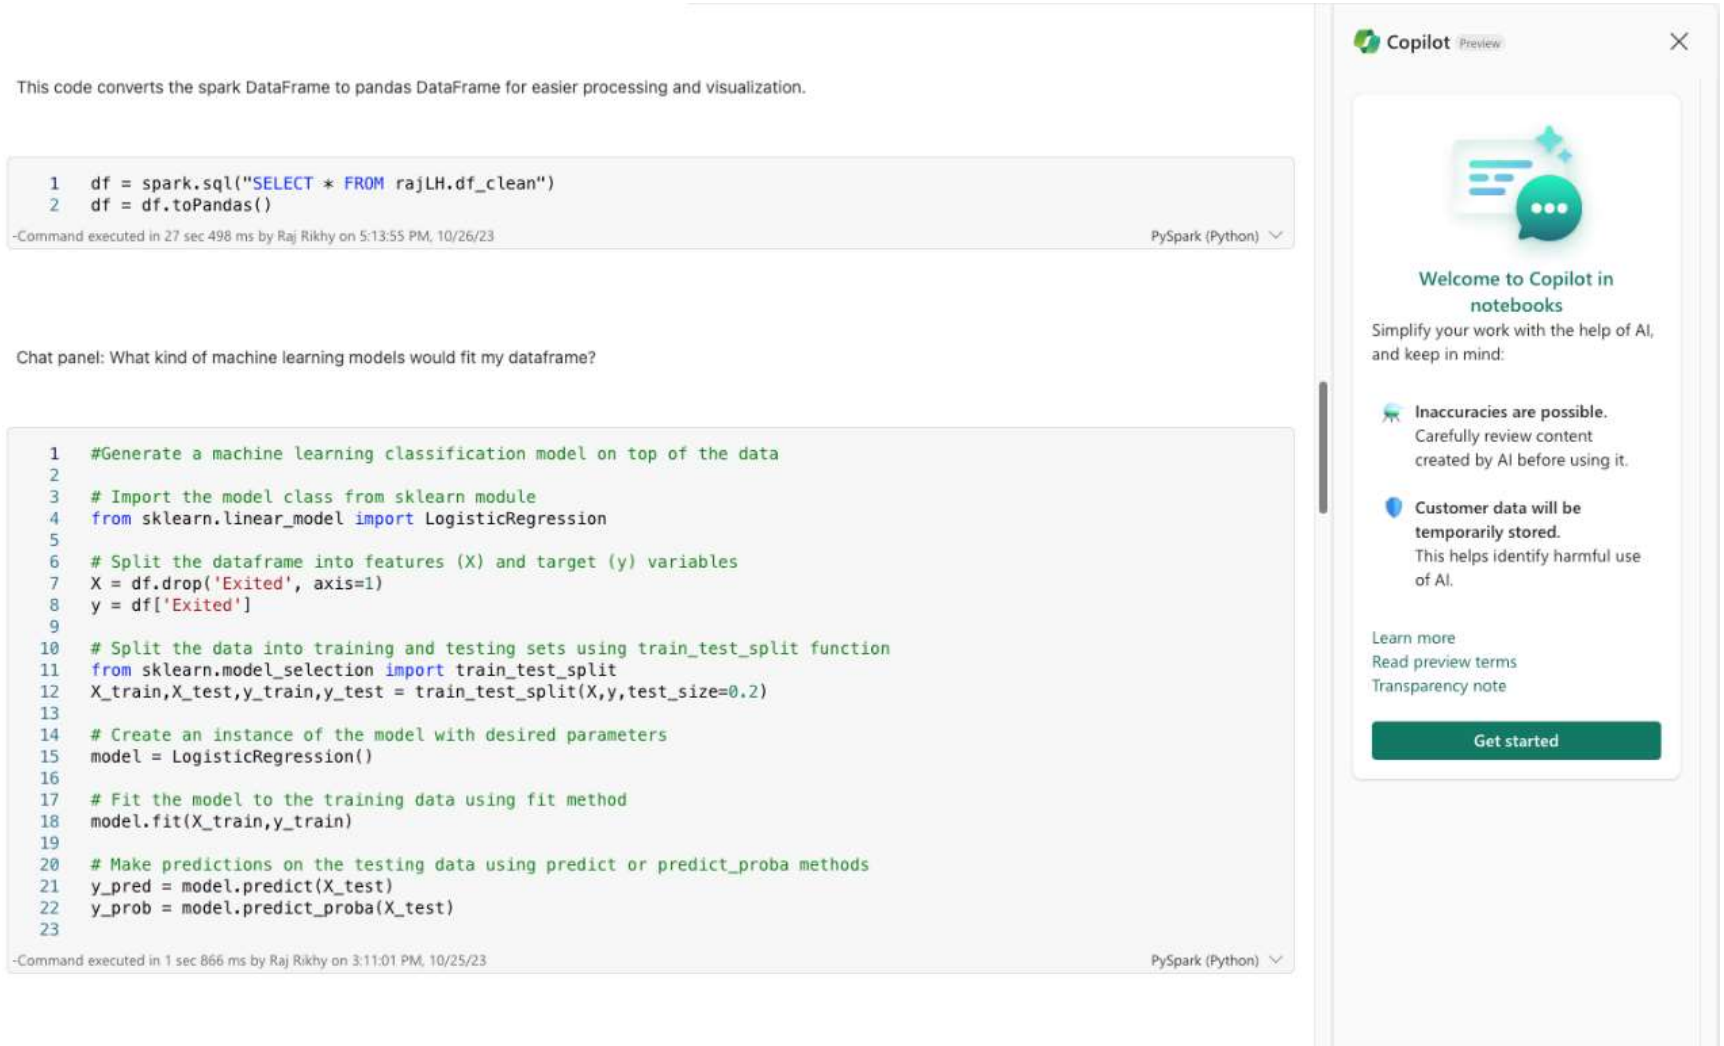

Copilot for Data Science

Copilot Chat Panel

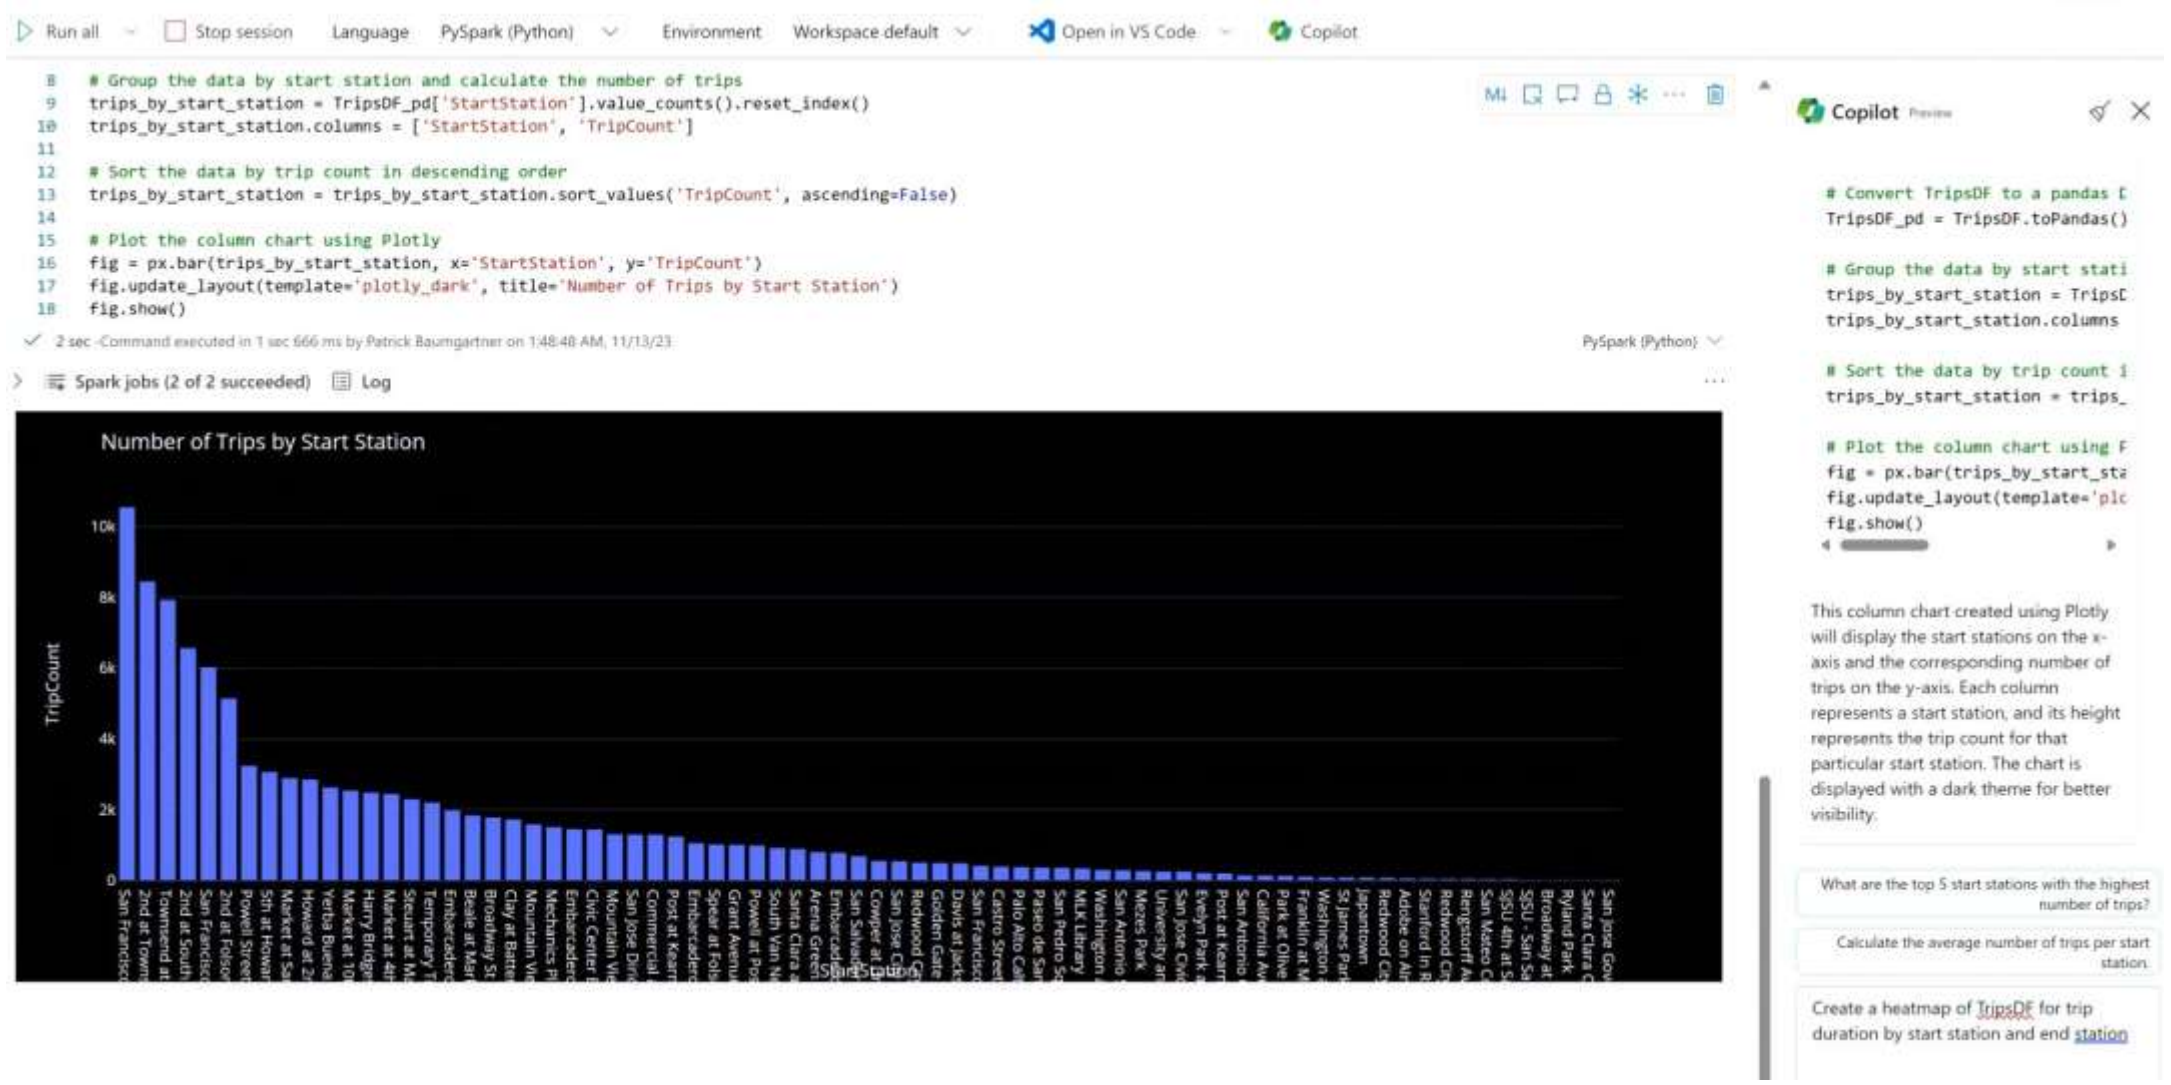

Copilot for Data Science

Chat Magics

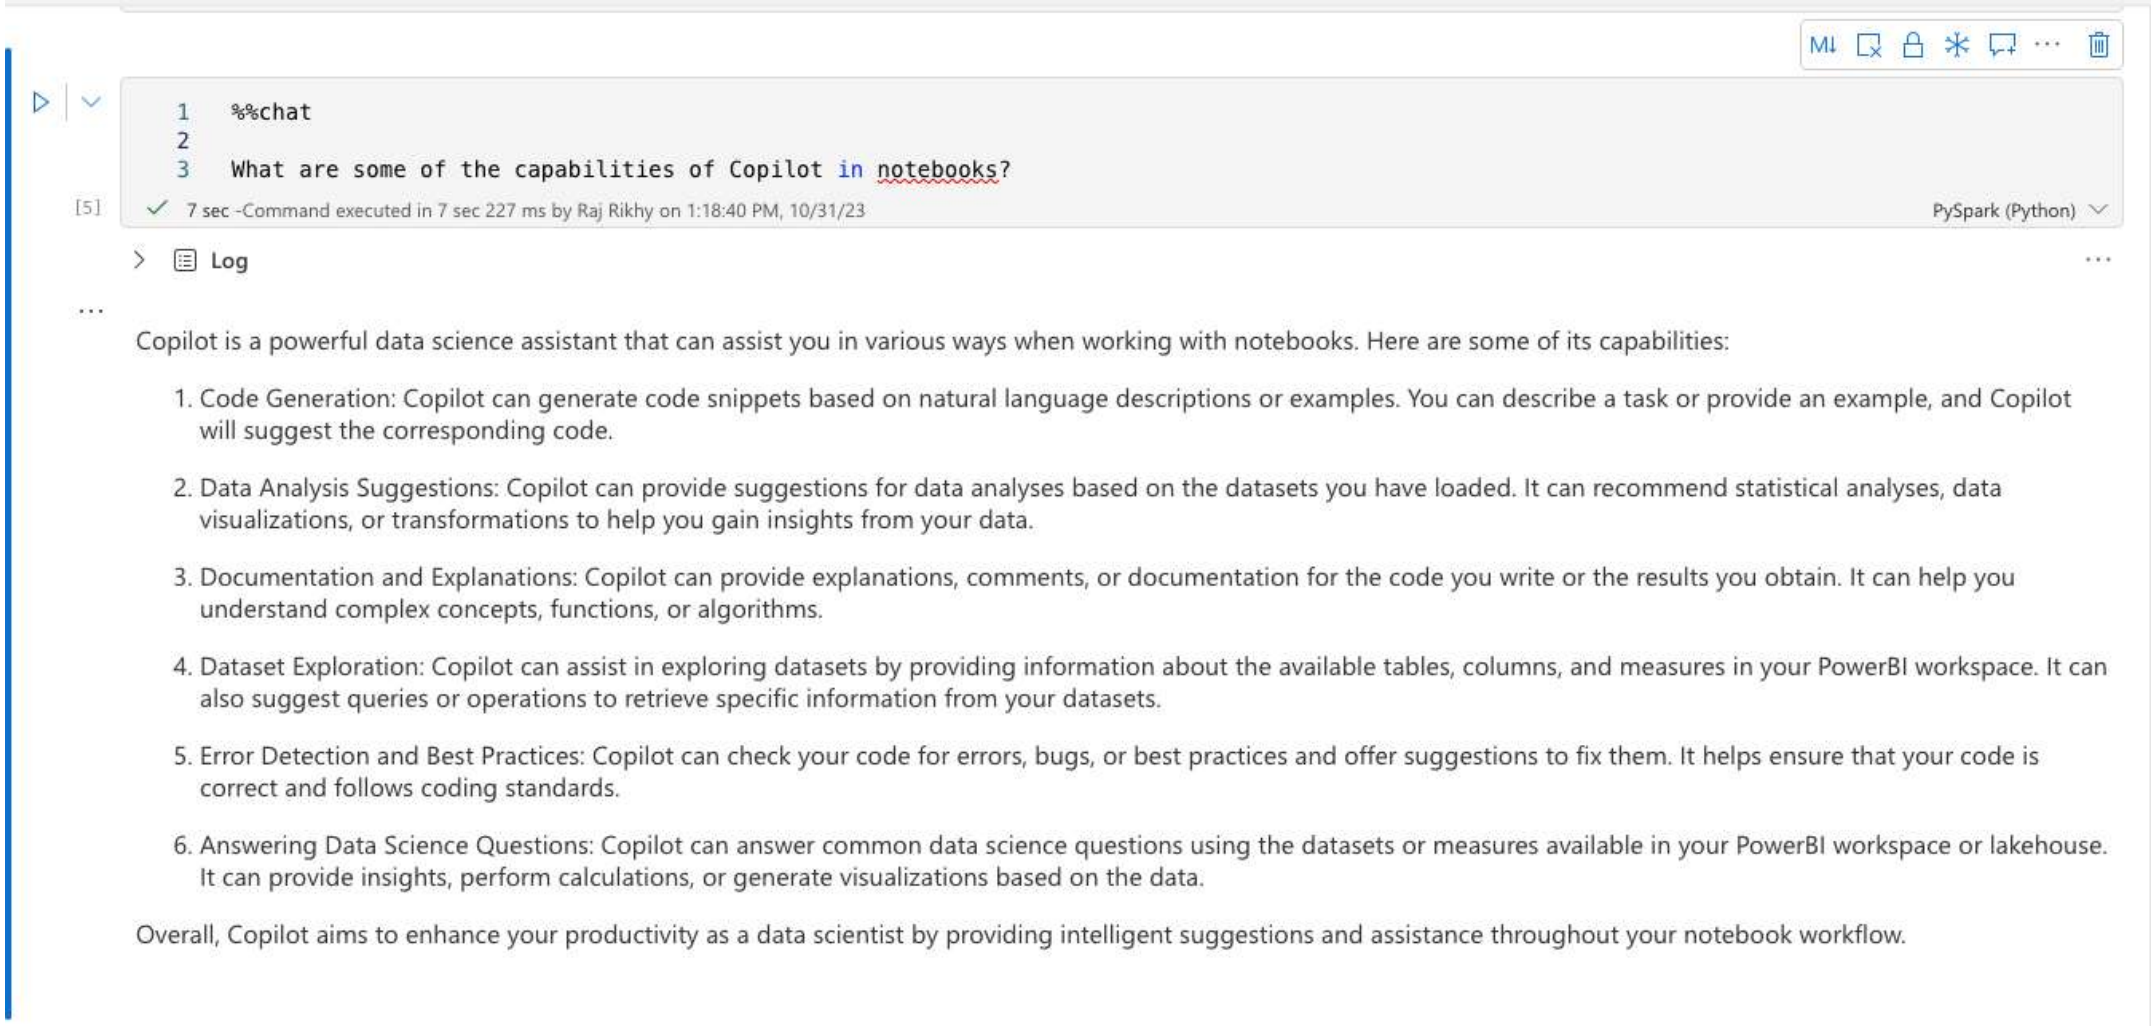

Copilot for DataFlow

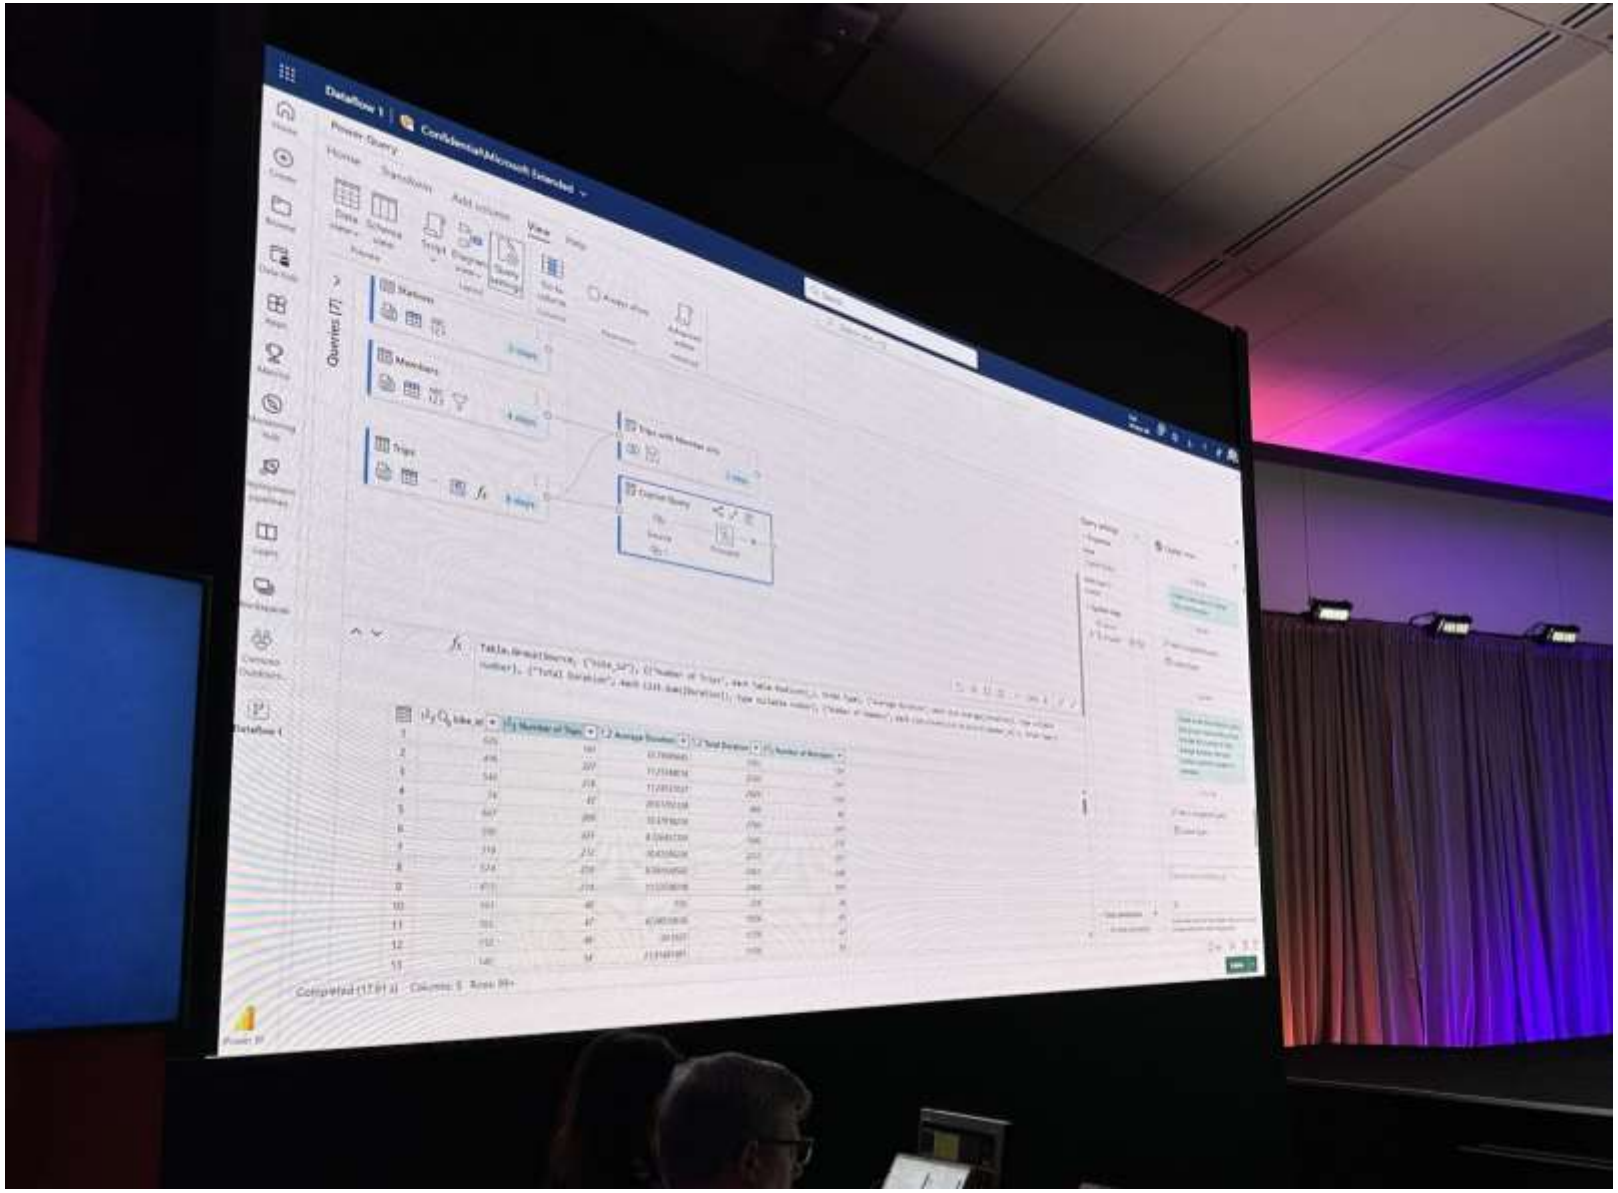

Quickly answer your data questions with Explore

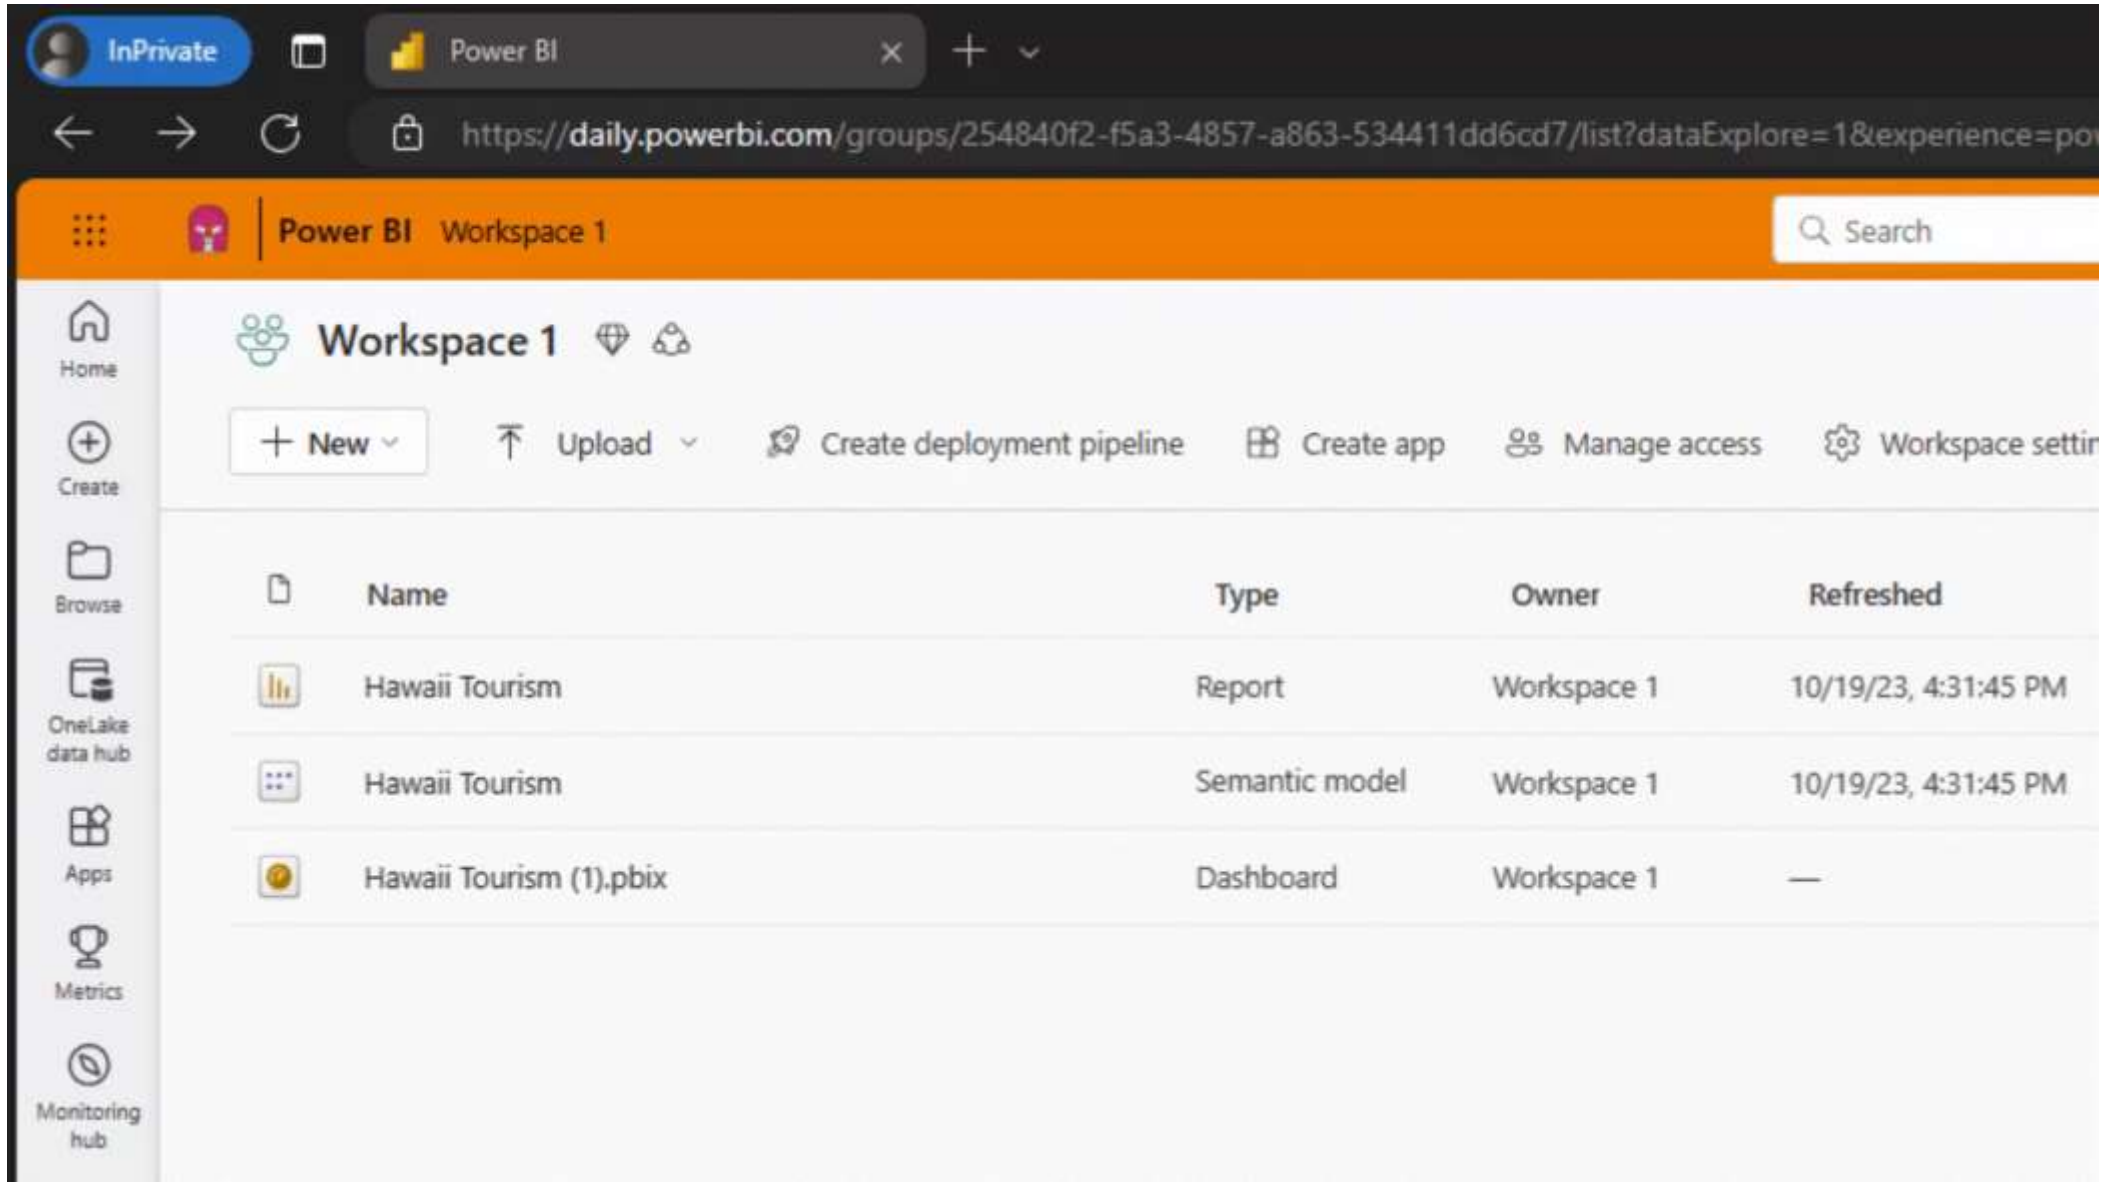

Mirroring


Coming soon
Mirroring in Microsoft Fabric
Seamlessly connect your data warehouses
and databases to OneLake
Azure Cosmos DB
Azure SQL DB
snowflake
MongoDB.
aka.ms/Fabric

Mirroring in Microsoft Fabric

- By just providing connection details, your db is instantly available in Fabric as a Mirrored database
- **Real-time data replication**

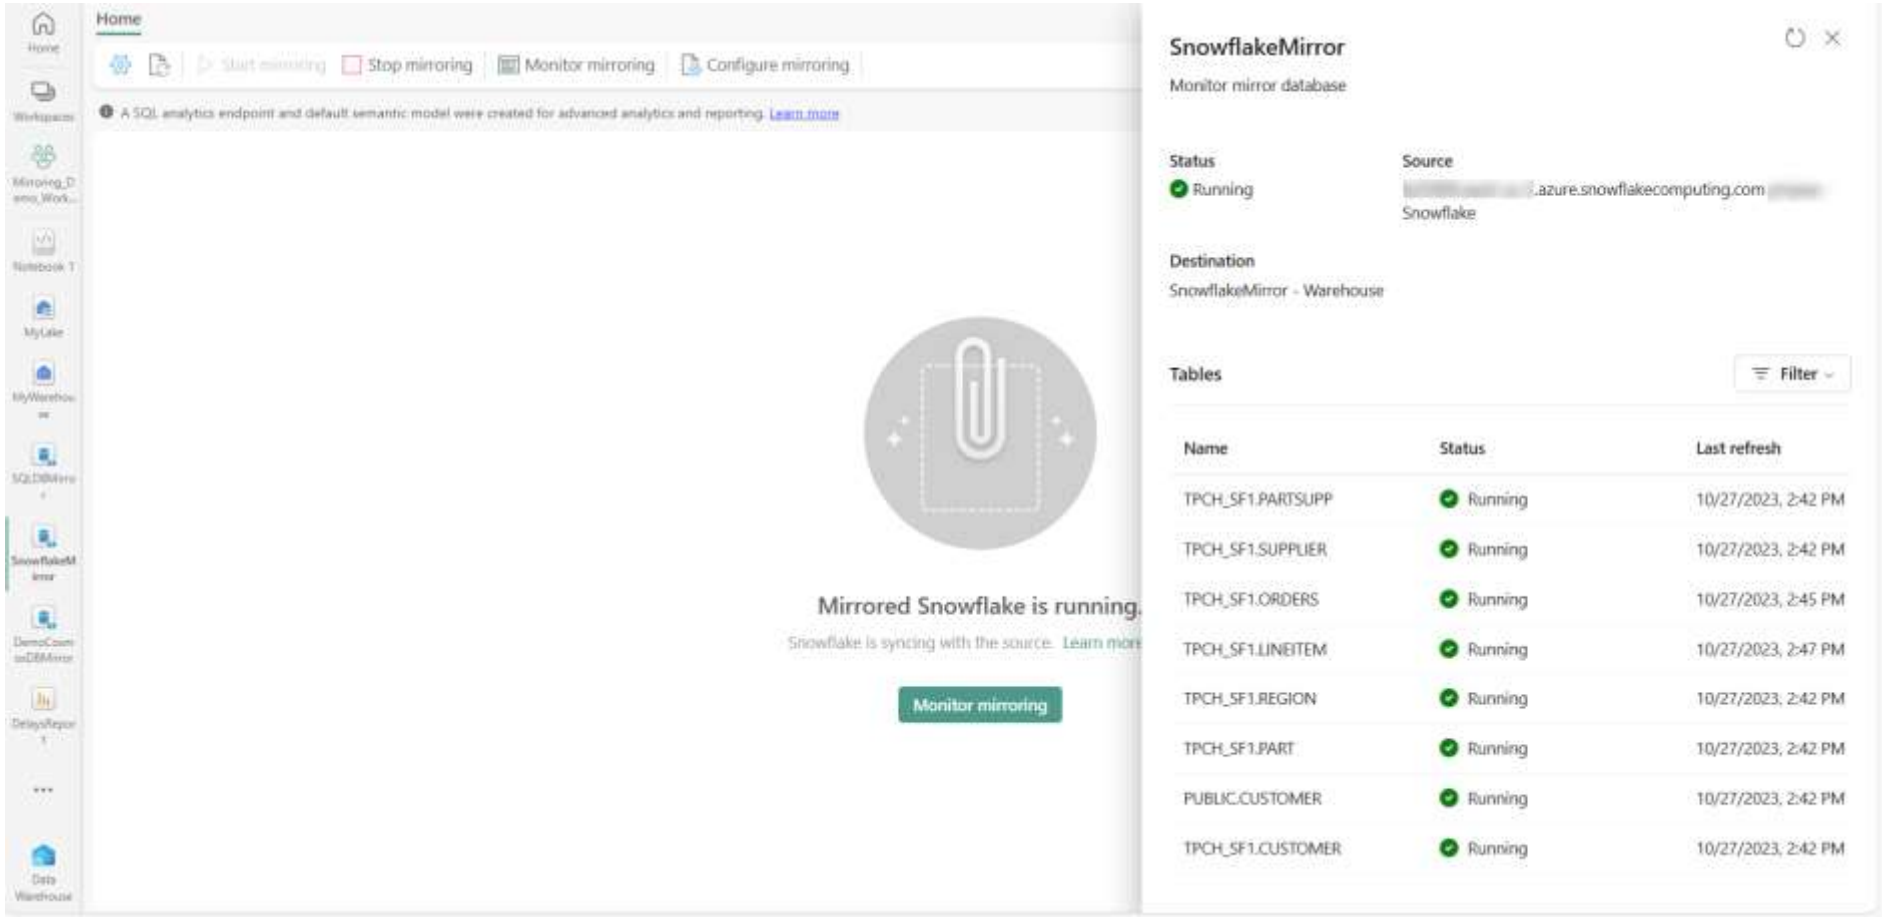

Cross-joining Mirrored db, Warehouses, Lakehouses

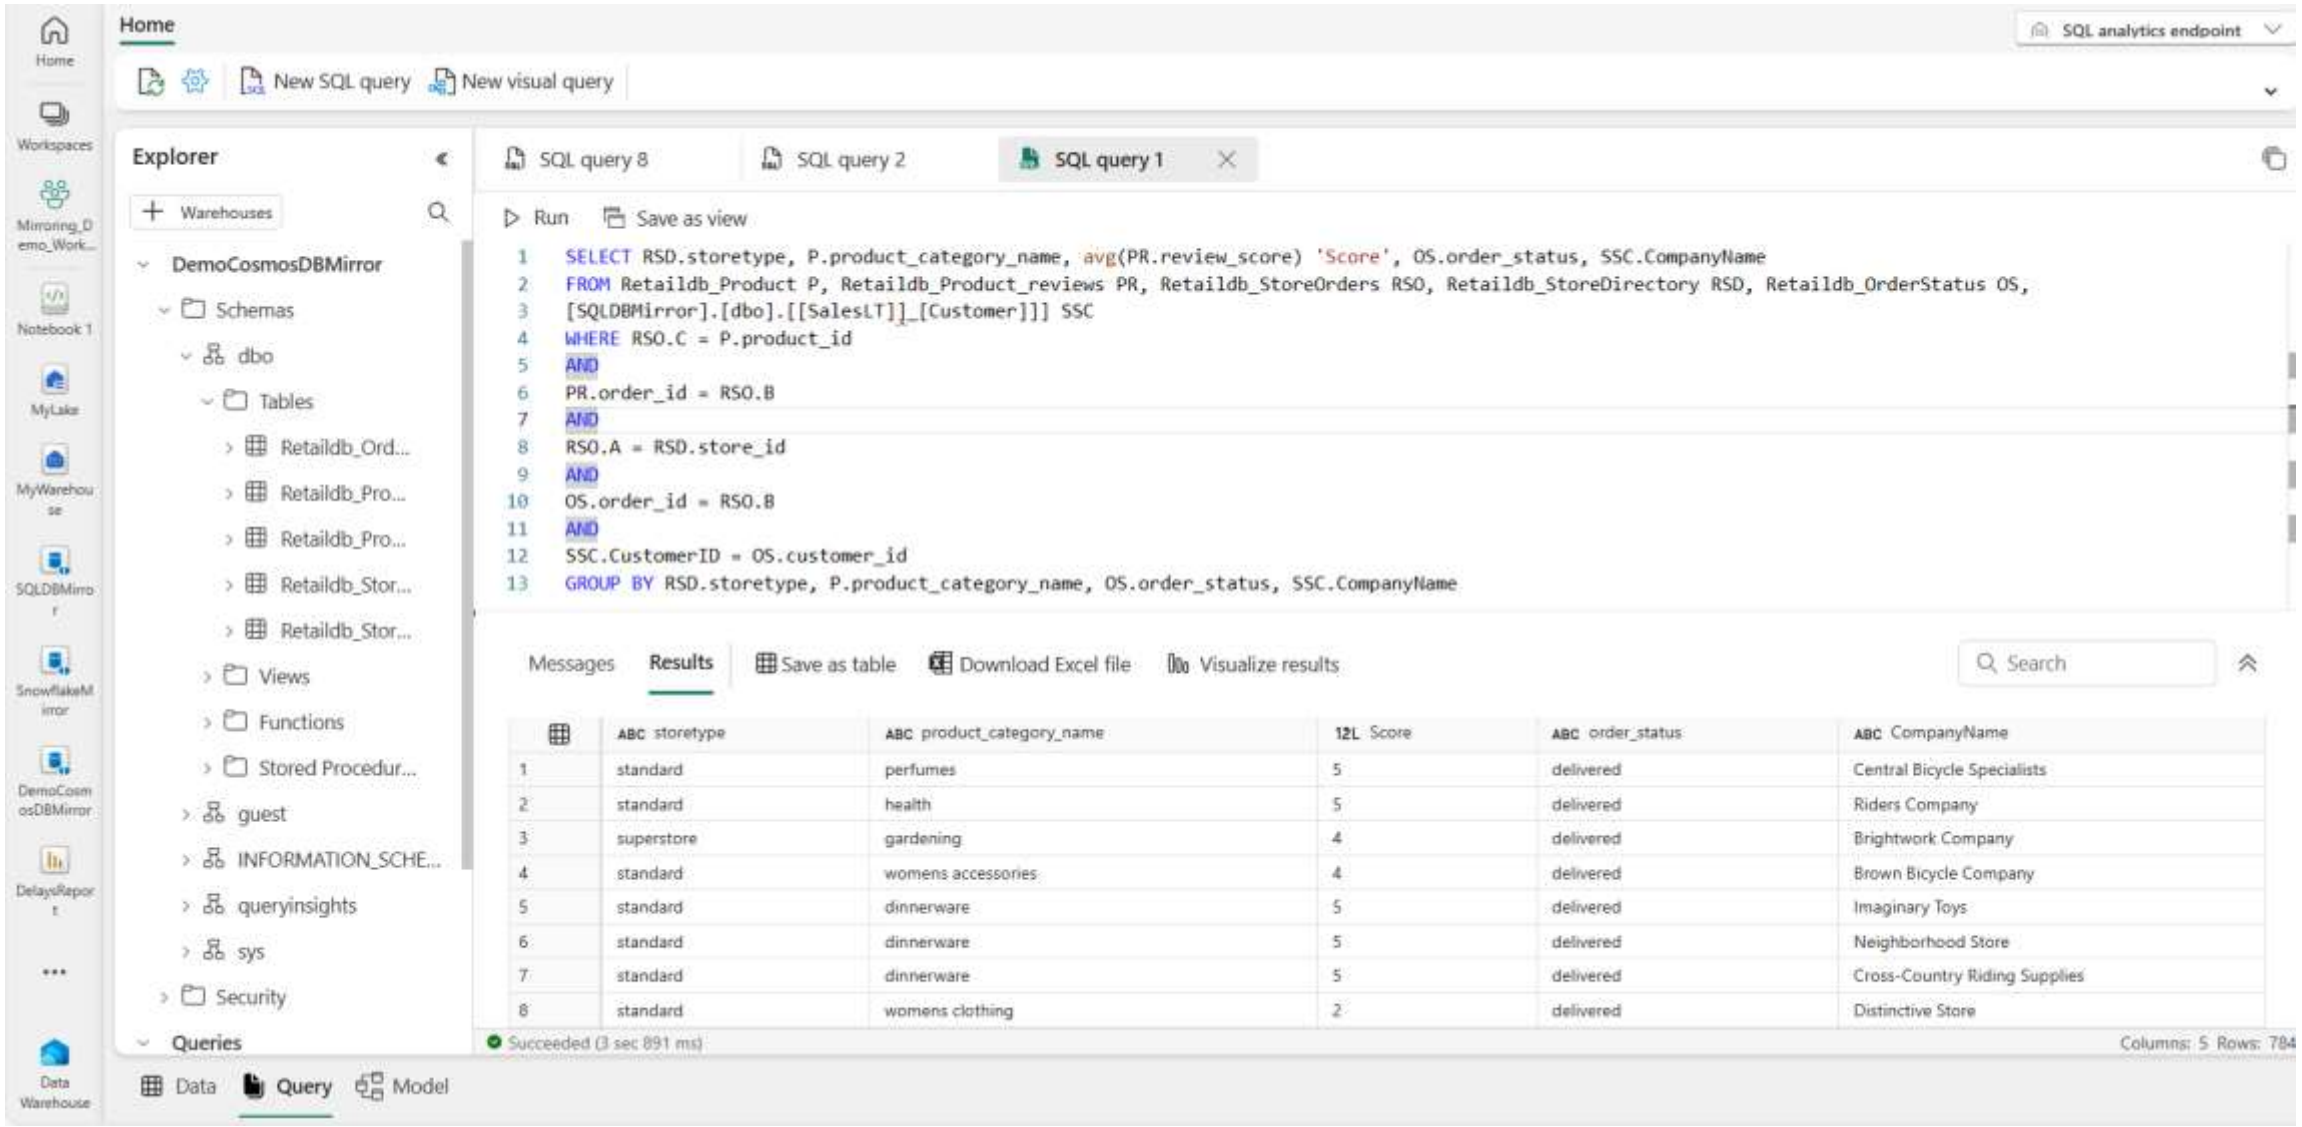

Mirroring Architecture

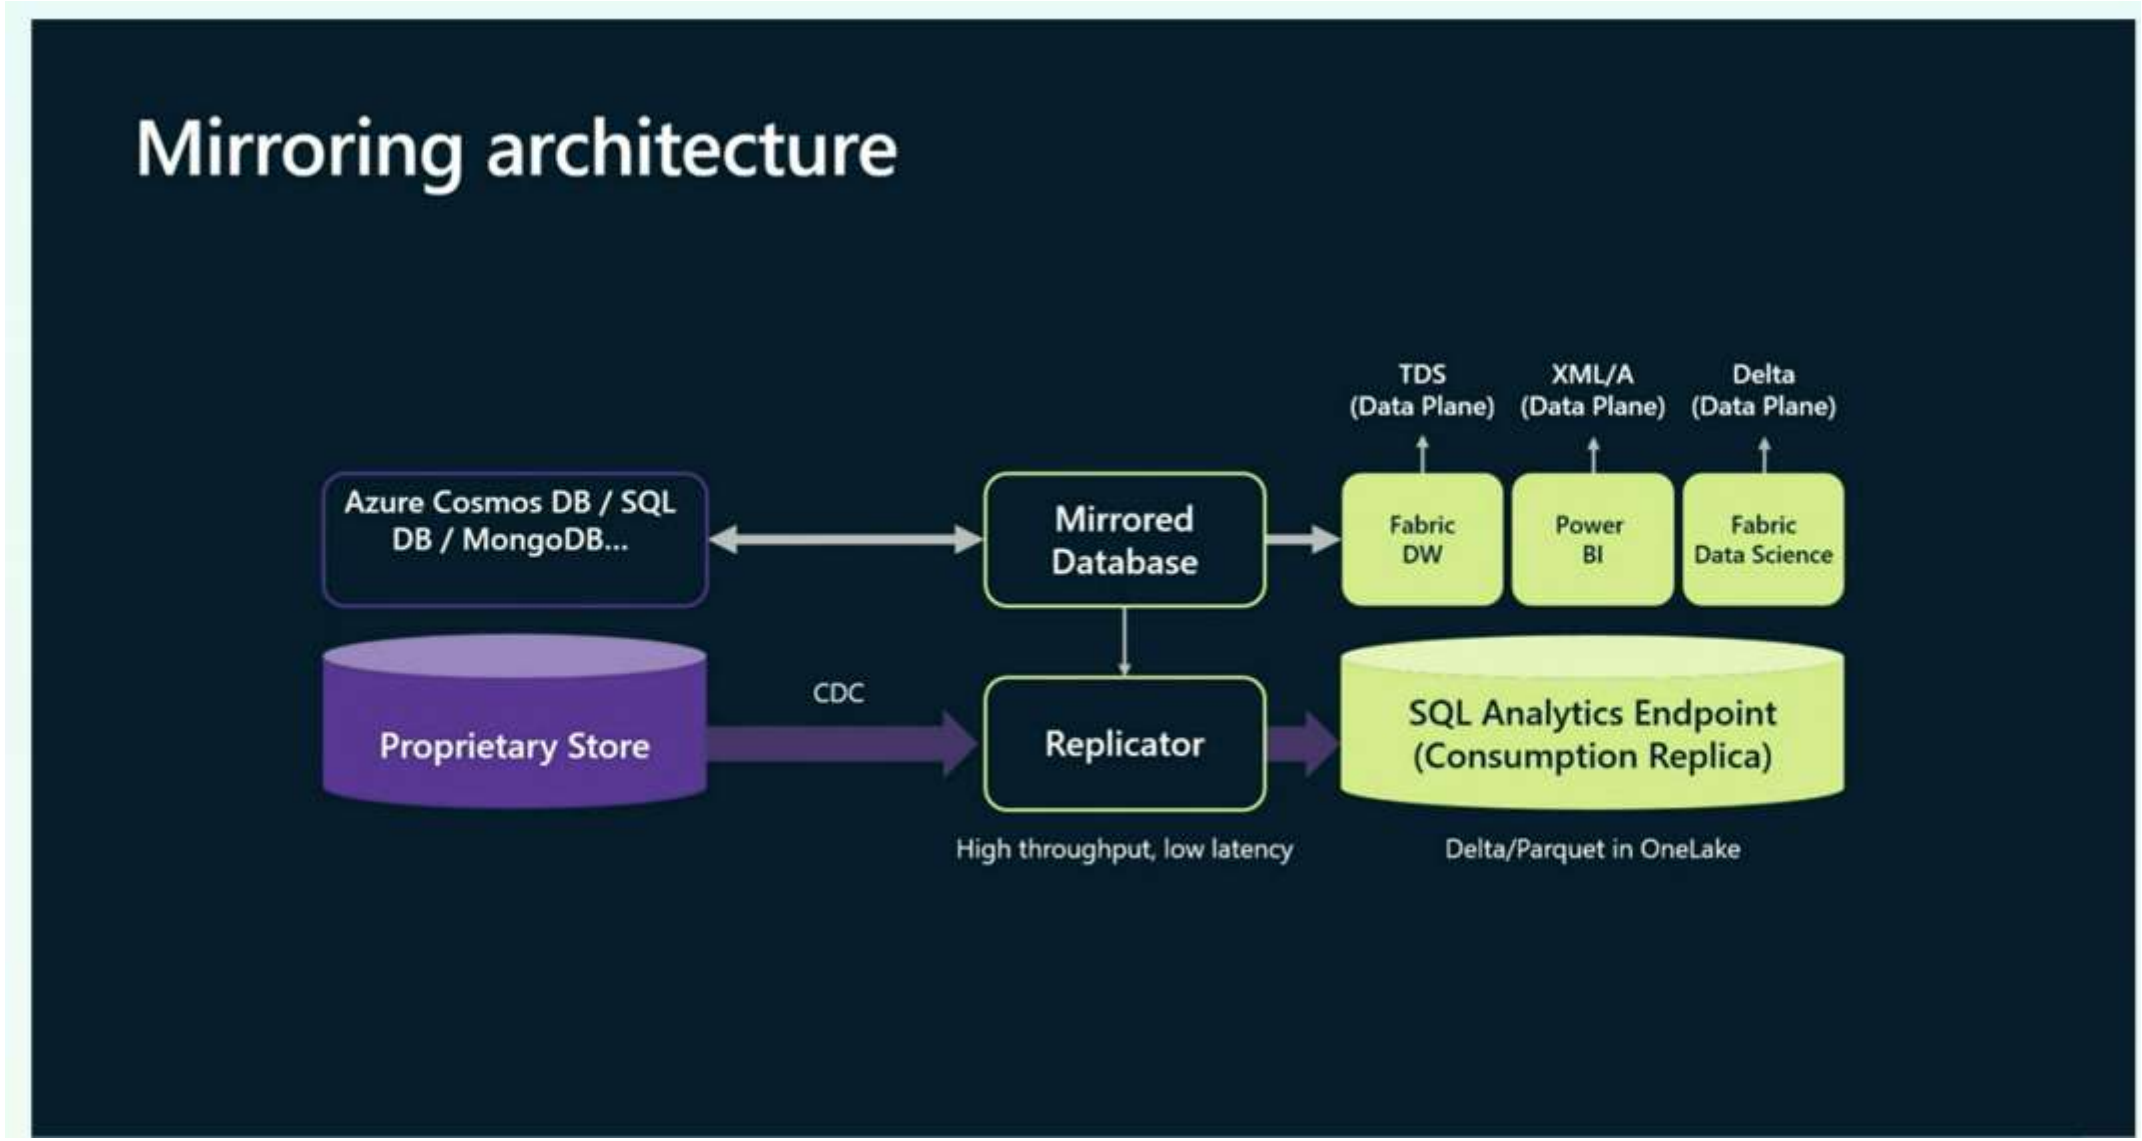

Data Lineage

Includes and spans over to mirrored systems !

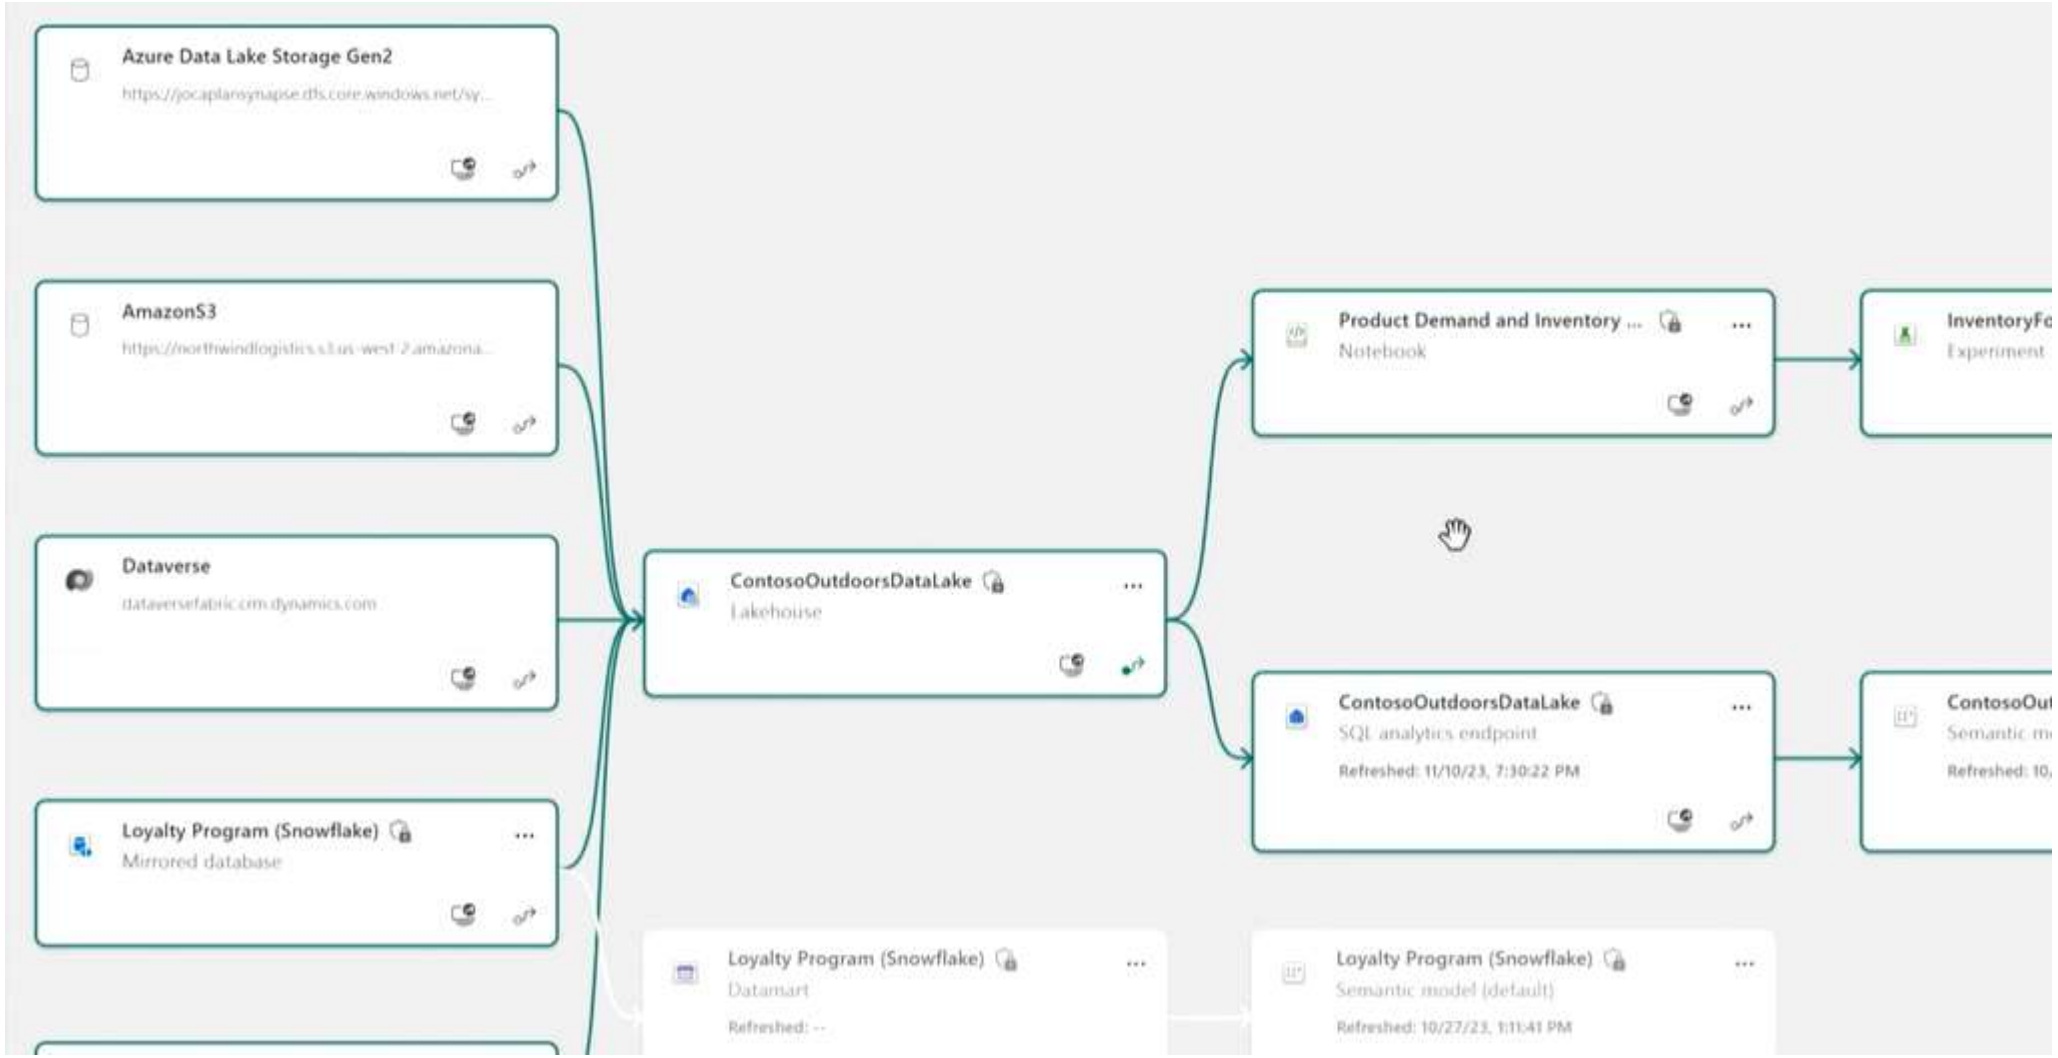

Mirroring Availability

Azure Cosmos DB, **Azure SQL DB** and **Snowflake** customers will be able to use Mirroring to mirror their data in OneLake and unlock all the capabilities of Fabric Warehouse, Direct Lake Mode, Notebooks and much more.

SQL Server, **Azure PostgreSQL**, **Azure MySQL**, **MongoDB** and other databases and data warehouses will be coming in CY24

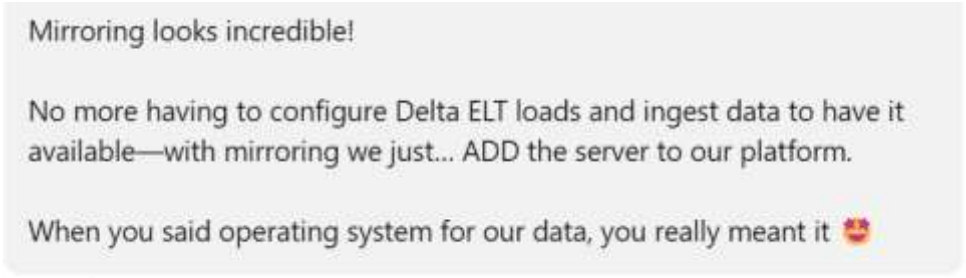

Mirroring Roadmap

Scenarios

Common analytics scenarios

Lakehouse

Data Warehouse

Data Science

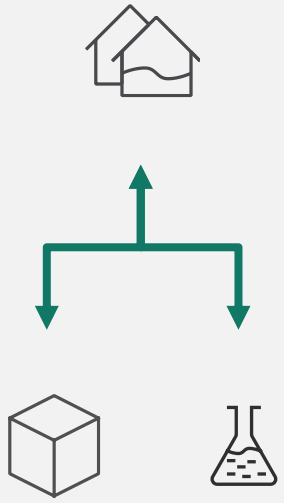

Real Time Analytics

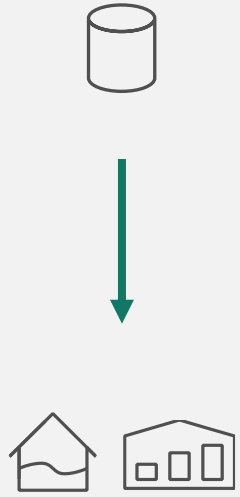

Lakehouse scenario

End-to-end analytics scenario

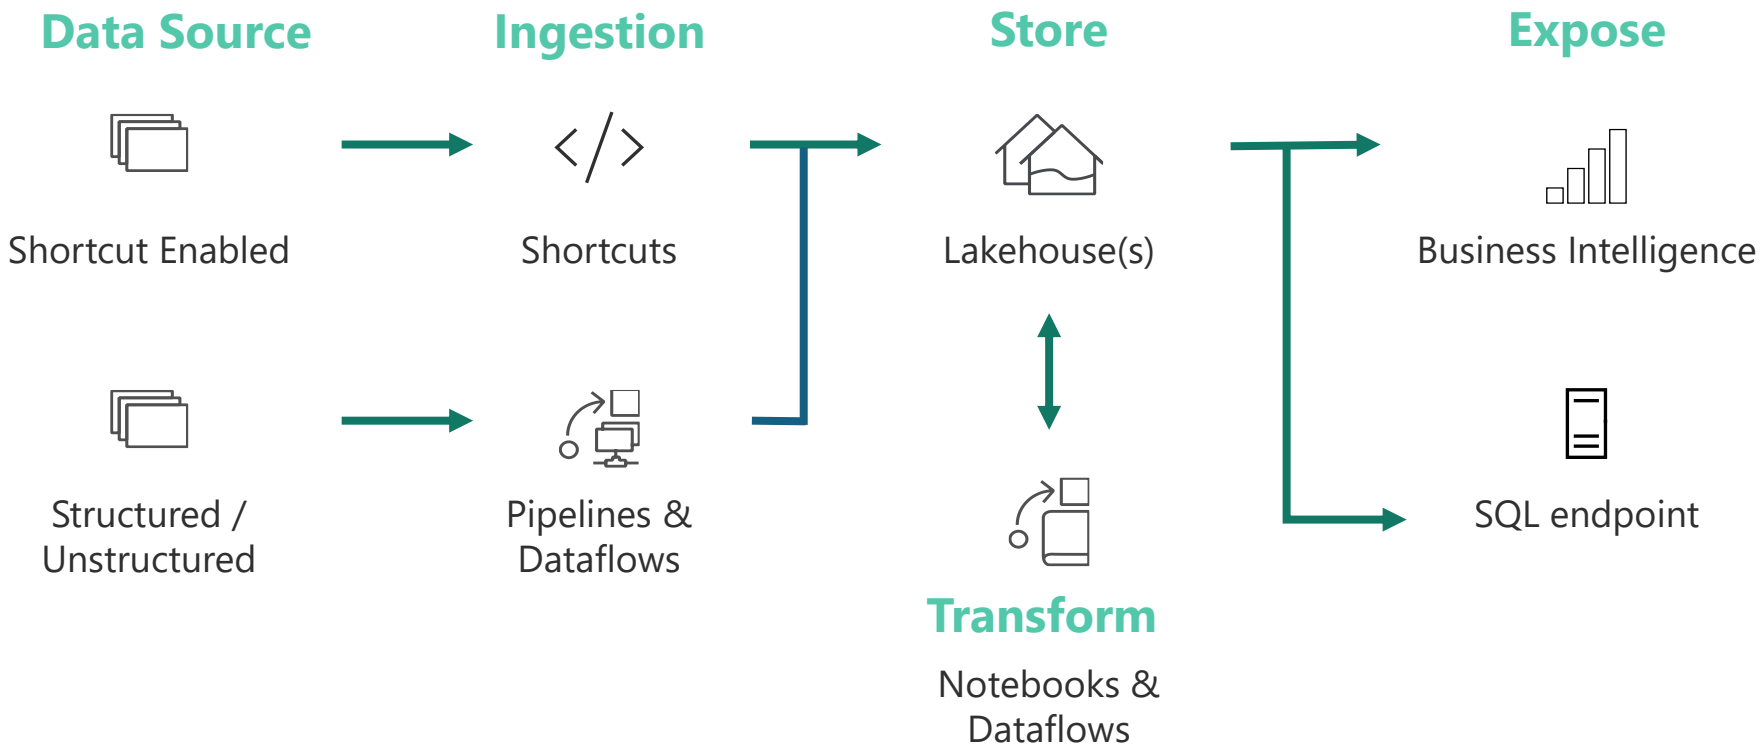

Data Warehouse scenario

End-to-end analytics scenario

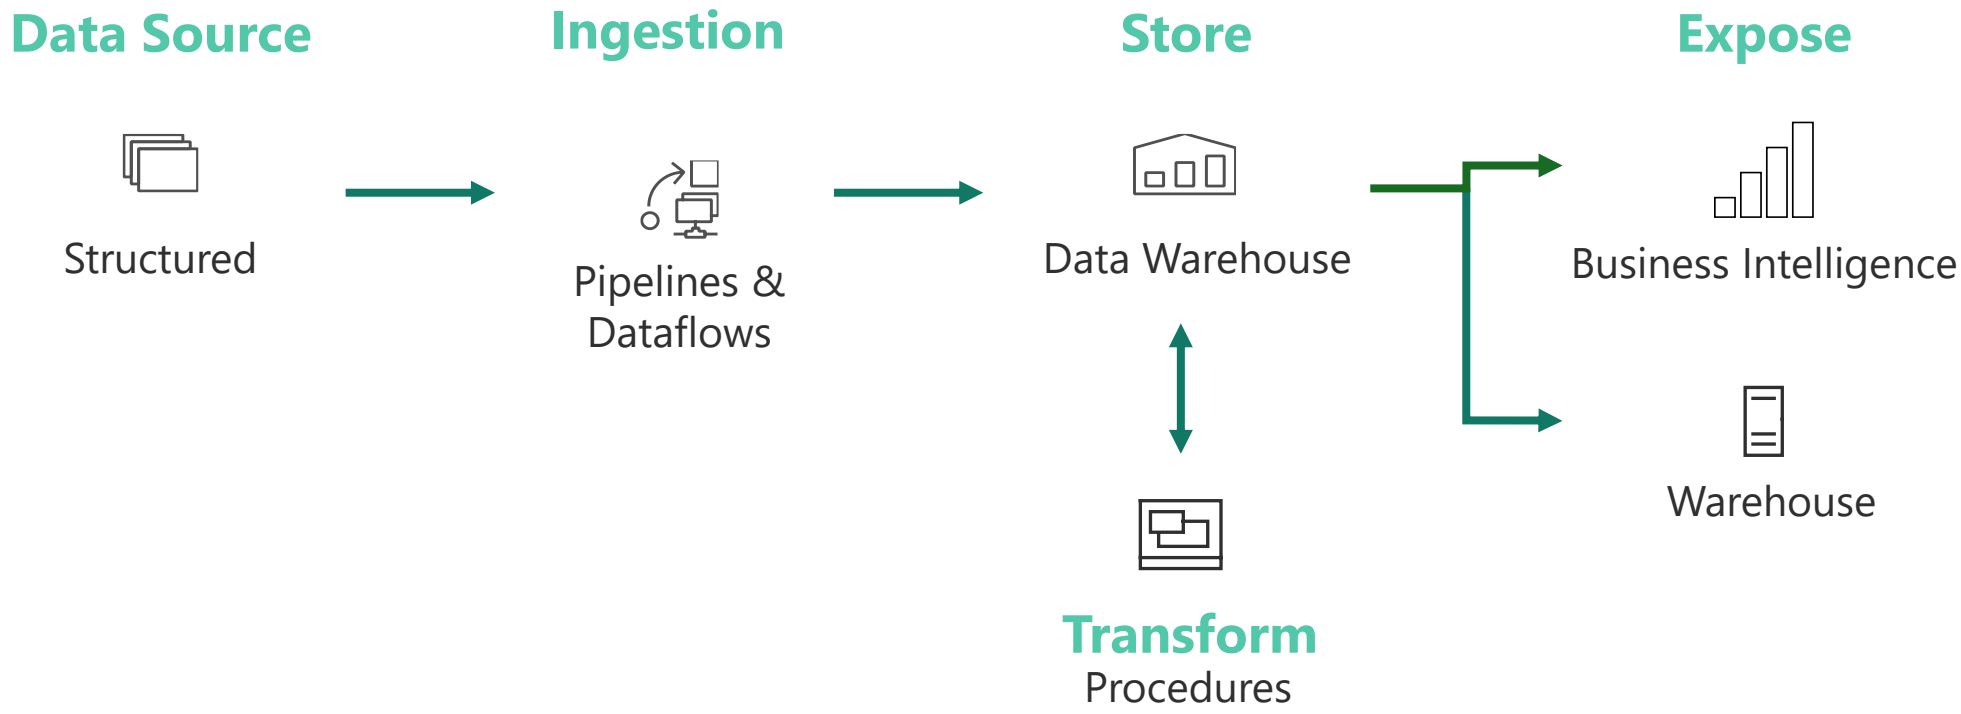

Data Science scenario

End-to-end analytics scenario

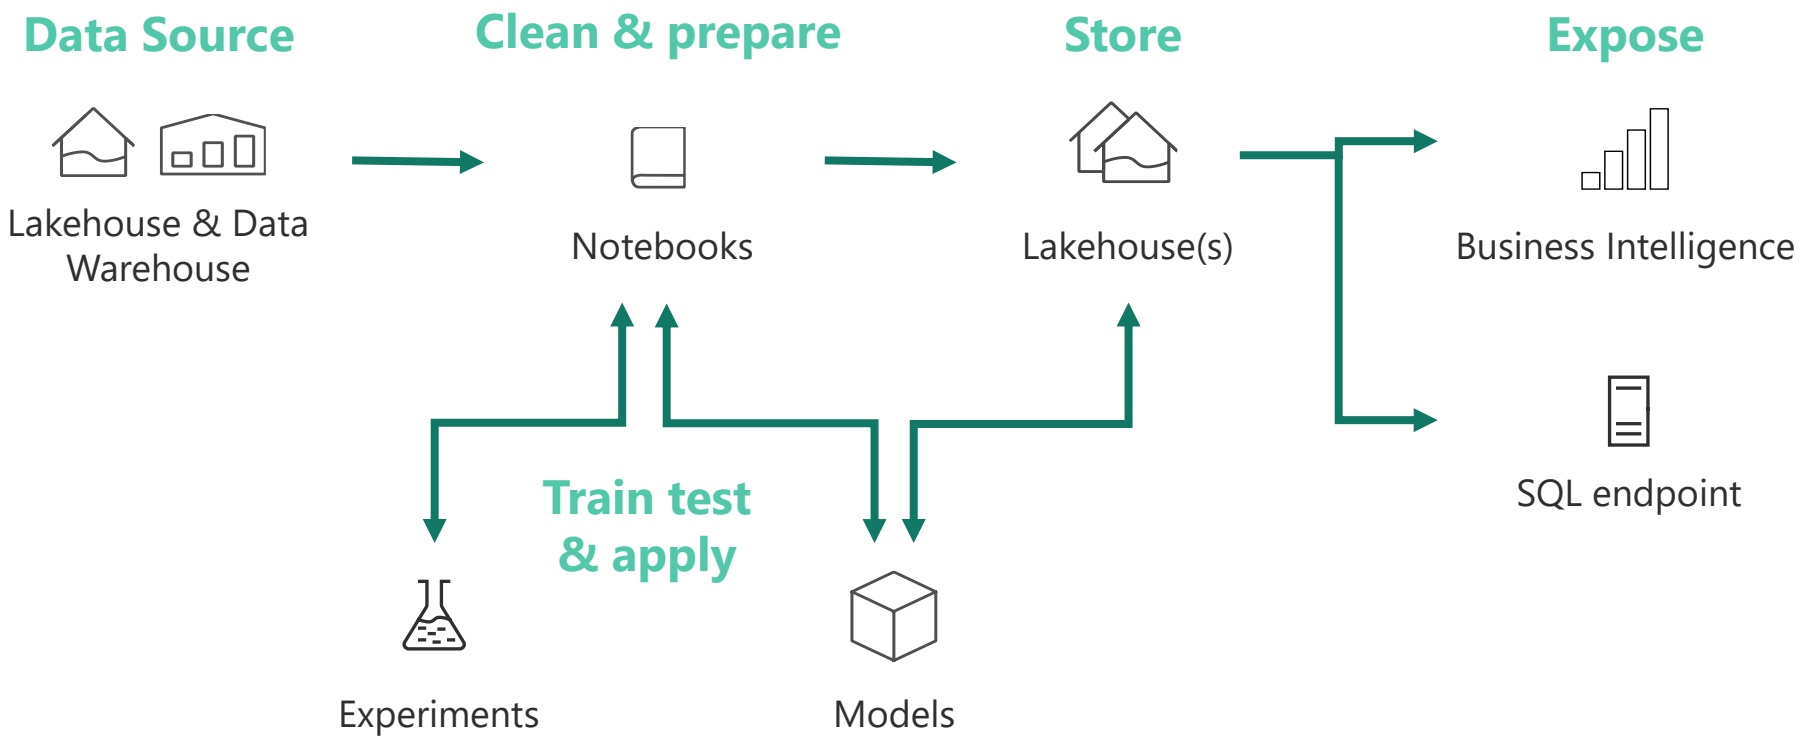

Real-time scenario

End-to-end analytics scenario

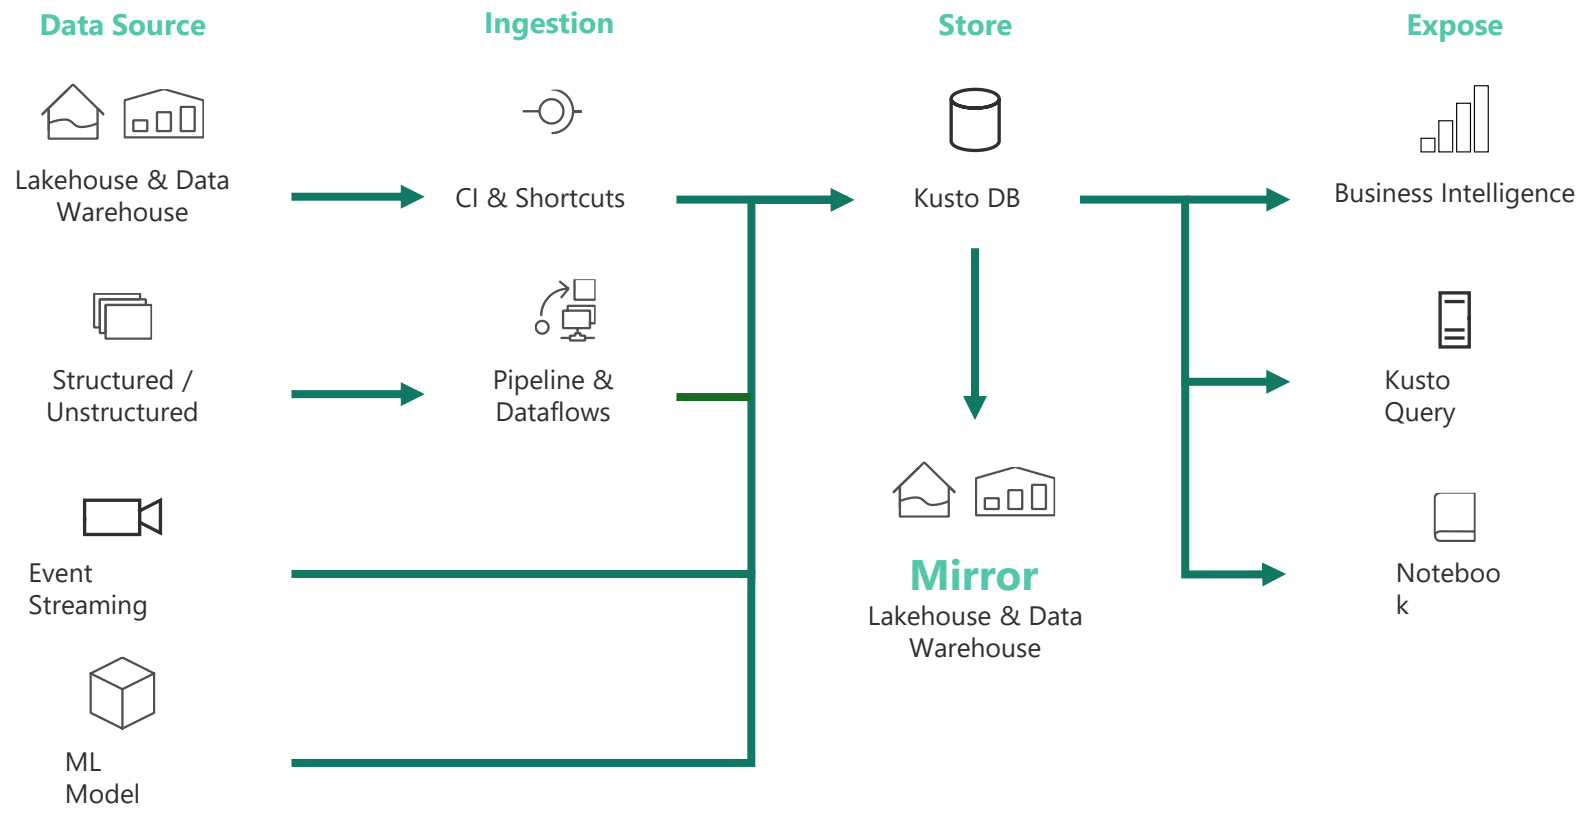

Spatial Analytics

Spatial Analytics

Esri announces integration with Microsoft to deliver leading-edge spatial analytics, allowing Microsoft Fabric customers valuable geospatial insights

Spatial Analytics

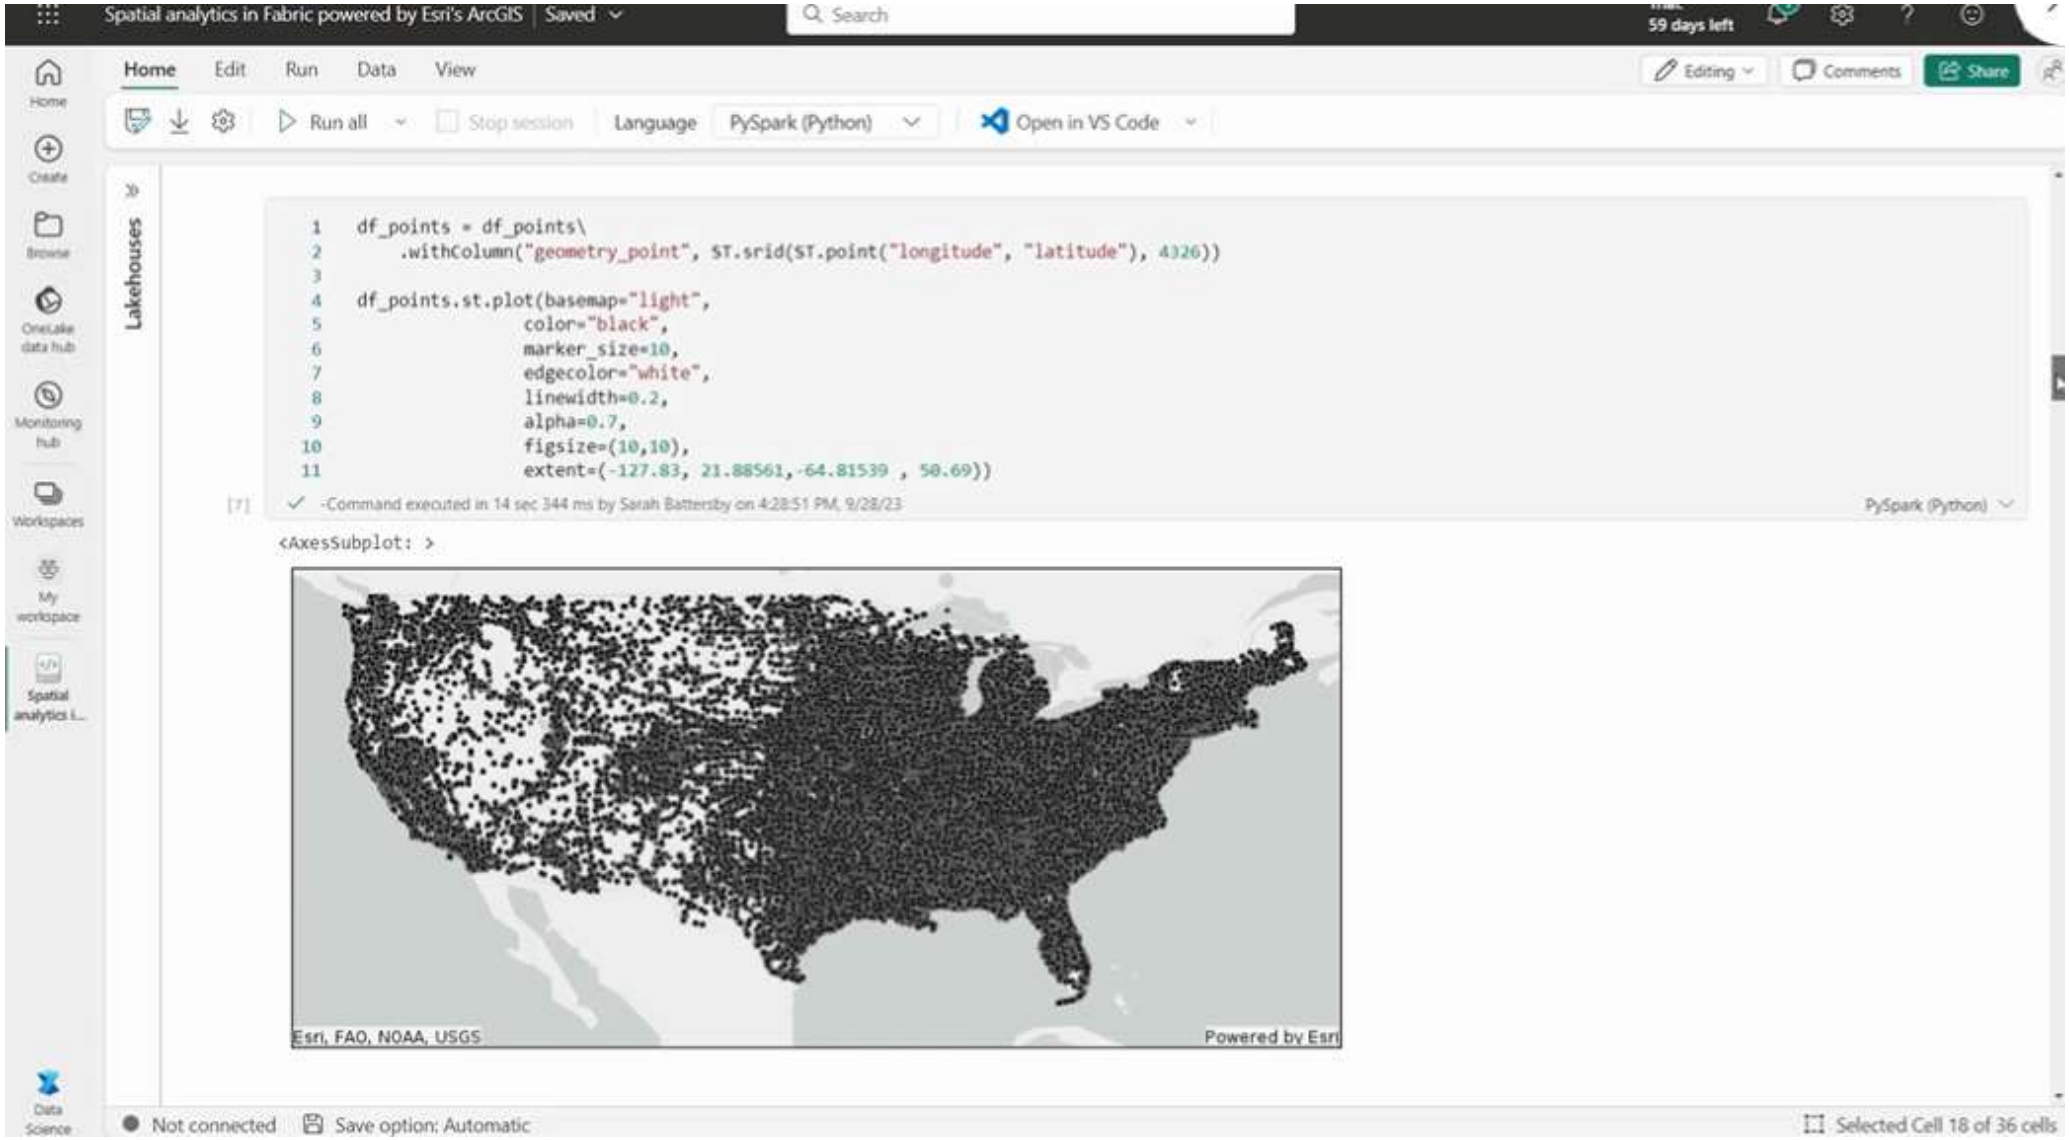

Tools

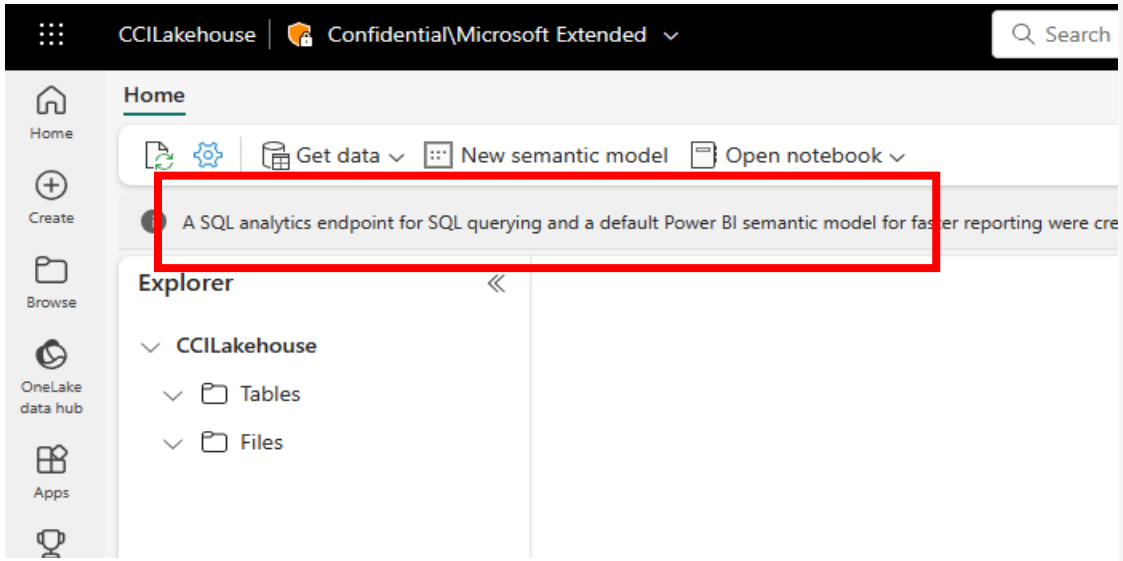
CCILakehouse
Confidential\Microsoft Extended
Search
Home
Get data
New semantic model
Open notebook
A SQL analytics endpoint for SQL querying and a default Power BI semantic model for faster reporting were cre
Explorer
CCILakehouse
Tables
Files

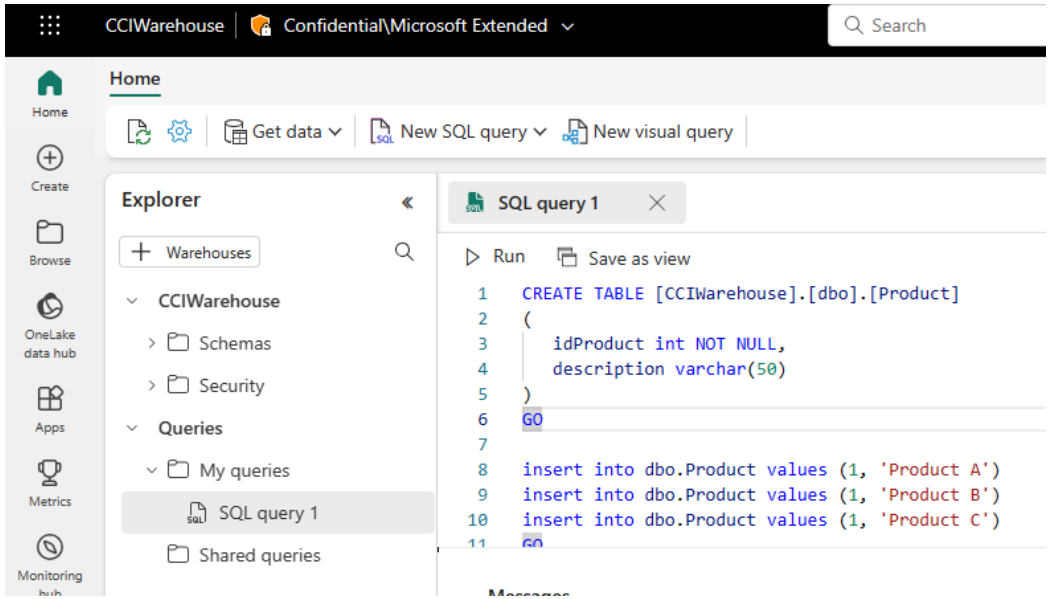
CCIWarehouse
Confidential\Microsoft Extended
Search
Home
Get data
New SQL query
New visual query
Explorer
Warehouses
CCIWarehouse
Schemas
Security
Queries
My queries
SQL query 1
Shared queries
SQL query 1
Run
Save as view
1 CREATE TABLE [CCIWarehouse].[dbo].[Product]
2 (
3 idProduct int NOT NULL,
4 description varchar(50)
5)
6 GO
7
8 insert into dbo.Product values (1, 'Product A')
9 insert into dbo.Product values (1, 'Product B')
10 insert into dbo.Product values (1, 'Product C')
11 GO

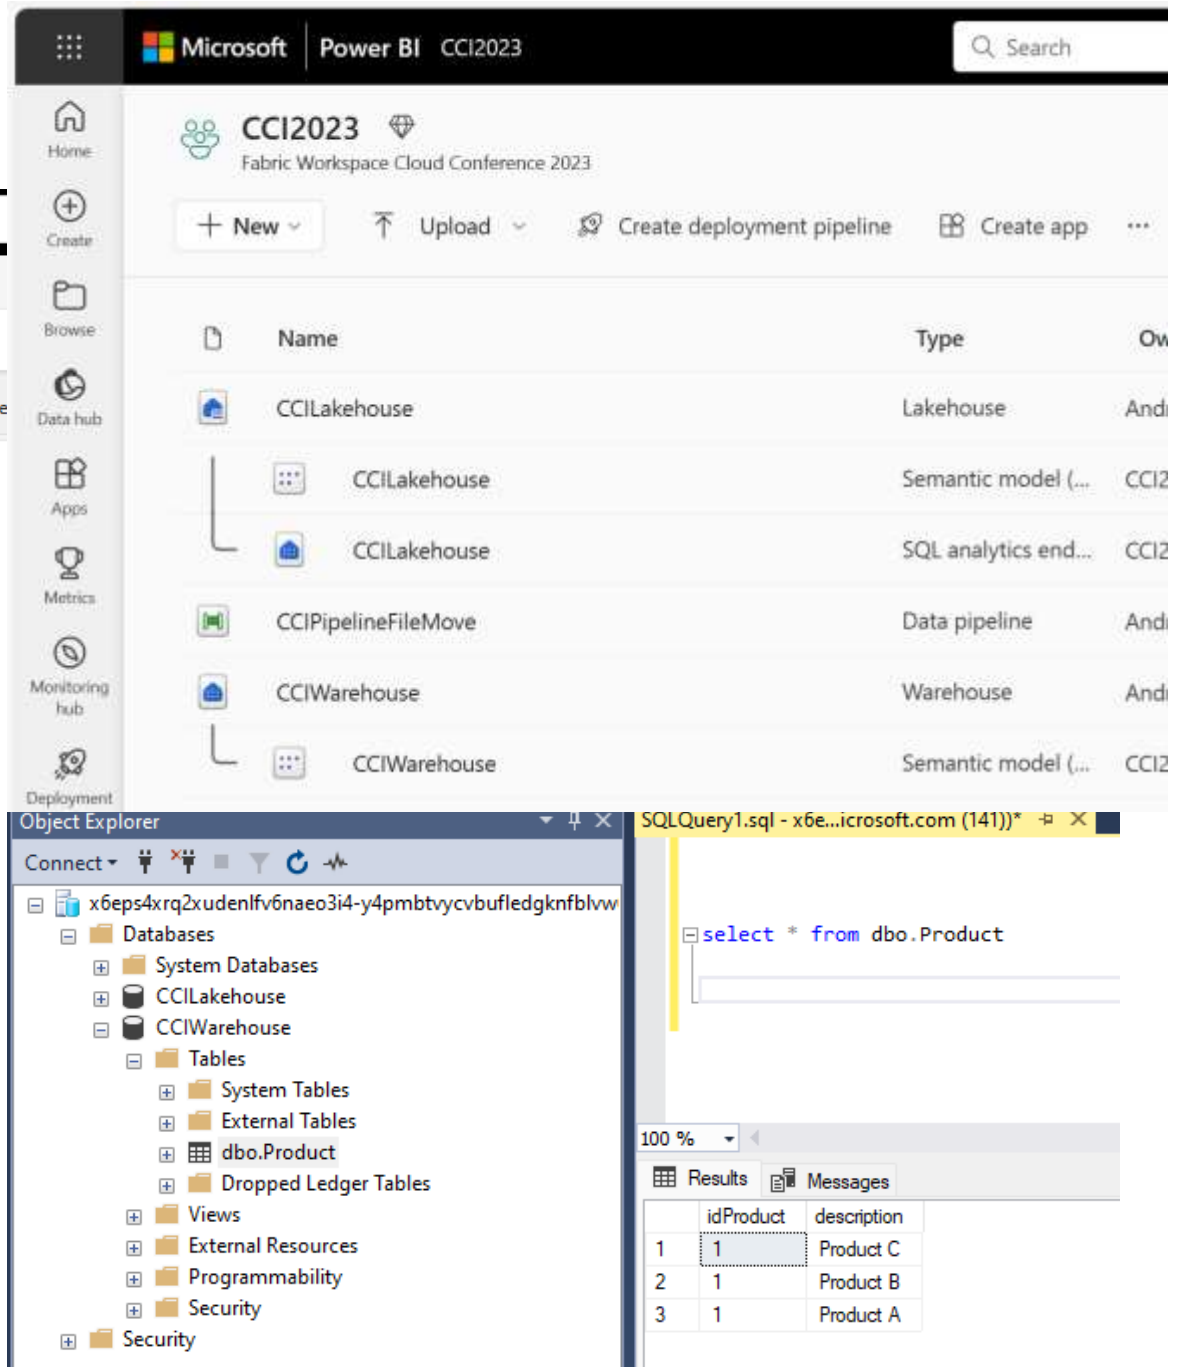
Microsoft
Power BI
CCI2023
Search
Home
Create
Browse
Data hub
Apps
Metrics
Monitoring hub
Deployment
CCI2023
Fabric Workspace Cloud Conference 2023
New
Upload
Create deployment pipeline
Create app
Name
Type
CCILakehouse
Lakehouse
CCILakehouse
Semantic model (...
CCILakehouse
SQL analytics end...
CCIPipelineFileMove
Data pipeline
CCIWarehouse
Warehouse
CCIWarehouse
Semantic model (...
Object Explorer
Connect
x6eps4xrq2xudenlfv6naeo3i4-y4pmbtvycvbufledgknfblvw
Databases
System Databases
CCILakehouse
CCIWarehouse
Tables
System Tables
External Tables
dbo.Product
Dropped Ledger Tables
Views
External Resources
Programmability
Security
Security
SQLQuery1.sql - x6e...icrosoft.com (141))*
select * from dbo.Product
100 %
Results
Messages
idProduct
description
1 1 Product C
2 1 Product B
3 1 Product A

Lakehouse / Warehouse

- Structured, unstructured, semi-structured data across all ranges of file formats
- Data Engineer (Spark, Scala)
- Data Scientist (Python)
- Bronze, Silver, Gold layer

- DirectQuery, Import, new DirectLake (!) connection from Power BI Semantic Models

- Structured data in tables
- Data Warehouse Developer (SQL)
- Gold layer (can be used)

- DirectQuery, Import, DirectLake connection from Power BI Semantic Models

DirectLake

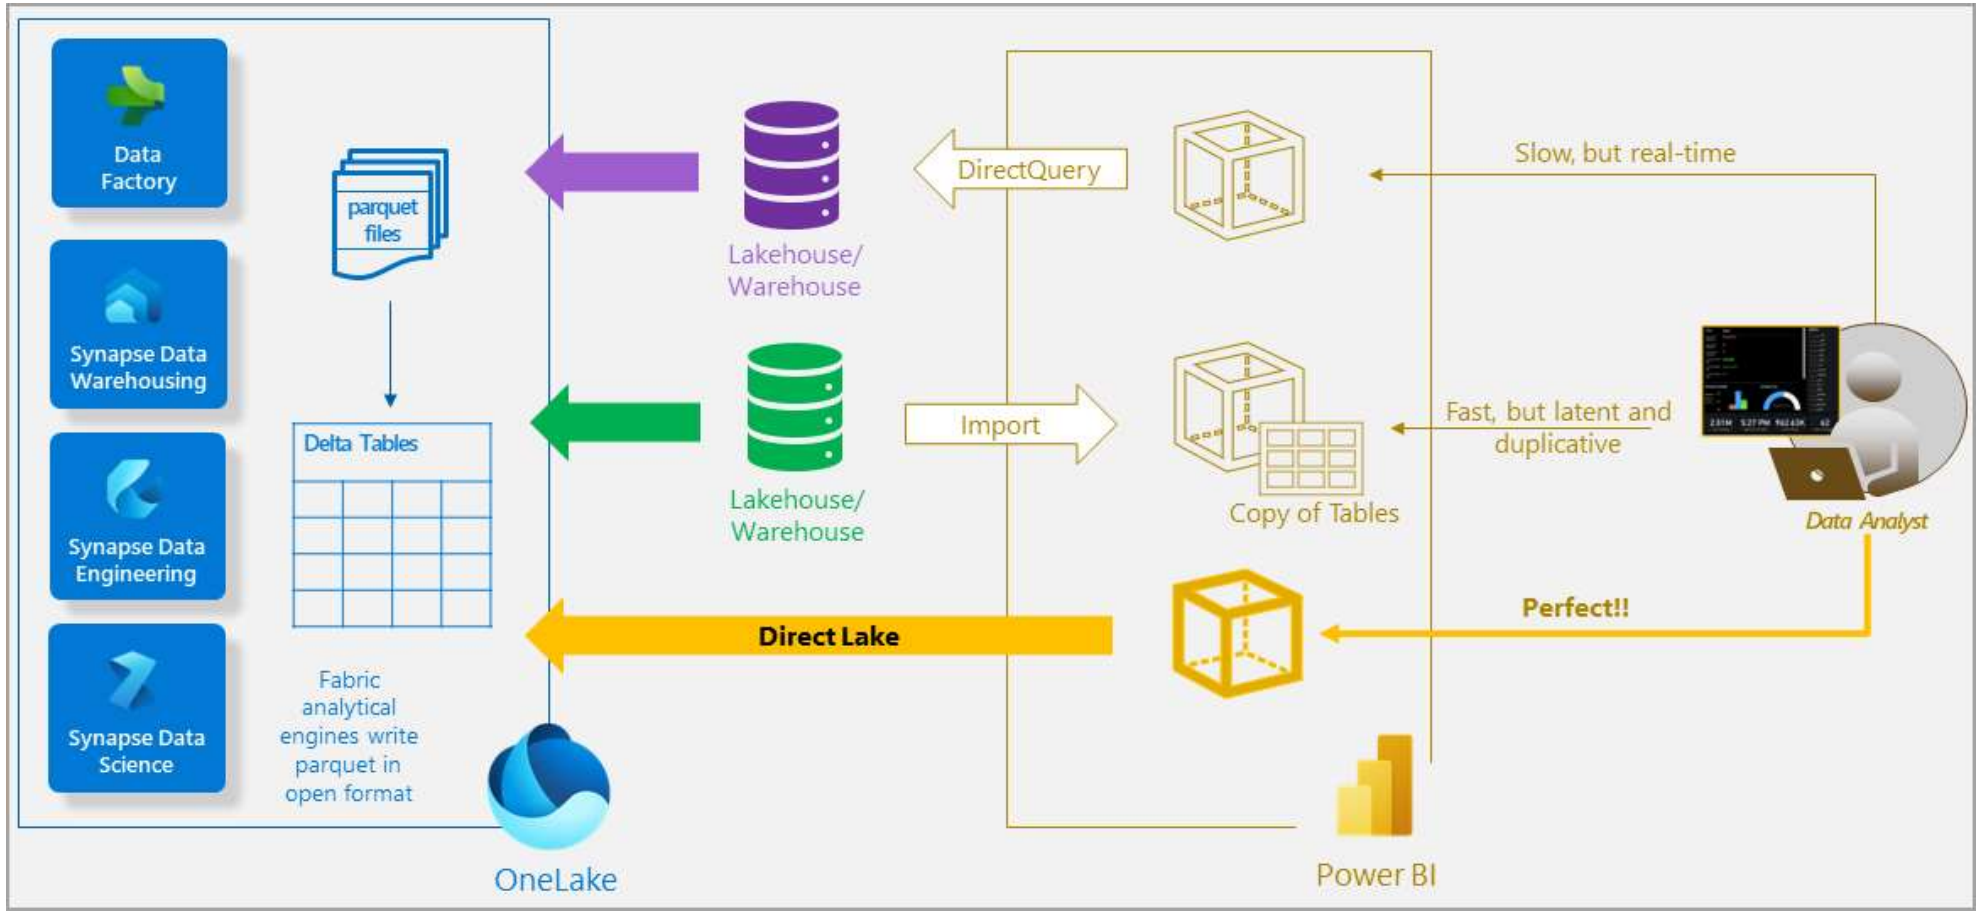
Data Factory
Synapse Data Warehousing
Synapse Data Engineering
Synapse Data Science
parquet files
Delta Tables
Fabric analytical engines write parquet in open format
OneLake
Lakehouse/ Warehouse
Lakehouse/ Warehouse
DirectQuery
Import
Direct Lake
Slow, but real-time
Fast, but latent and duplicative
Perfect!!
Copy of Tables
Data Analyst
Power BI

Lakehouse

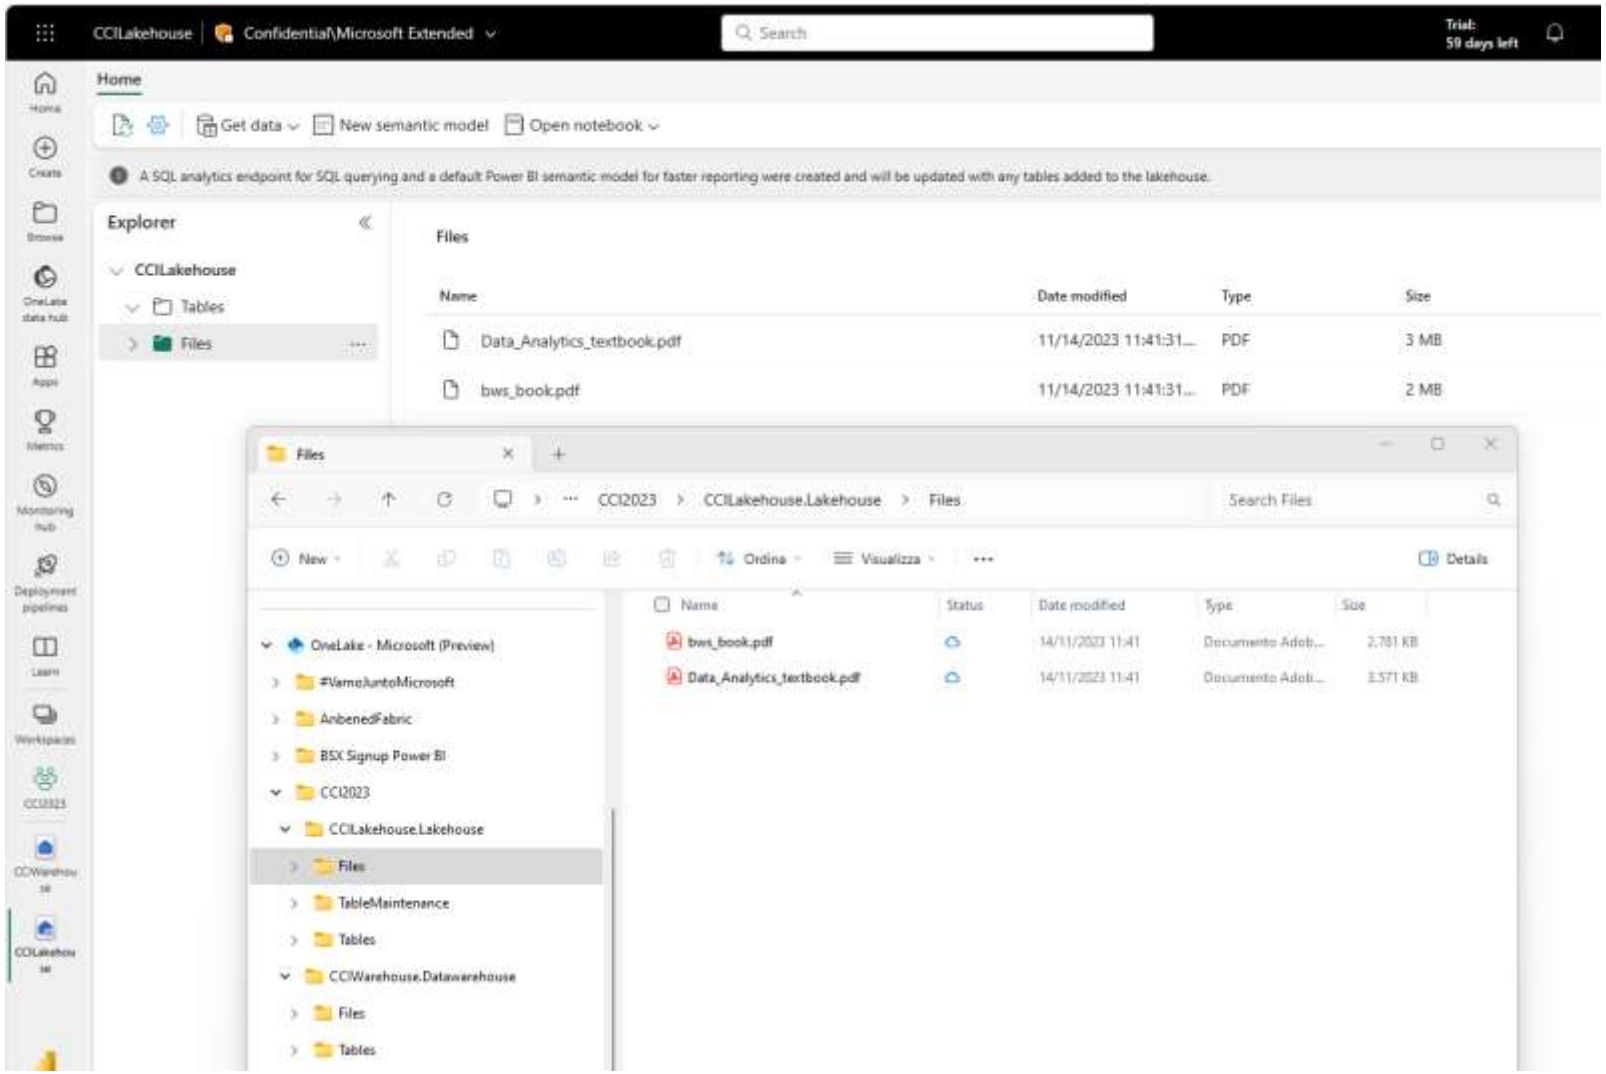
CCILakehouse
Confidential\Microsoft Extended
Search
Trial:
59 days left
Home
Get data
New semantic model
Open notebook
A SQL analytics endpoint for SQL querying and a default Power BI semantic model for faster reporting were created and will be updated with any tables added to the lakehouse.
Explorer
CCILakehouse
Tables
Files
Files
Name
Date modified
Type
Size
Data_Analytics_textbook.pdf
11/14/2023 11:41:31...
PDF
3 MB
bws_book.pdf
11/14/2023 11:41:31...
PDF
2 MB
Files
CCI2023
CCILakehouse.Lakehouse
Files
Search Files
New
Ordina
Visualizza
Details
OneLake - Microsoft (Preview)
#VamoJuntoMicrosoft
AnbenedFabric
BSX Signup Power BI
CCI2023
CCILakehouse.Lakehouse
Files
TableMaintenance
Tables
CCIWarehouse.Datawarehouse
Files
Tables
Name
Status
Date modified
Type
Size
bws_book.pdf
14/11/2023 11:41
Documento Adob...
2.781 KB
Data_Analytics_textbook.pdf
14/11/2023 11:41
Documento Adob...
3.571 KB

Lakehouse

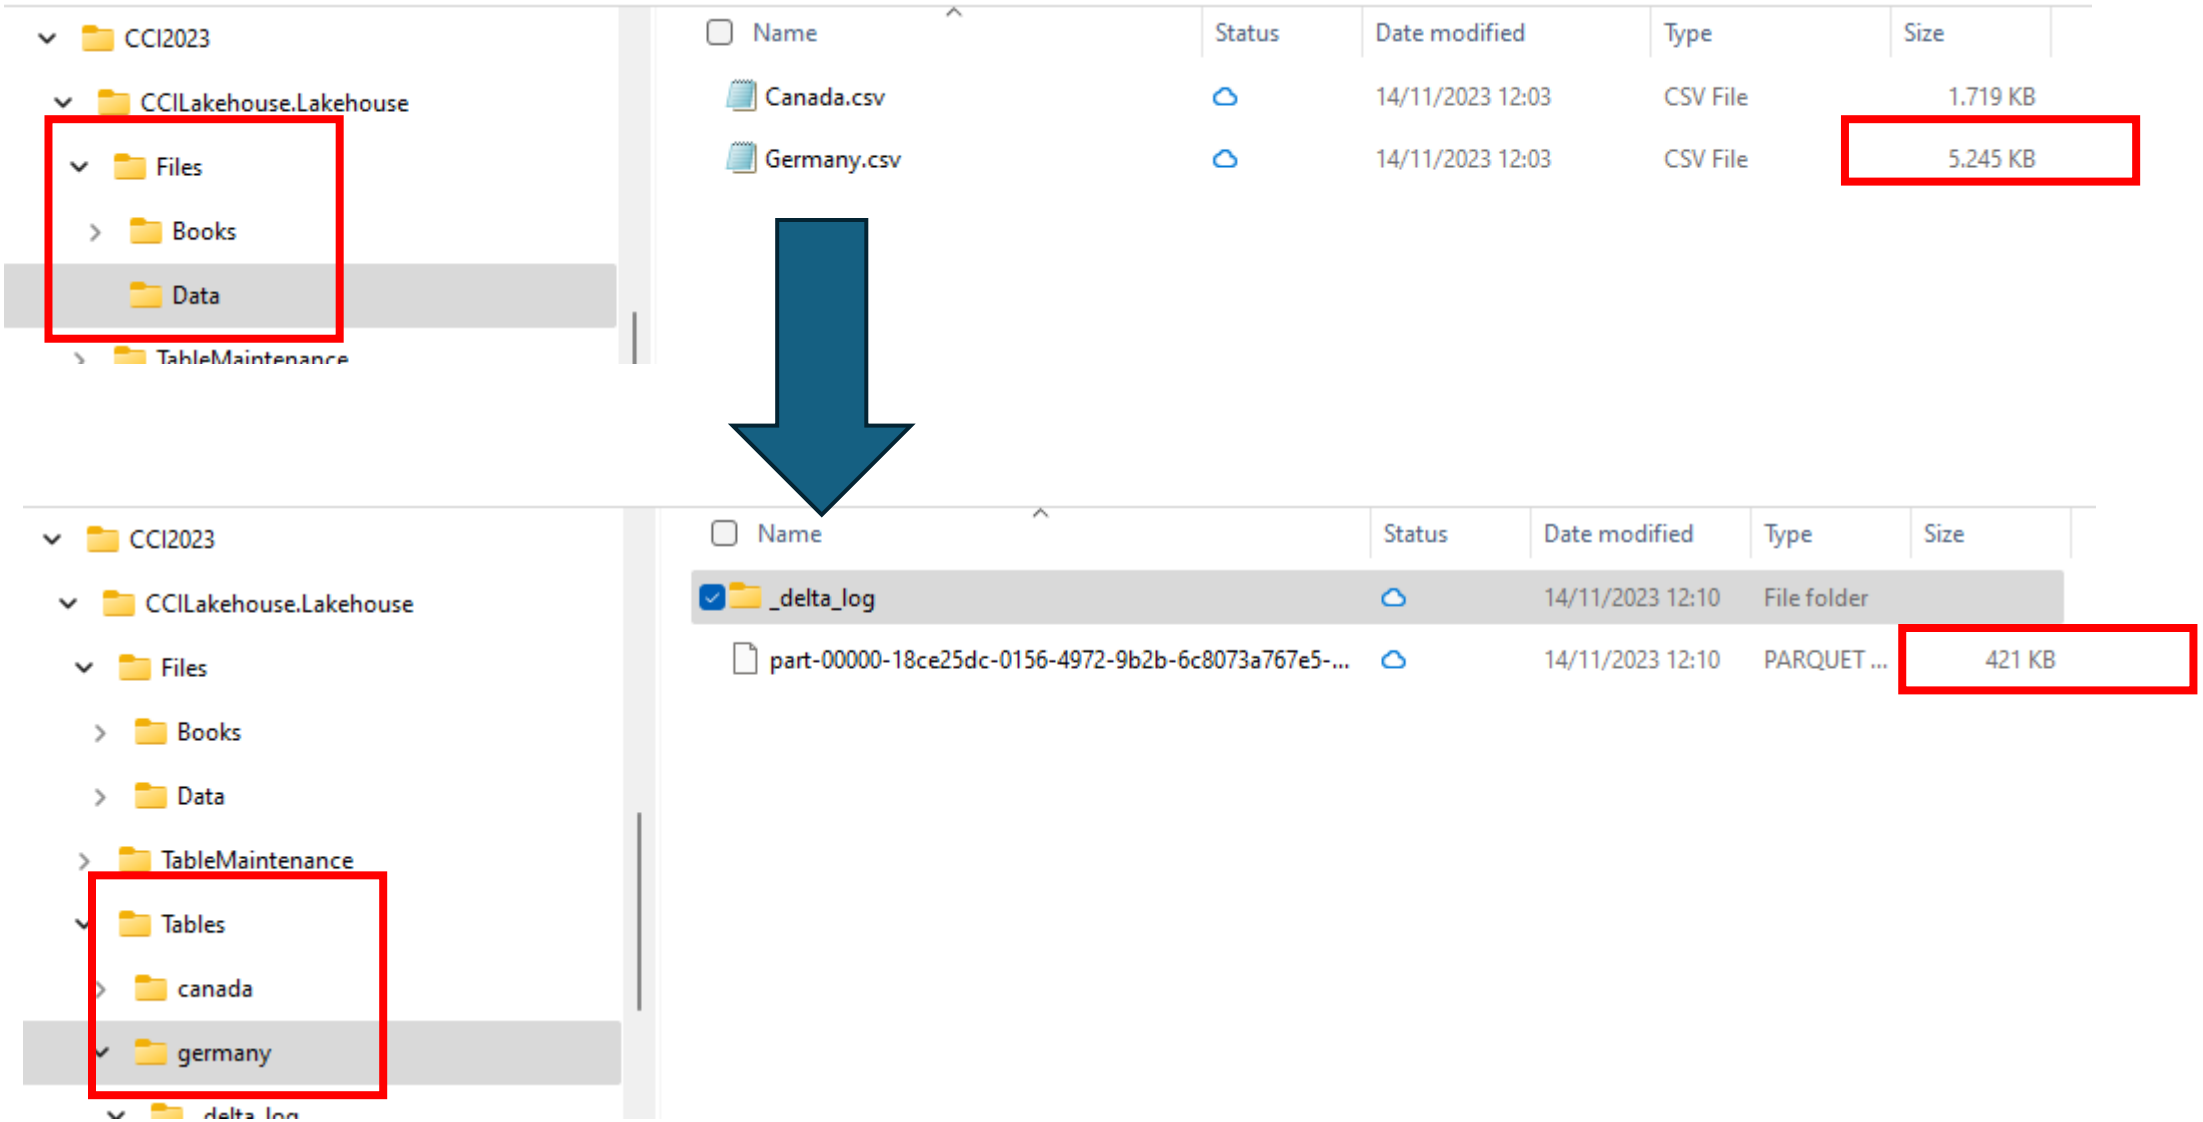
CCI2023
CCILakehouse.Lakehouse
Files
Books
Data
TableMaintenance
Name
Status
Date modified
Type
Size
Canada.csv
14/11/2023 12:03
CSV File
1.719 KB
Germany.csv
14/11/2023 12:03
CSV File
5.245 KB
CCI2023
CCILakehouse.Lakehouse
Files
Books
Data
TableMaintenance
Tables
canada
germany
delta_log
Name
Status
Date modified
Type
Size
_delta_log
14/11/2023 12:10
File folder
part-00000-18ce25dc-0156-4972-9b2b-6c8073a767e5-...
14/11/2023 12:10
PARQUET ...
421 KB

Lakehouse

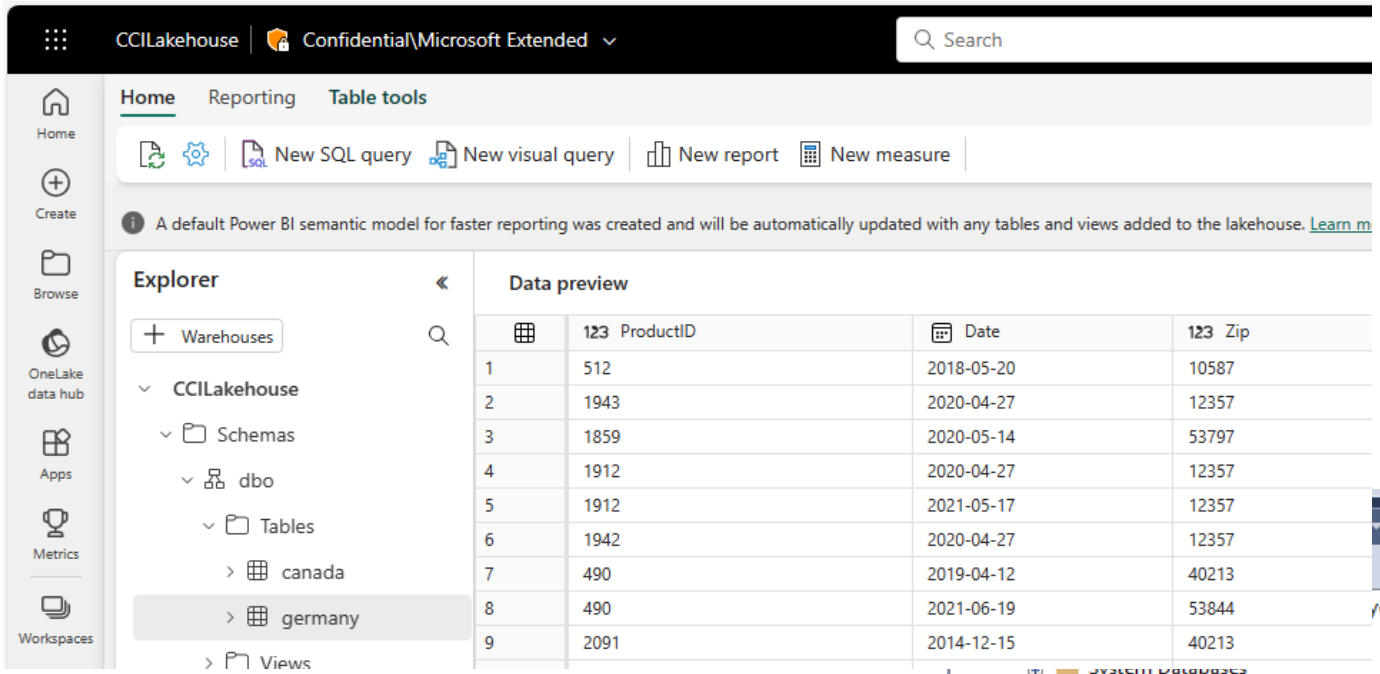
CCILakehouse
Confidential\Microsoft Extended
Search
Home
Reporting
Table tools
New SQL query
New visual query
New report
New measure
A default Power BI semantic model for faster reporting was created and will be automatically updated with any tables and views added to the lakehouse.
Explorer
Warehouses
CCILakehouse
Schemas
dbo
Tables
canada
germany
Views
Data preview
ProductID
Date
Zip
1 512 2018-05-20 10587
2 1943 2020-04-27 12357
3 1859 2020-05-14 53797
4 1912 2020-04-27 12357
5 1912 2021-05-17 12357
6 1942 2020-04-27 12357
7 490 2019-04-12 40213
8 490 2021-06-19 53844
9 2091 2014-12-15 40213

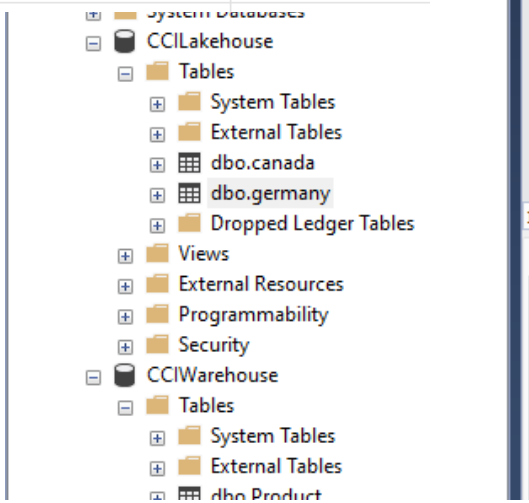
CCILakehouse
Tables
System Tables
External Tables
dbo.canada
dbo.germany
Dropped Ledger Tables
Views
External Resources
Programmability
Security
CCIWarehouse
Tables
System Tables
External Tables

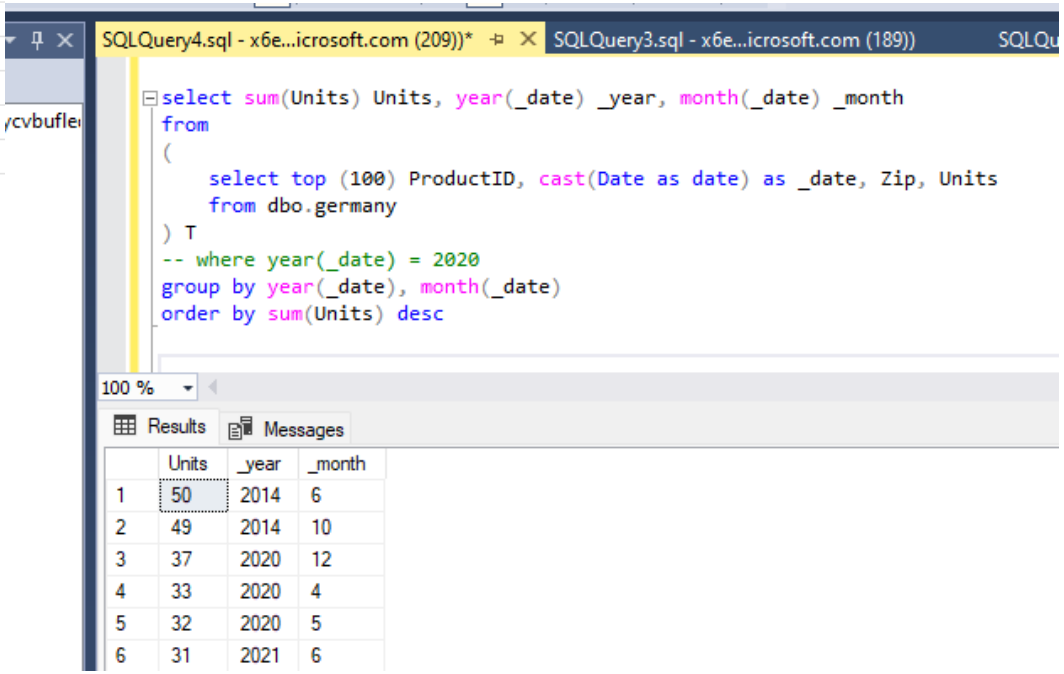
SQLQuery4.sql - x6e...icrosoft.com (209))*
SQLQuery3.sql - x6e...icrosoft.com (189))
select sum(Units) Units, year(_date) _year, month(_date) _month
from
(
select top (100) ProductID, cast(Date as date) as _date, Zip, Units
from dbo.germany
) T
-- where year(_date) = 2020
group by year(_date), month(_date)
order by sum(Units) desc
100 %
Results
Messages
Units _year _month
1 50 2014 6
2 49 2014 10
3 37 2020 12
4 33 2020 4
5 32 2020 5
6 31 2021 6

Warehouse

- Auto-scale out
- Fault tolerant
- Distributed compute engine

- Enhanced version of Serverless SQL Pool engine

- Native storage uses Delta Lake
  - Transaction Log
  - Metadata
  - Parquet file format

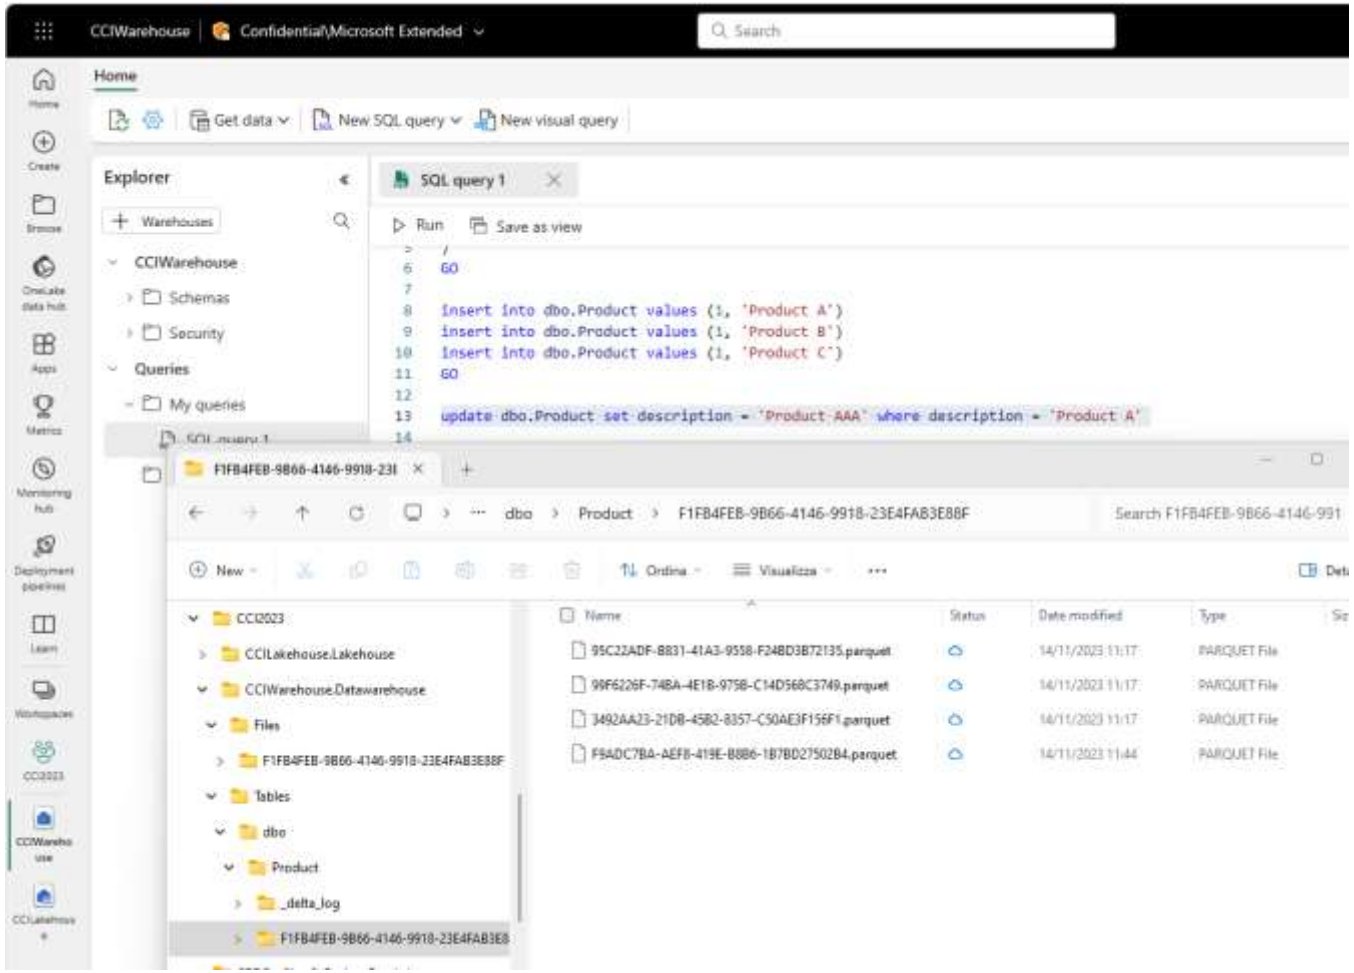

Loading

- SQL COPY INTO → from Azure Storage
- SQL CREATE TABLE AS … SELECT
- SQL INSERT INTO… VALUES / SELECT
- Pyodbc → from Notebooks

- Data Factory Data Pipeline

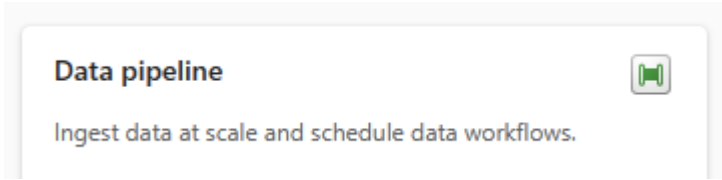

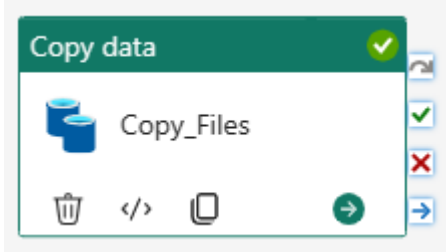

- Dataflows Gen2 (Power Query)

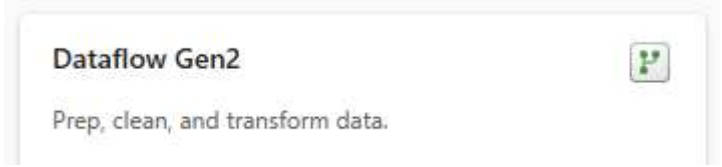

DEMO

If you are an existing Power BI Premium per capacity customer, you can already access Microsoft Fabric by simply turning on Fabric in your admin portal.

Economics

Microsoft Fabric Reserved Instances are here!

F2 reserved capacity that comes with all the workloads, all running at blazing speed, costs just $156/month.

Region: Central US | Currency: United States – Dollar ($) USD | Display pricing by: Month

SKU	Capacity unit (CU)	Pay-as-you-go	Reservation
F 2	2	$262.80/month	$156.334/month ~41% savings
F 4	4	$525.60/month	$312.667/month ~41% savings
F 8	8	$1,051.20/month	$625.334/month ~41% savings
F 16	16	$2,102.40/month	$1,250.667/month ~41% savings
F 32	32	$4,204.80/month	$2,501.334/month ~41% savings
F 64	64	$8,409.60/month	$5,002.667/month ~41% savings
F 128	128	$16,819.20/month	$10,005.334/month ~41% savings
F 256	256	$33,638.40/month	$20,010.667/month ~41% savings
F 512	512	$67,276.80/month	$40,021.334/month ~41% savings

Reservations

Reservation pricing → pre-commit to Fabric Capacity Units in one-year increments

- save up to 40.5 percent over the pay-as-you-go prices (excluding Power BI Capacity SKUs)
- You can pay for all your consumption, across every Fabric workload, with a single, unified bill and use the Capacity Metrics app to understand and track your usage

- You can choose to pay monthly or upfront
- You can increase the size of Fabric capacity reservations by purchasing additional reservations or by exchanging them for new ones
- You can exchange existing Azure Synapse Analytics Dedicated SQL pool (formerly SQL DW) reservations to Fabric capacity reservations
- Reservation for Fabric capacity doesn't cover storage or networking charges associated with the Microsoft Fabric usage; it only covers Fabric capacity usage

Notes

Regulatory compliance

Data residency

- Fabric will be available in every Azure region
- Data at rest: compliant with EUDB and other single-geo data residency regulations
- Multi-geo capacities allow control over content storage location in most Azure data centers world-wide

Microsoft pledges support for EU Data Boundary

Choose a data store

Data warehouse and lakehouse properties

	Data warehouse	Lakehouse	Power BI Datamart	KQL Database
Data volume	Unlimited	Unlimited	Up to 100 GB	Unlimited
Type of data	Structured	Unstructured,semi-structured,structured	Structured	Unstructured, semi-structured, structured
Primary developer persona	Data warehouse developer, SQL engineer	Data engineer, data scientist	Citizen developer	Citizen Data scientist, Data engineer, Data scientist, SQL engineer
Primary developer skill set	SQL	Spark(Scala, PySpark, Spark SQL, R)	No code, SQL	No code, KQL, SQL
Data organized by	Databases, schemas, and tables	Folders and files, databases, and tables	Database, tables, queries	Databases, schemas, and tables
Read operations	Spark,T-SQL	Spark,T-SQL	Spark,T-SQL,Power BI	KQL, T-SQL, Spark, Power BI
Write	T-SQL	Spark(Scala, PySpark, Spark SQL,	Dataflows,	KQL, Spark, connector

https://learn.microsoft.com/en-us/fabric/get-started/decision-guide-data-store

Copy activity, dataflow, or Spark

Copy activity, dataflow, and Spark properties

	Pipeline copy activity	Dataflow Gen 2	Spark
Use case	Data lake and data warehouse migration, data ingestion, lightweight transformation	Data ingestion, data transformation, data wrangling, data profiling	Data ingestion, data transformation, data processing, data profiling
Primary developer persona	Data engineer, data integrator	Data engineer, data integrator, business analyst	Data engineer, data scientist, data developer
Primary developer skill set	ETL, SQL, JSON	ETL, M, SQL	Spark (Scala, Python, Spark SQL, R)
Code written	No code, low code	No code, low code	Code
Data volume	Low to high	Low to high	Low to high
Development interface	Wizard, canvas	Power query	Notebook, Spark job definition
Sources	30+ connectors	150+ connectors	Hundreds of Spark libraries

https://learn.microsoft.com/en-us/fabric/get-started/decision-guide-pipeline-dataflow-spark

Fabric Community

https://ideas.fabric.microsoft.com/

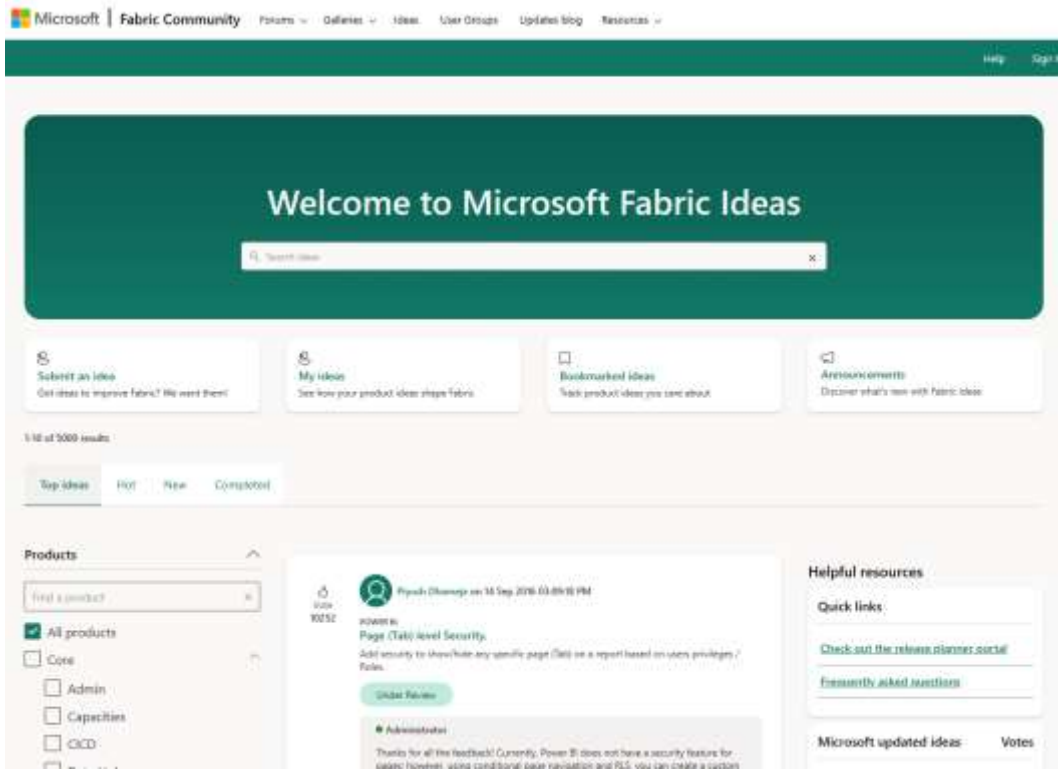

Learn

An official **Microsoft certification** for implementing analytics solutions with Microsoft Fabric and Power BI

Exam DP-600: Implementing Analytics Solutions Using Microsoft Fabric (beta) - Certifications | Microsoft Learn

Microsoft Fabric Community Resources

• Community Call to Action

- ✓ Try Microsoft Fabric for free: https://aka.ms/try-fabric
- ✓ Join the Fabric community: https://aka.ms/fabriccommunity
- ✓ Share and vote for ideas to improve Fabric: https://aka.ms/fabricideas
- ✓ Read and comment our blog: https://aka.ms/fabricblog

Learn More about Microsoft Fabric

- Product announcement: https://aka.ms/fabric
- Digital Event at Build (videos): https://aka.ms/build-with-analytics
- Product website: https://aka.ms/microsoft-fabric
- Documentation: https://aka.ms/fabric-docs
- Fabric e-book: https://aka.ms/fabric-get-started-ebook
- Microsoft Learn: https://aka.ms/learn-fabric
- End-to-end scenario tutorials: https://aka.ms/fabric-tutorials
- Fabric Notes: https://aka.ms/fabric-notes

End-to-end tutorials

- Lakehouse tutorial
  - https://learn.microsoft.com/en-us/fabric/data-engineering/tutorial-lakehouse-introduction
- Data Science tutorial
  - https://learn.microsoft.com/en-us/fabric/data-science/tutorial-data-science-introduction
- Real-Time Analytics tutorial
  - https://learn.microsoft.com/en-us/fabric/real-time-analytics/tutorial-introduction
- Data warehouse tutorial
  - https://learn.microsoft.com/en-us/fabric/data-warehouse/tutorial-introduction
- Power BI tutorial
  - https://learn.microsoft.com/en-us/power-bi/fundamentals/fabric-get-started
- Data Factory tutorial
  - https://learn.microsoft.com/en-us/fabric/data-factory/tutorial-end-to-end-introduction

GRAZIE!

SPEAKER

Andrea Benedetti

Sr Cloud Architect, Microsoft

/in/abenedetti

@anBenedetti

https://github.com/anbened