

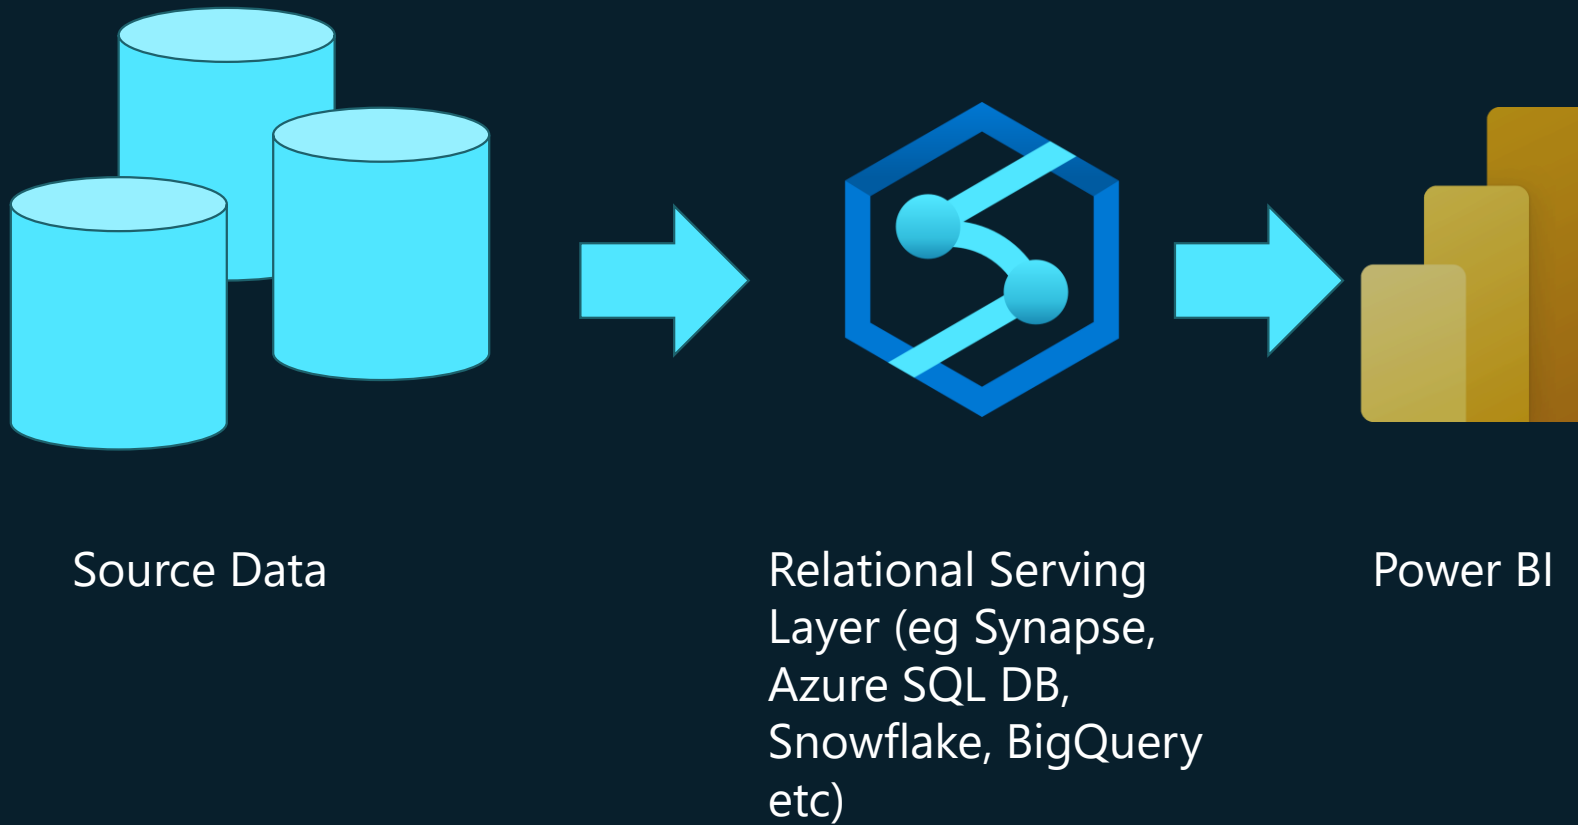


Direct Lake

Andrea Benedetti

**What Is
Direct Lake?**

Traditional Microsoft BI Architecture



Storage Modes (today)

| SMALLER MODELS | Time to Import Data | Model Size | Query Speed |
|----------------|---------------------|------------|-------------|
| Direct Query | - | - | ? |
| Import | 😊 | 😊 | 😊 |

| LARGE MODELS | Time to Import Data | Model Size | Query Speed |
|--------------|---------------------|------------|-------------|
| Direct Query | - | - | ? |
| Import | 😐 | 😐 | 😊 |

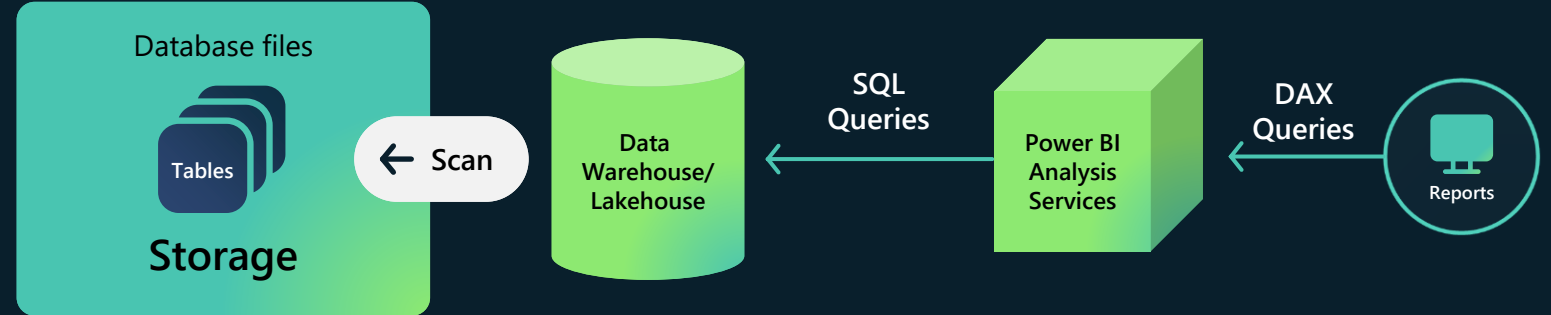
Storage Modes

| SMALLER MODELS | Time to Import Data | Model Size | Query Speed |
|----------------|---------------------|------------|-------------|
| Direct Query | - | - | ? |
| Import | 😊 | 😊 | 😊 |
| Direct Lake | 😊 | 😊 | 😊 |

| LARGE MODELS | Time to Import Data | Model Size | Query Speed |
|--------------|---------------------|------------|-------------|
| Direct Query | - | - | ? |
| Import | 😐 | 😐 | 😊 |
| Direct Lake | 😊 | 😊 | 😊 |

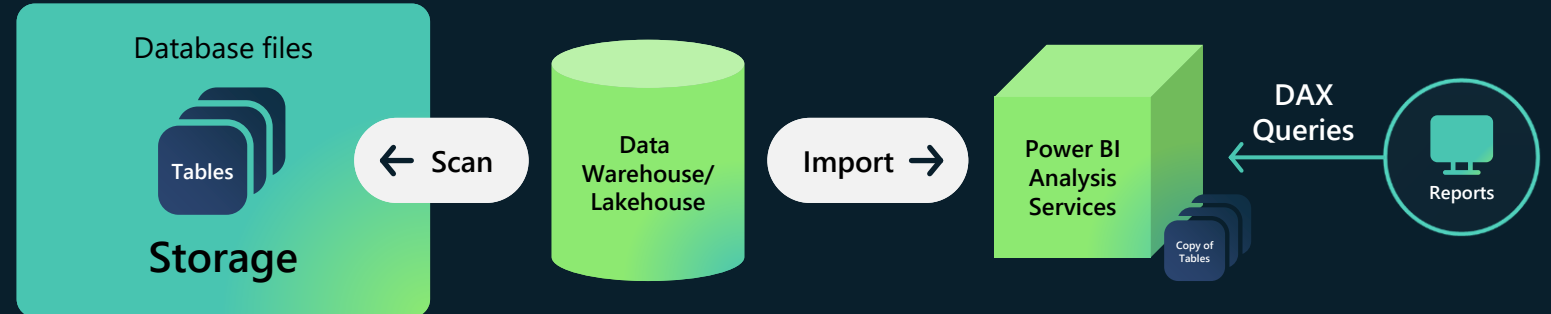
"Direct Query Mode"

Slow, but real time



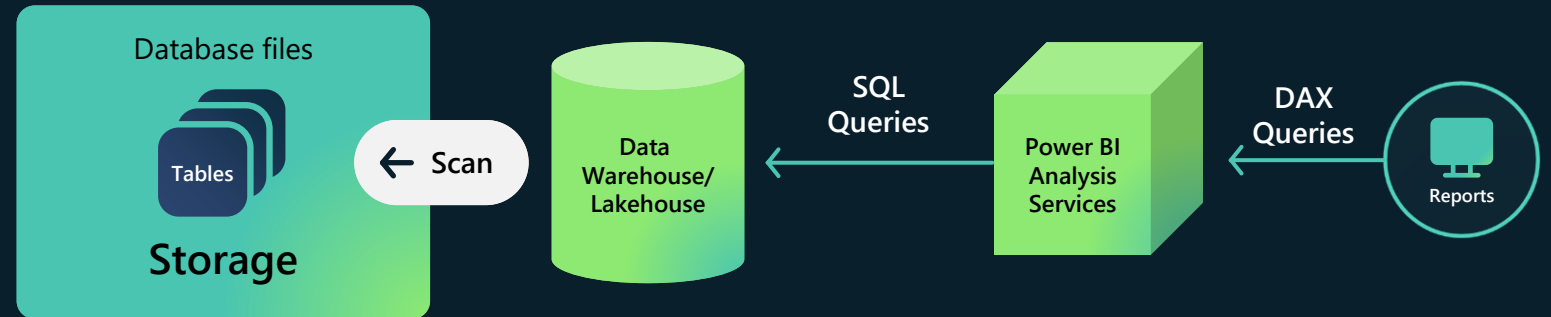
"Import Mode"

Latent & duplicative but fast



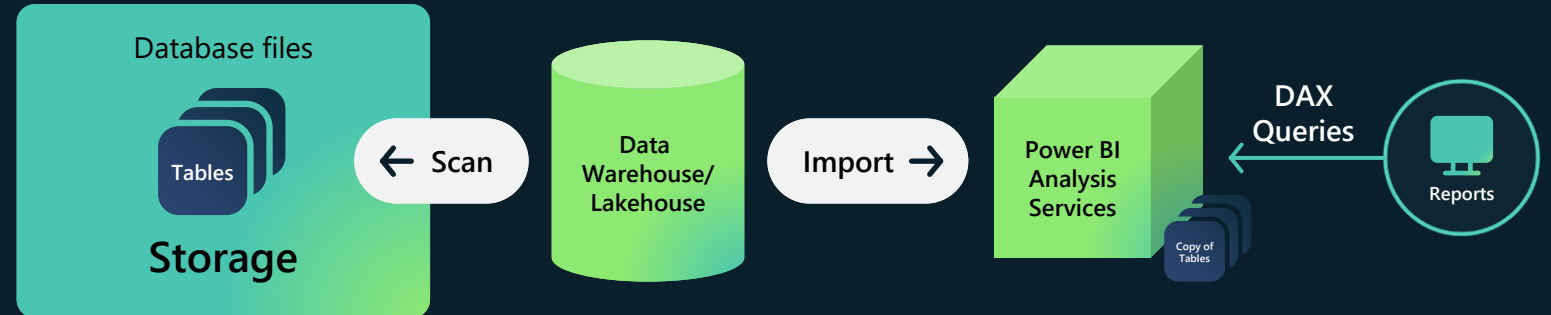
"Direct Query Mode"

Slow, but real time




"Import Mode"

Latent & duplicative but fast

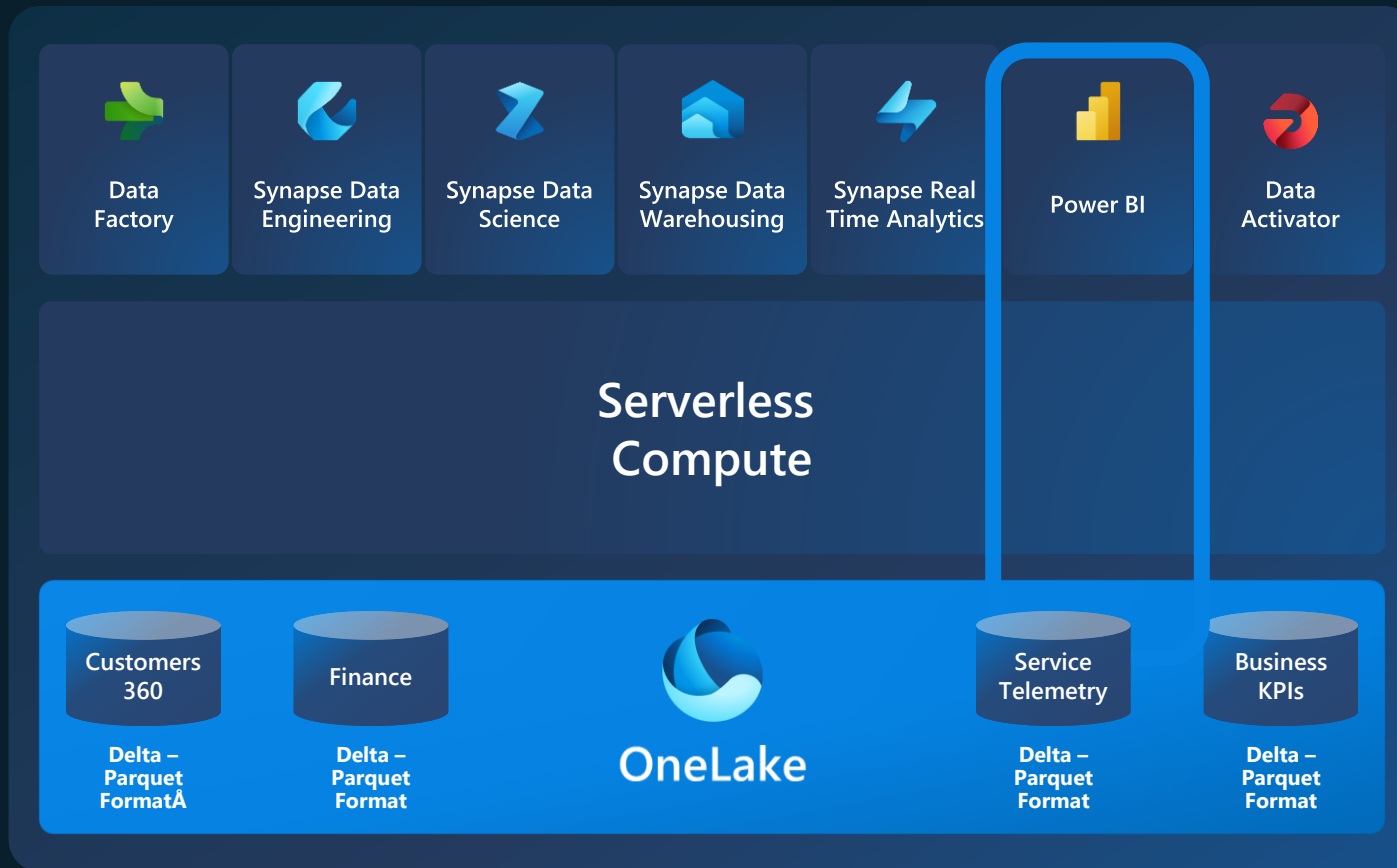


"Direct Lake Mode"

Perfect!



One Copy – Direct Lake



All the compute engines store their data automatically in OneLake

The data is stored in a single common format

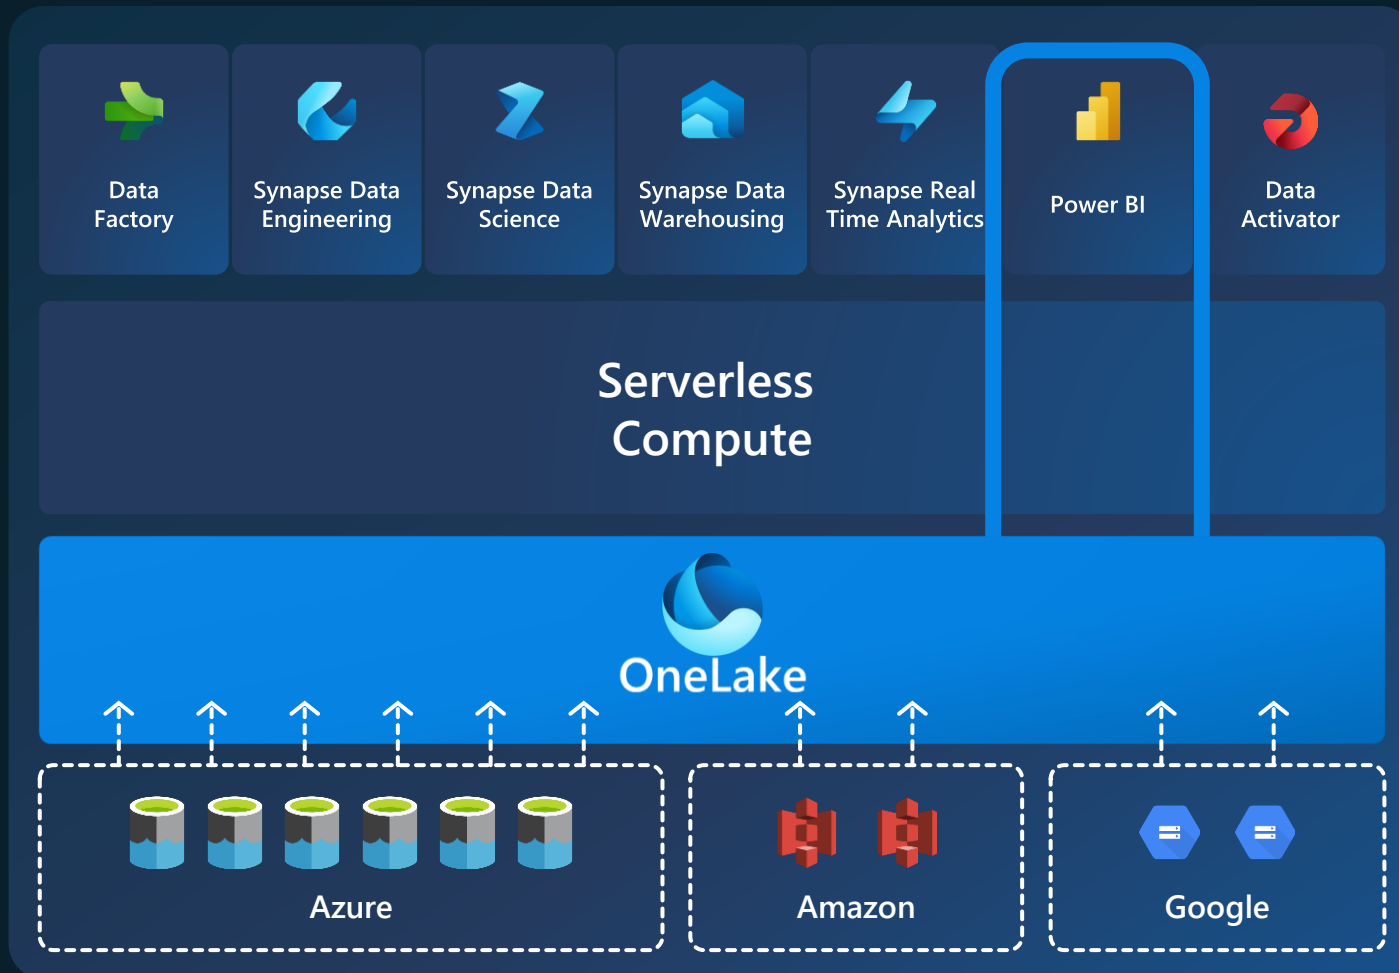
Delta – Parquet, an open standards format, is the storage format for all tabular data in Analytics vNext

Once data is stored in the lake, it is directly accessible by all the engines without needing any import/export

All the compute engines have been fully optimized to work with Delta Parquet as their native format

Shared universal security model is enforced across all the engines

One Copy – Direct Lake



Sharing data in OneLake is as easy as sharing files in OneDrive, removing the needs for data duplication

With **shortcuts**, data throughout OneLake can be composed together without any data movement

Shortcuts also allow instant linking of data already existing in Azure and in other clouds, without any data duplication and movement, making **OneLake the first multi-cloud data lake**

With support for industry standard APIs, OneLake data can be directly accessed by any application or service

Fundamentals

- Only one data source can be used with Direct Lake
- Direct Lake semantic model starts life with no data in memory
- Data is *paged* into semantic model on demand triggered by query
- Tables can have resident and non-resident columns
- Column data can get evicted for multiple reasons
- Direct Lake fallback to SQL Server for *suitable* sub-queries
- “Framing” for data consistency in Power BI reports

Direct Lake limitations (for now)

- No calculated columns or calculated tables
- No composite models
 - Although calculation groups and field parameters are now allowed
- Can only be used with tables, not views
- Can only be used with security defined in the semantic model
- Web authoring experience (or 3rd party tool)
- Not all data types supported
 - No structured data types, binary or GUID columns
 - DateTime relationships not supported
 - String length limited to 4000 characters

Direct Lake prerequisites

SKU Requirements

- Power BI Premium P
- Microsoft Fabric F SKUs only, including Trial

Not supported on:

- Power BI Pro
- Premium Per User
- Power BI Embedded A/EM Skus

Direct Lake prerequisites



Lakehouse



Warehouse

Direct Lake



Anatomy of Parquet

Why Parquet?

- Open source/open data format
- Column-oriented format is optimized for data storage and retrieval
- Efficient data compression and encoding especially data in bulk
- Is lingua franca for data storage format
 - Databricks, Microsoft - delta lake and parquet
 - Snowflake - iceberg and parquet/orc

Anatomy of a Parquet File

CSV, XML, JSON..... Parquet

| StoreID | | DateTime | | ProductID | | Value |
|---------|---|------------|---|-----------|---|-------|
| StoreA | , | 2023-01-01 | , | SKU001 | , | 10 |
| StoreA | , | 2023-01-02 | , | SKU001 | , | 15 |
| StoreA | , | 2023-01-03 | , | SKU001 | , | 12 |

Anatomy of a Parquet File

CSV, **XML**, JSON..... Parquet

```
<sale>
```

```
  <StoreID>StoreA</StoreID>
```

```
  <DateTime>2023-01-01</DateTime>
```

```
  <ProductID>SKU001</ProductID>
```

```
  <Value>10</Value>
```

```
</sales>
```

```
<sale> ... </sale>
```

Anatomy of a Parquet File

CSV, XML, **JSON**..... Parquet

```
{sales[
  {
    StoreID: "StoreA" ,
    DateTime: "2023-01-01" ,
    ProductID: "SKU001" ,
    Value:10
  },
  {...}
]}
```

Anatomy of a Parquet File

CSV, XML, JSON..... **Parquet**

Header:

RowGroup1:

StoreID : StoreA, StoreA, StoreA

DateTime : 2023-01-01, 2023-01-02, 2023-01-03

ProductID: SKU001, SKU001, SKU001

Value : 10, 15, 12

RowGroup2:

...

Footer:

Anatomy of a Parquet File – Dictionary IDs

CSV, XML, JSON..... **Parquet**

Header:

RowGroup1:

StoreID : 1, 1, 1

DateTime : 1, 2, 3

ProductID : 1, 1, 1

Value : 1, 2, 3

RowGroup2:

...


Footer:

V-Order

V-Order

- V-Order is a Microsoft-proprietary optimisation for writing data in parquet files (as used in Delta tables)
- V-Order uses the same algorithms used by Power BI Import mode semantic models to compress data
- V-Ordered Delta tables can be read by any tool that can read Delta
- Direct Lake will perform better on V-Ordered Delta tables
- Direct Lake will work on all Delta tables, even without V-Order

VORDER compressed Parquet



x3.2

Demo

Framing

Framing

- What is framing
 - "point in time" way of tracking what data can be queried by Direct Lake
- Why is this important
 - Data consistency for some Power BI Reports
 - Delta-lake data is transient for many reasons
- ETL Process
 - Ingest data to delta lake tables
 - Transform as needed using preferred tool
 - When ready, perform *Framing* operation on semantic model
- Framing is near instant and acts like a cursor
 - Determines the set of .parquet files to use/ignore for *transcoding* operations

Framing



Framing - Options

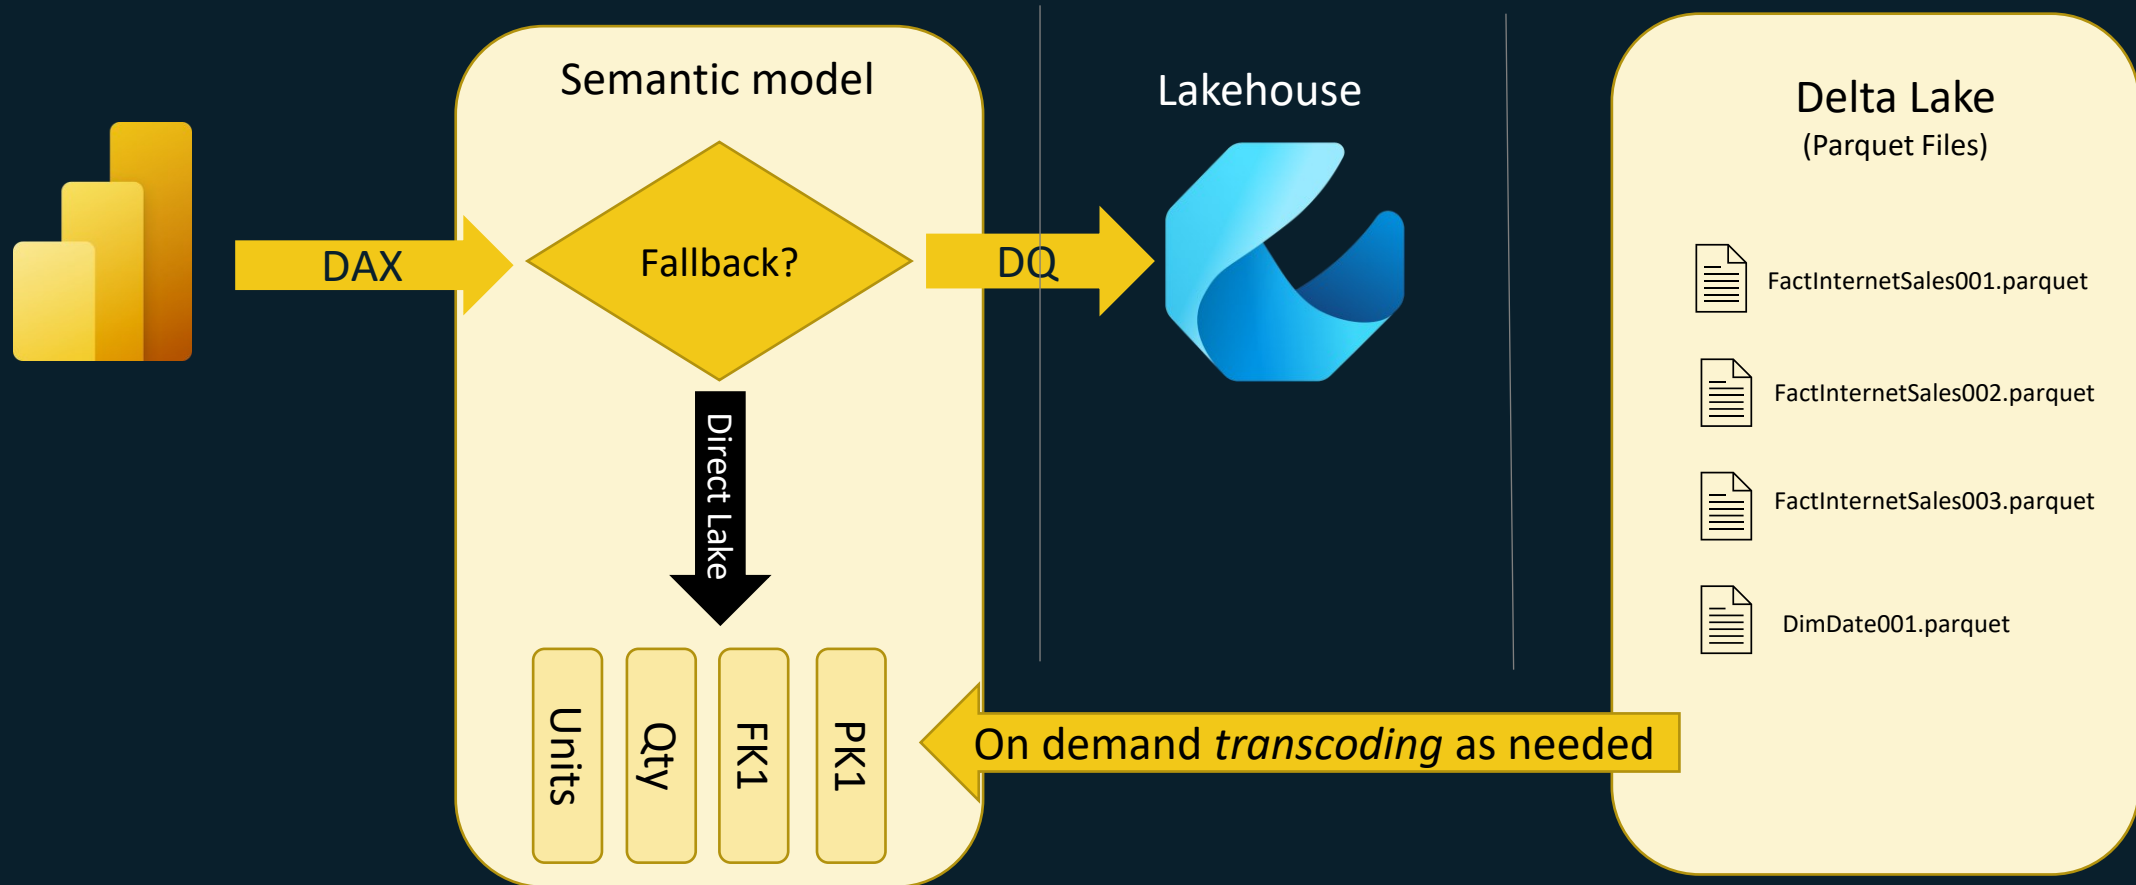
- Automatic
 - Default - can be turned off
 - Triggered each time Delta table gets modified
- Via Fabric Service
 - Manually by refreshing the semantic model
 - Configure a schedule
- Via Notebook
 - Use Semantic-link to call reframe using native method
 - Execute_tmsl for fine grain reframing
 - Consider cache-warming as option

Framing - Options

- SSMS (TMSL)
- Rest API
- Pipeline
- Notebooks (sempy)
- Power Automate etc.

DirectQuery Fallback

Fallback to DirectQuery



Fallback to DirectQuery

- You are using features that prevent Direct Lake
- Views are not allowed because they don't have corresponding tables stored in a Lakehouse
- RLS or OLS is defined in a Warehouse
 - Security rules take high priority when defined

Fallback to DirectQuery – data volumes

- There are limits on how much data can be used for Direct Lake
- These limits vary by capacity SKU size
- If you exceed these limits, Direct Lake will use DirectQuery
 - Query performance may be noticeably worse
- Fabric checks limits during reframing process
- Can be turned On/Off using Direct Lake Behaviour property

Fallback guardrails

| Fabric/Power BI SKUs | Parquet files per table | Row groups per table | Rows per table (millions) | Max model size on disk/OneLake (GB) | Max memory (GB) |
|----------------------|-------------------------|----------------------|---------------------------|-------------------------------------|-----------------|
| F2 | 1,000 | 1,000 | 300 | 10 | 3 |
| F4 | 1,000 | 1,000 | 300 | 10 | 3 |
| F8 | 1,000 | 1,000 | 300 | 10 | 3 |
| F16 | 1,000 | 1,000 | 300 | 20 | 5 |
| F32 | 1,000 | 1,000 | 300 | 40 | 10 |
| F64/FT1/P1 | 5,000 | 5,000 | 1,500 | Unlimited | 25 |
| F128/P2 | 5,000 | 5,000 | 3,000 | Unlimited | 50 |
| F256/P3 | 5,000 | 5,000 | 6,000 | Unlimited | 100 |
| F512/P4 | 10,000 | 10,000 | 12,000 | Unlimited | 200 |
| F1024/P5 | 10,000 | 10,000 | 24,000 | Unlimited | 400 |
| F2048 | 10,000 | 10,000 | 24,000 | Unlimited | 400 |

[Learn about Direct Lake in Power BI and Microsoft Fabric](#)

Detecting fallback to DirectQuery

- Performance Analyzer, Profiler traces and/or Log Analytics will show what happens for individual queries
- Limits on data volumes can be checked with Python notebooks (Delta Analyzer) and in some cases DMVs

Controlling fallback to DirectQuery

- The **DirectLakeBehavior** property sets fallback behaviour
- Automatic (default): allows fallback to DirectQuery if data can't be loaded into memory
- DirectLakeOnly: allows use of DirectLake but prevents fallback and returns an error instead of using DirectQuery
- DirectQueryOnly: forces all queries to use DirectQuery mode

Demo

Direct Lake Behavior

Performance

Performance considerations

- Reframing
- Cold Cache
- Warm Cache
- SQL Fallback

Performance considerations – Reframing

- Time to evict columns and load certain objects
- Loads Delta metadata and some metadata from parquet files

Performance considerations – Cold Cache

- The time needed to page data into memory from One Lake
- Number/layout of data across Parquet files
- Cache warming tricks

Performance considerations – Warm Cache

- Query Plans
 - Direct Lake Behaviour property
 - Other optimisations
- Encoding
 - All data is HASH encoded – no option to use VALUE encoding
- Segment data profile
 - Number and layout of data within segments can impact scan performance
 - Depends greatly on filters used per query

Performance numbers – sample model

- With V-Order
- No V-Order
- Column partitioned by Date
- V-Order and Z-Order

Performance - some numbers

| | V-Order | No V-Order | Partitioned by DateKey | V-Order & Z-Order |
|----------------|---------------|---------------|------------------------|-------------------|
| Rows | 1,000,000,000 | 1,000,000,000 | 1,000,000,000 | 1,000,000,000 |
| Columns | 10 | 10 | 10 | 10 |
| V-Order | TRUE | | TRUE | TRUE |
| Z-Order | | | | DateKey |
| Parquet Size | 7.1GB | 11.6GB | 8.4GB | 6.9GB |
| Files | 14 | 200 | 807 | 6 |
| Row Groups | 26 | 200 | 807 | 24 |
| | | | | |
| Model Size | | | | |
| Data | 7.1GB | 14.9GB | 6.6GB | 6.9GB |
| Total (memory) | 9.6GB | 17.5GB | 9.1GB | 9.4GB |

Performance numbers – Cold cache

| | V-Order | No V-Order | Partitioned by DateKey | V-Order & Z-Order |
|--------|---------|------------|---------------------------|----------------------|
| Test 1 | 2m 24s | 7m 35s | 7m 16s | 2m 18s |
| Test 2 | 2m 26s | 8m 34s | 7m 30s | 2m 17s |
| Test 3 | 2m 27s | 7m 46s | 7m 27s | 2m 18s |

| Parquet Size | 7.1GB | 11.6GB | 8.4GB | 6.9GB |
|--------------|-------|--------|-------|-------|
| Files | 14 | 200 | 807 | 6 |
| Row Groups | 26 | 200 | 807 | 24 |
| | | | | |
| Model Size | | | | |
| Data | 7.1GB | 14.9GB | 6.6GB | 6.9GB |
| Total | 9.6GB | 17.5GB | 9.1GB | 9.4GB |

Performance numbers – Warm Cache

Group by Weekday
Filter on Category
and Month
Sum Quantity

Query 1

Group by Month
Filter on Category
Compare using
PREVIOUSMONTH
Sum Quantity

Query 2

Filter by Month
All countries ranked
by distinct User ID

Query 3

Total Time

FE CPU

| | V-Order | No V-Order | Partitioned by DateKey | V-Order & Z-Order |
|--|---------|------------|---------------------------|----------------------|
| | | | | |
| | 223 | 863 | 47 | 203 |
| | 1,449 | 12,840 | 550 | 1,148 |
| | | | | |
| | | | | |
| | 1,594 | 2,891 | 94 | 1,379 |
| | 11,890 | 39,125 | 1,070 | 10,754 |
| | | | | |
| | | | | |
| | 4,817 | 4,129 | 5,851 | 10,845 |
| | 29,937 | 43,933 | 42,523 | 18,867 |

Summary

- Will my reports run faster with Direct Lake?
- Do I have to use Direct Lake with Fabric?
- Incremental Refresh?
- Aggregations?

Thanks